


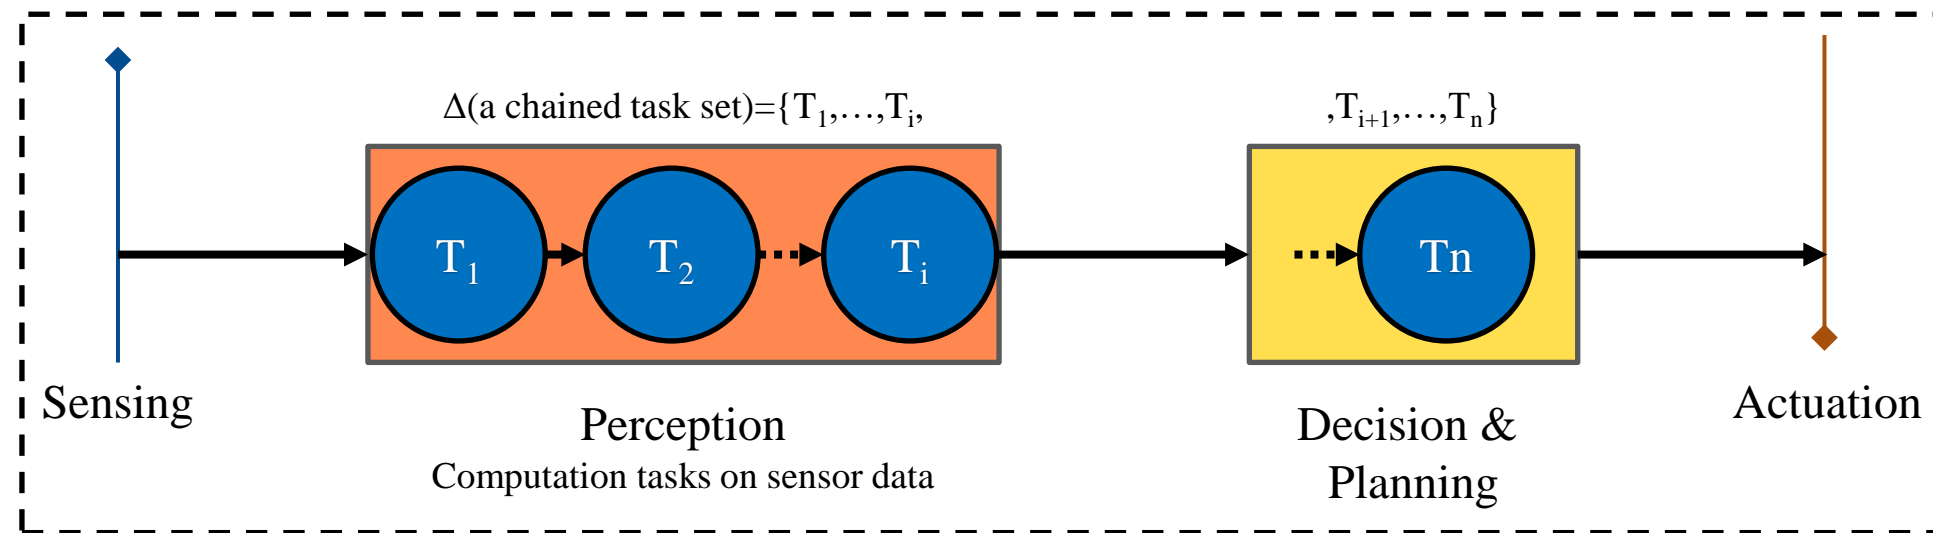
Predictable Data-driven Resource Management: an Implementation using Autoware on Autonomous Platforms

Soroush Bateni and Cong Liu

Department of Computer Science
Erik Jonsson School of Engineering & Computer Science
The University of Texas at Dallas


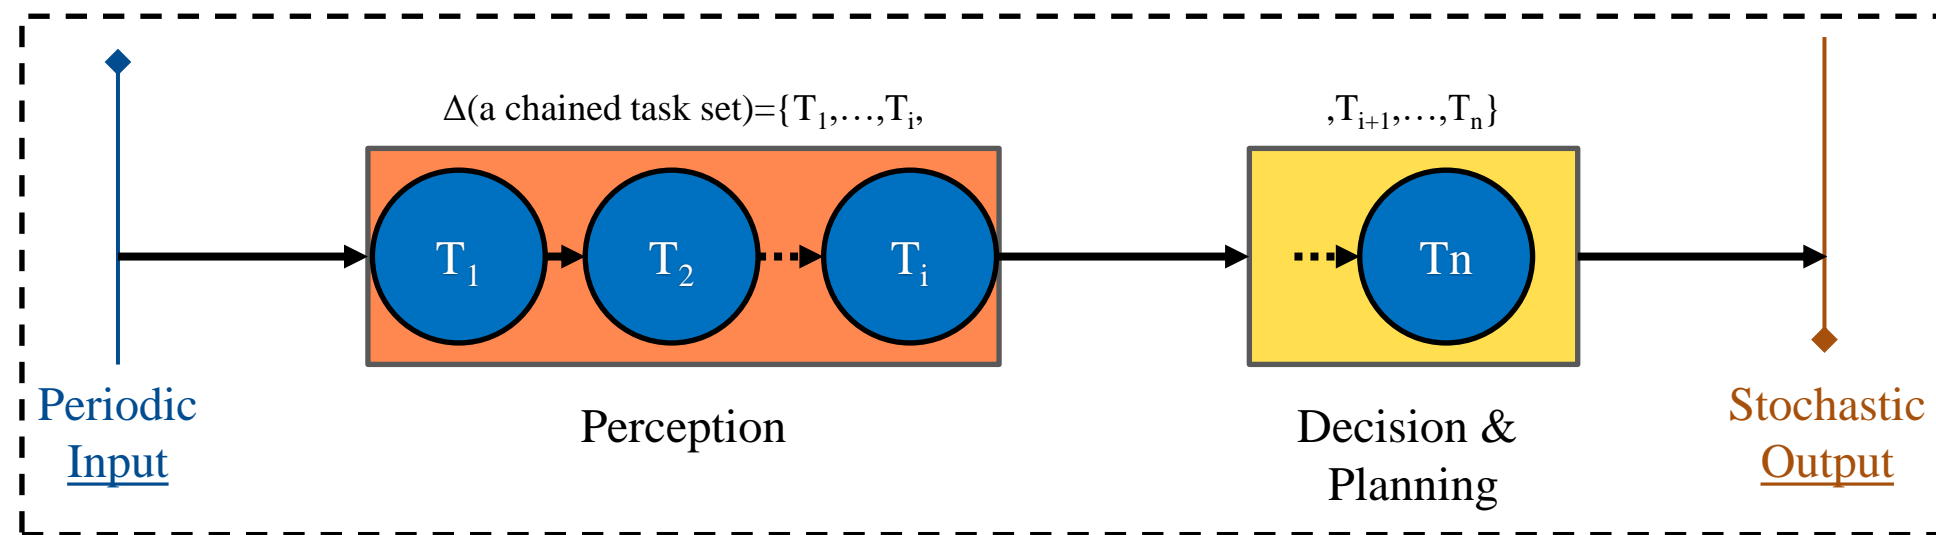


Autonomous Embedded Systems




Embedded Platform

Autonomous Embedded Systems




Embedded Platform

Autonomous Embedded Systems


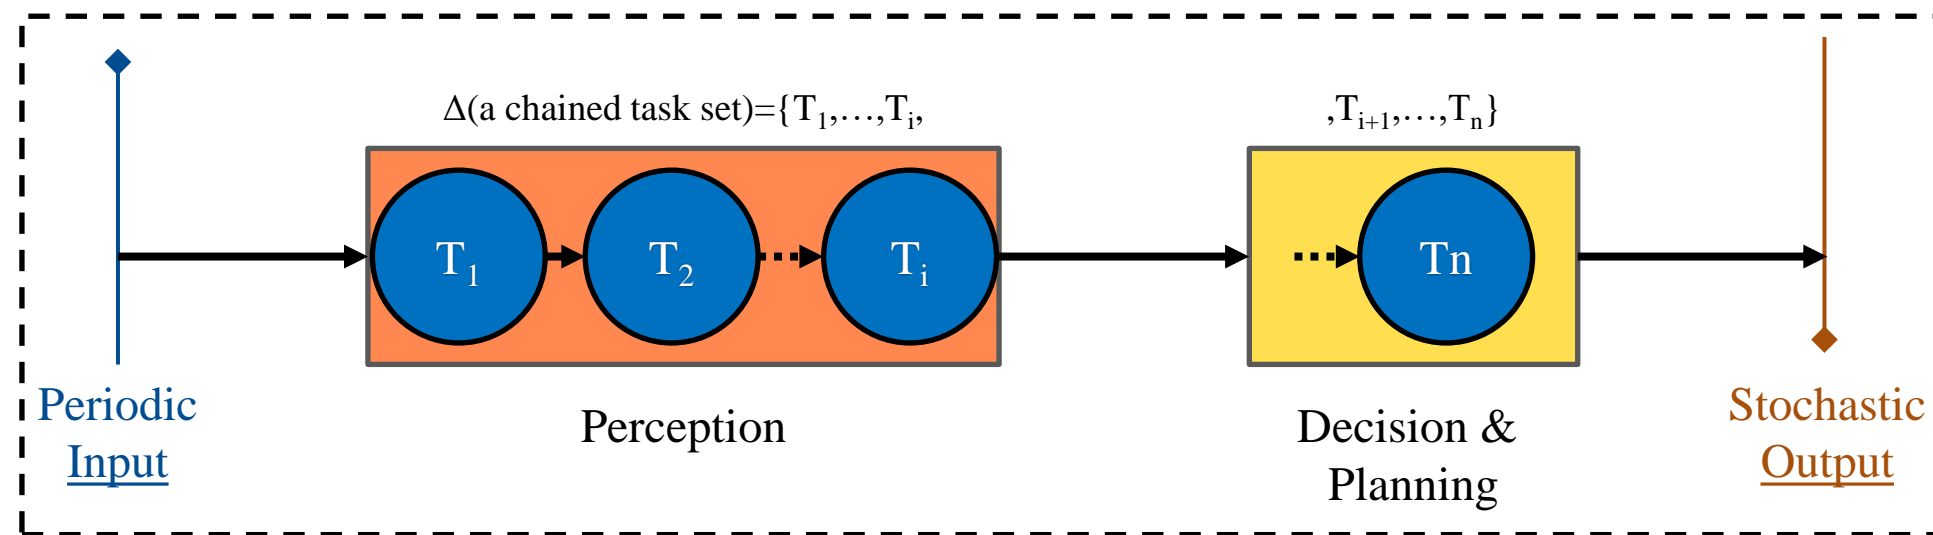


For example, a camera can release images every 33ms.



Embedded Platform

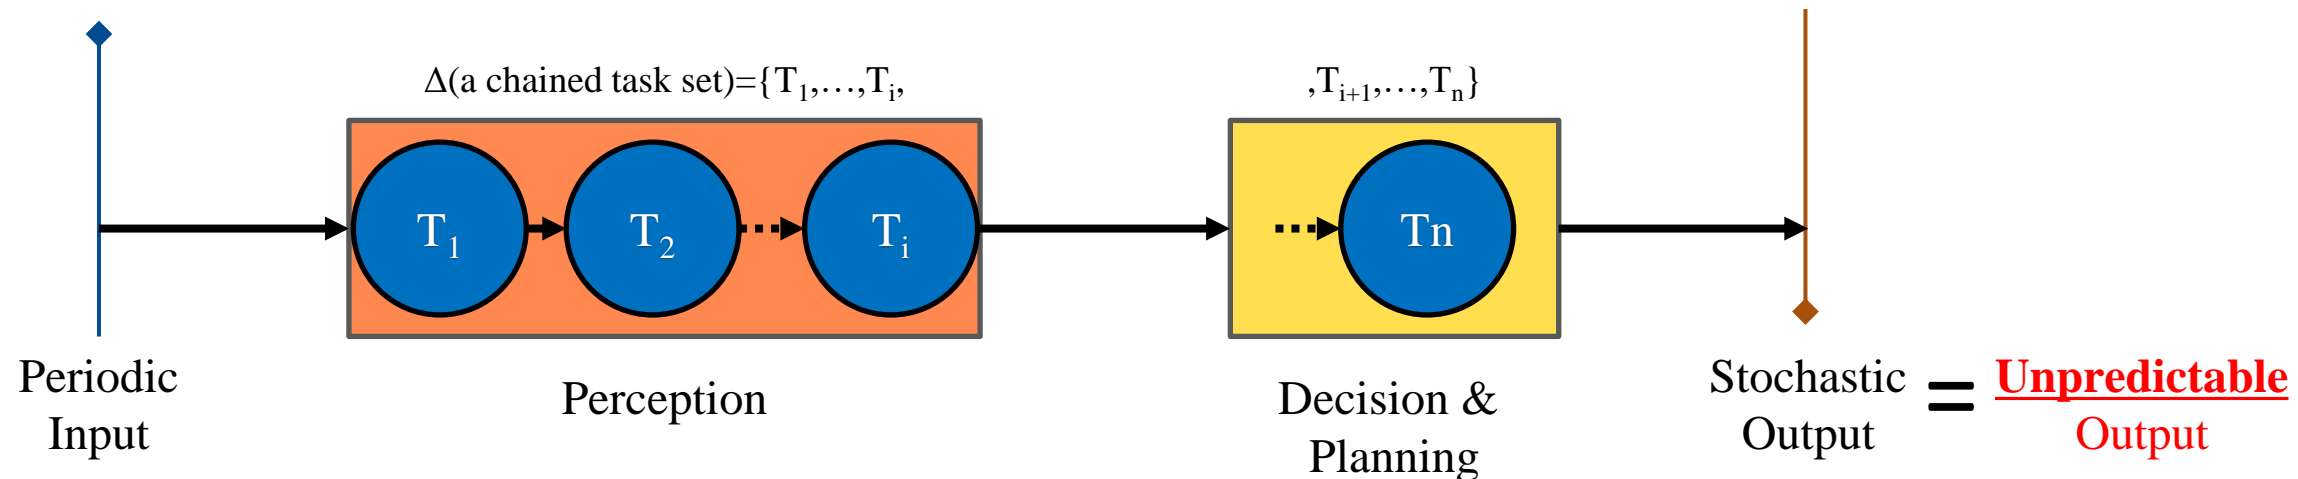
Autonomous Embedded Systems



Embedded Platform

Output depends on when the task set finishes.

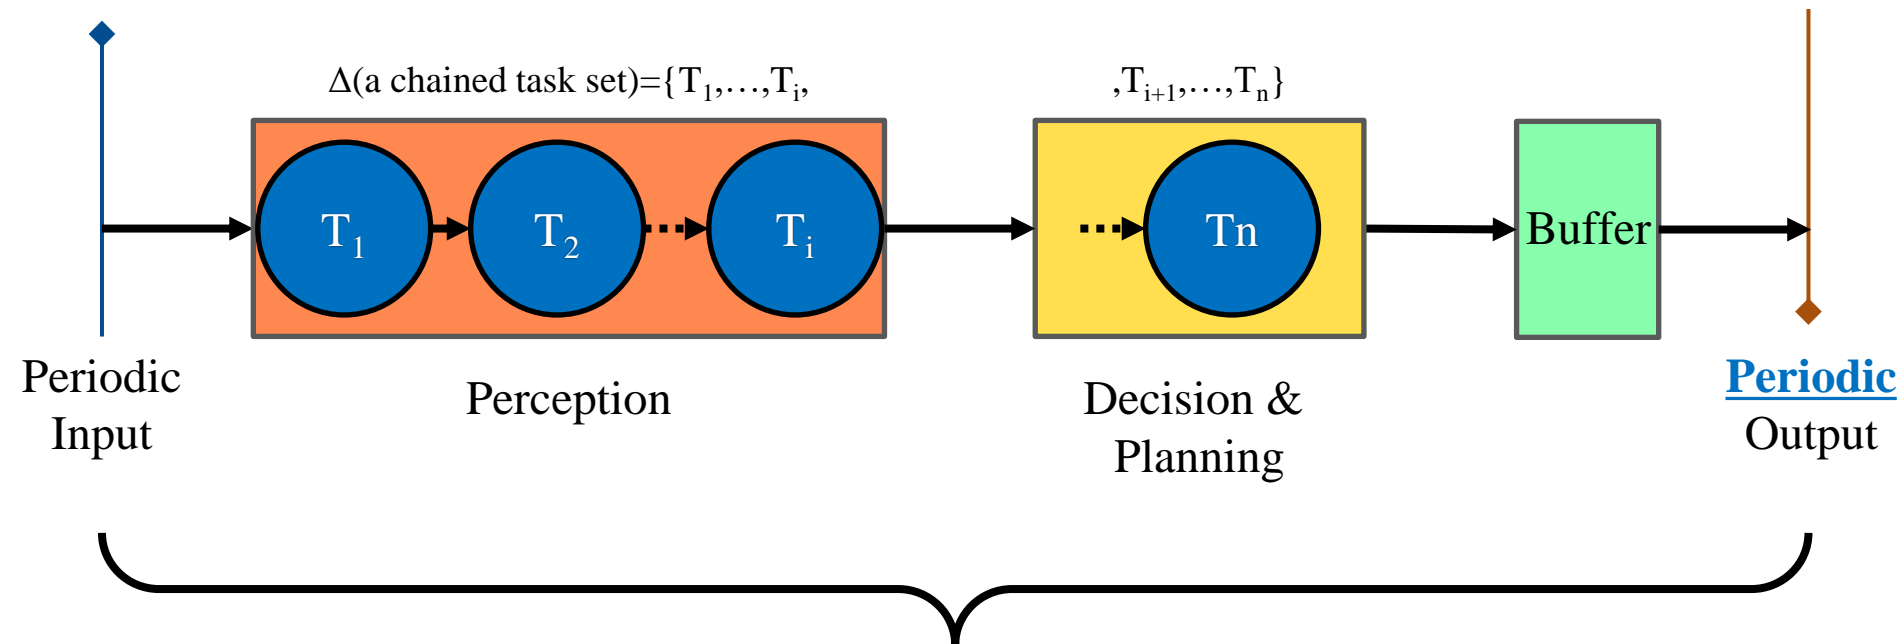
Logical Execution Time: Motivation



- ❖ Unpredictable output is undesirable in AES¹
 - ❖ Harder to verify and certify
 - ❖ Harder to schedule
 - ❖ Can be classified as unexpected behavior pattern, specifically prohibited in most standards
 - A self-driving car sometimes takes 30ms to react and sometimes 25ms

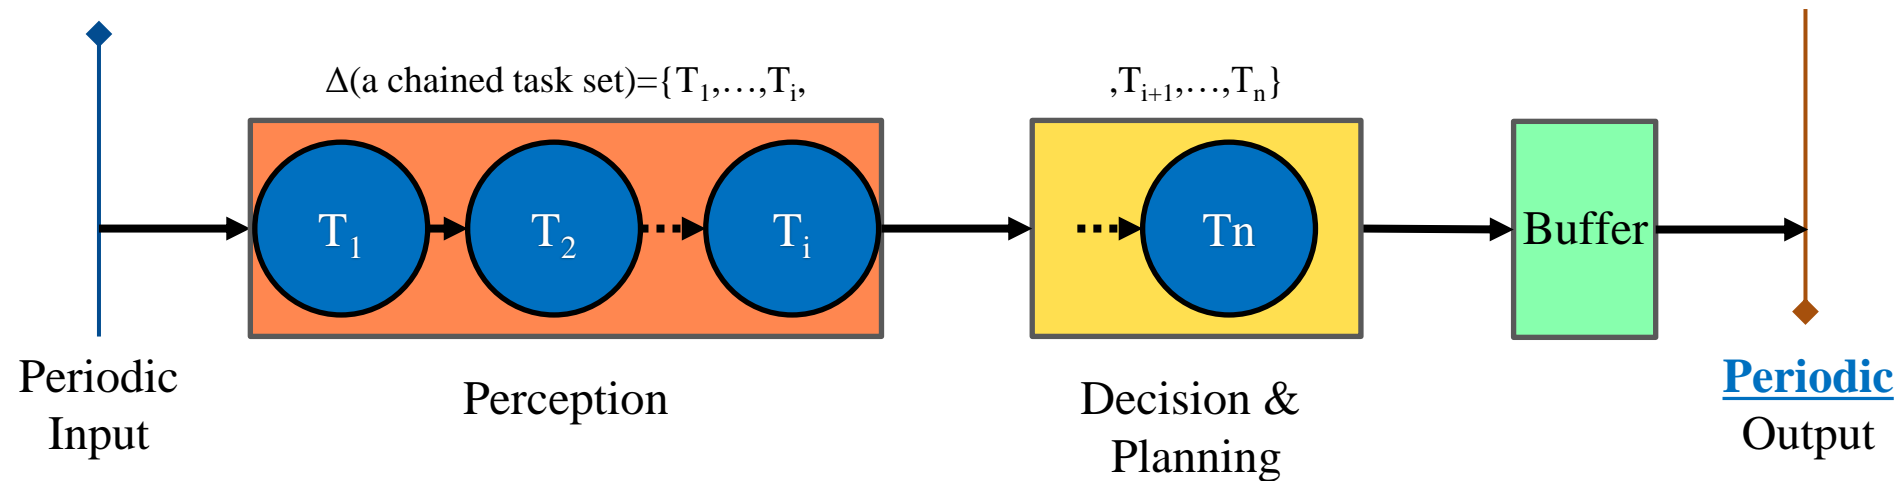
¹ Kirsch, Christoph M., and Ana Sokolova. "The logical execution time paradigm." Advances in Real-Time Systems. Springer, Berlin, Heidelberg, 2012. 103-120.

Logical Execution Time: **Intuition**



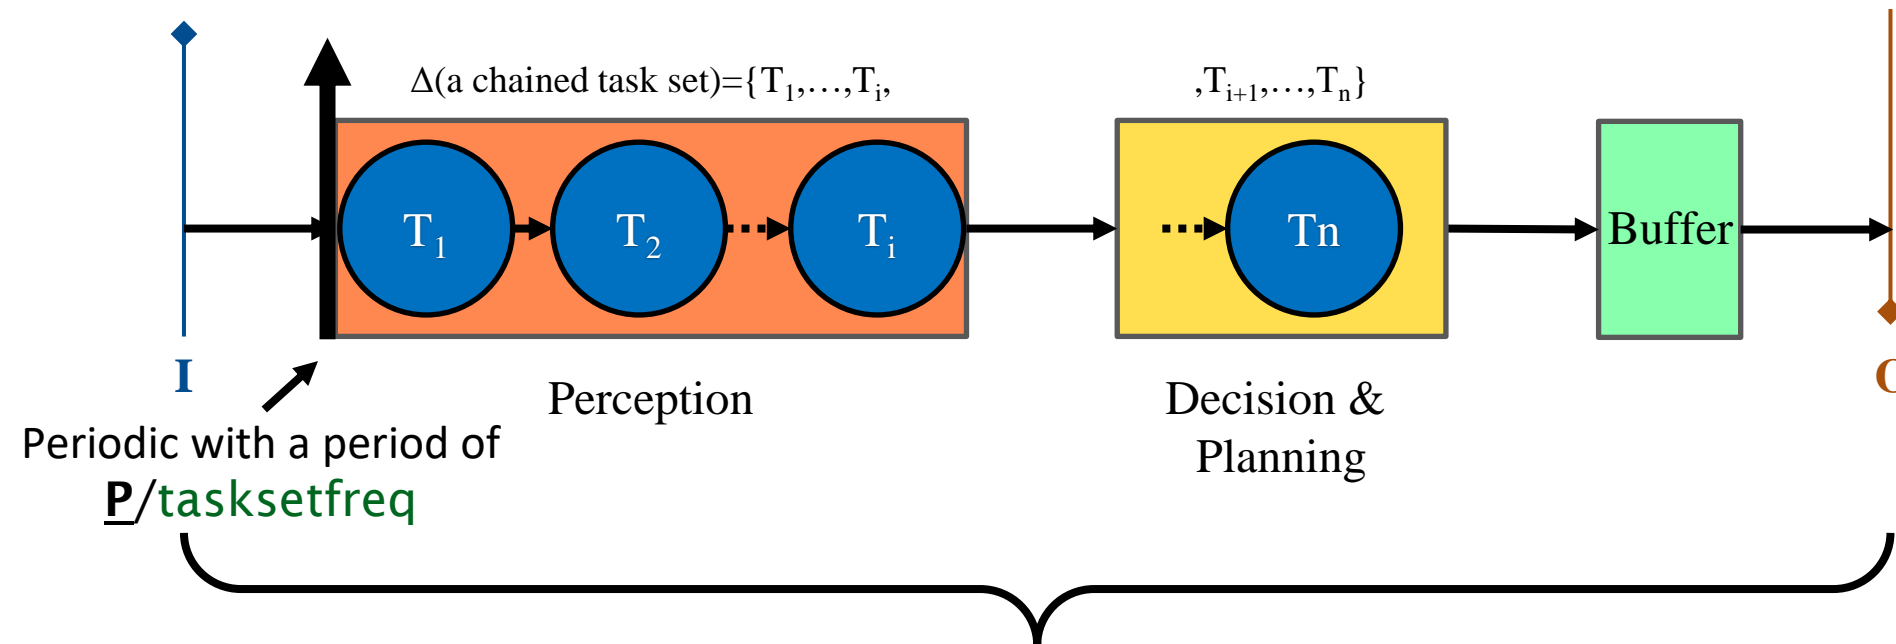
Task system always takes **LET (fixed time)** to execute.

Logical Execution Time: **Intuition**



The periodic output must be met and not violated.

Logical Execution Time: Example

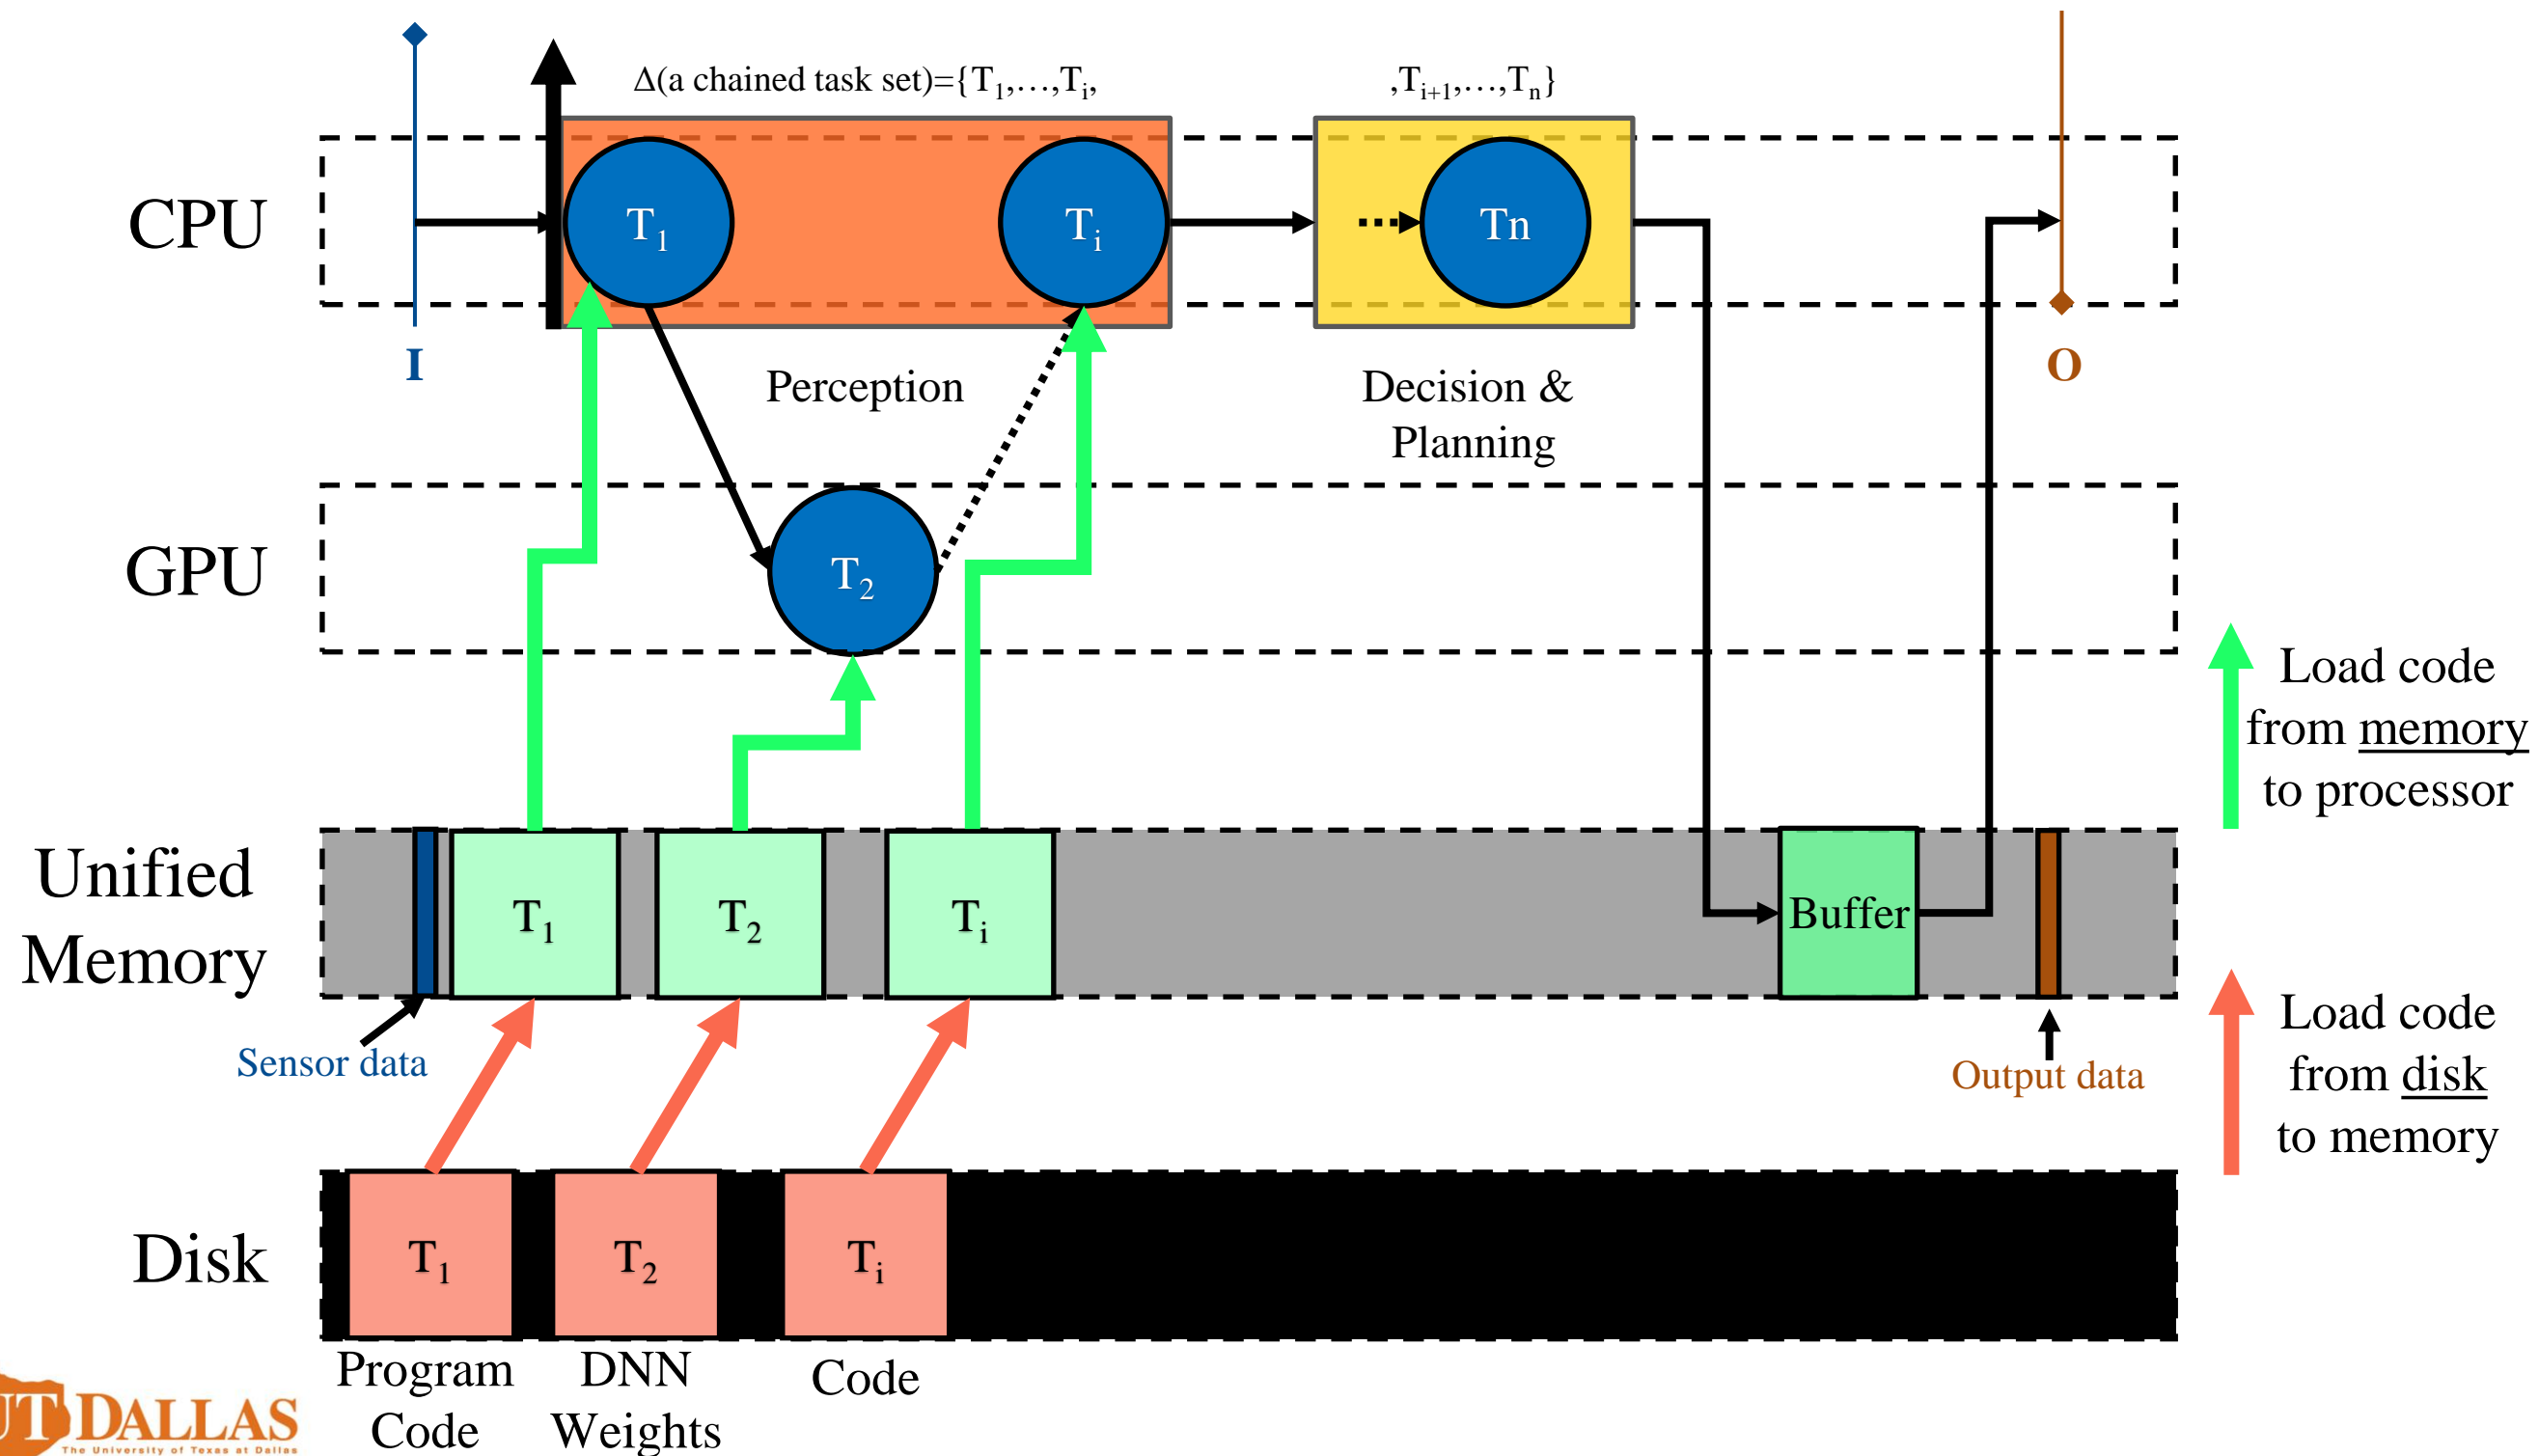


Example Abstract Code

```

mode m() Period 20 {
    inputfreq  do  $P_I(d_I)$       ➤  $\underline{P}/\text{inputfreq}$ : input interval (period)
    tasksetfreq do  $\Delta(d_\Delta)$  ➤  $\underline{P}/\text{tasksetfreq}$ : period of task set  $\Delta$ 
    outputfreq do  $P_O(d_O)$     ➤  $\underline{P}/\text{outputfreq}$ : output interval (period)
}
    
```

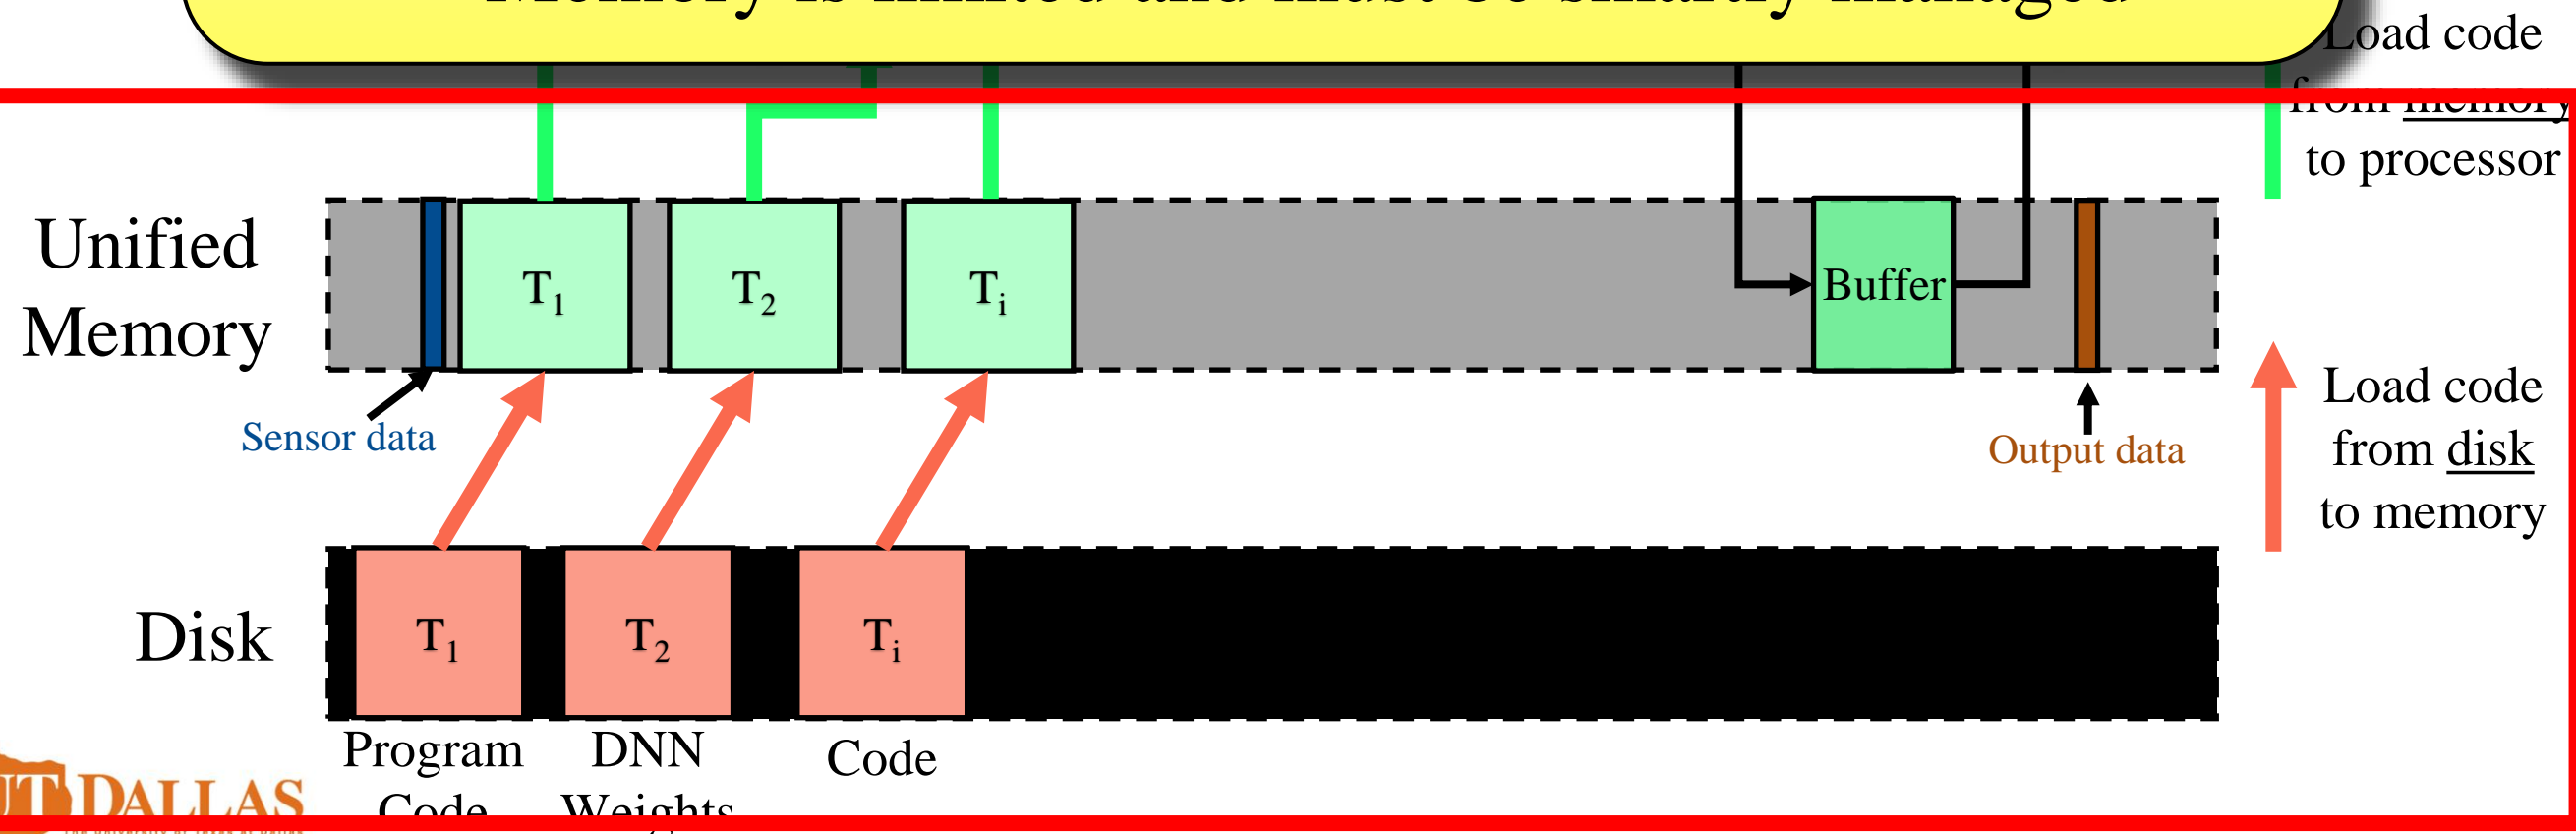
Integrated CPU/GPU Architecture



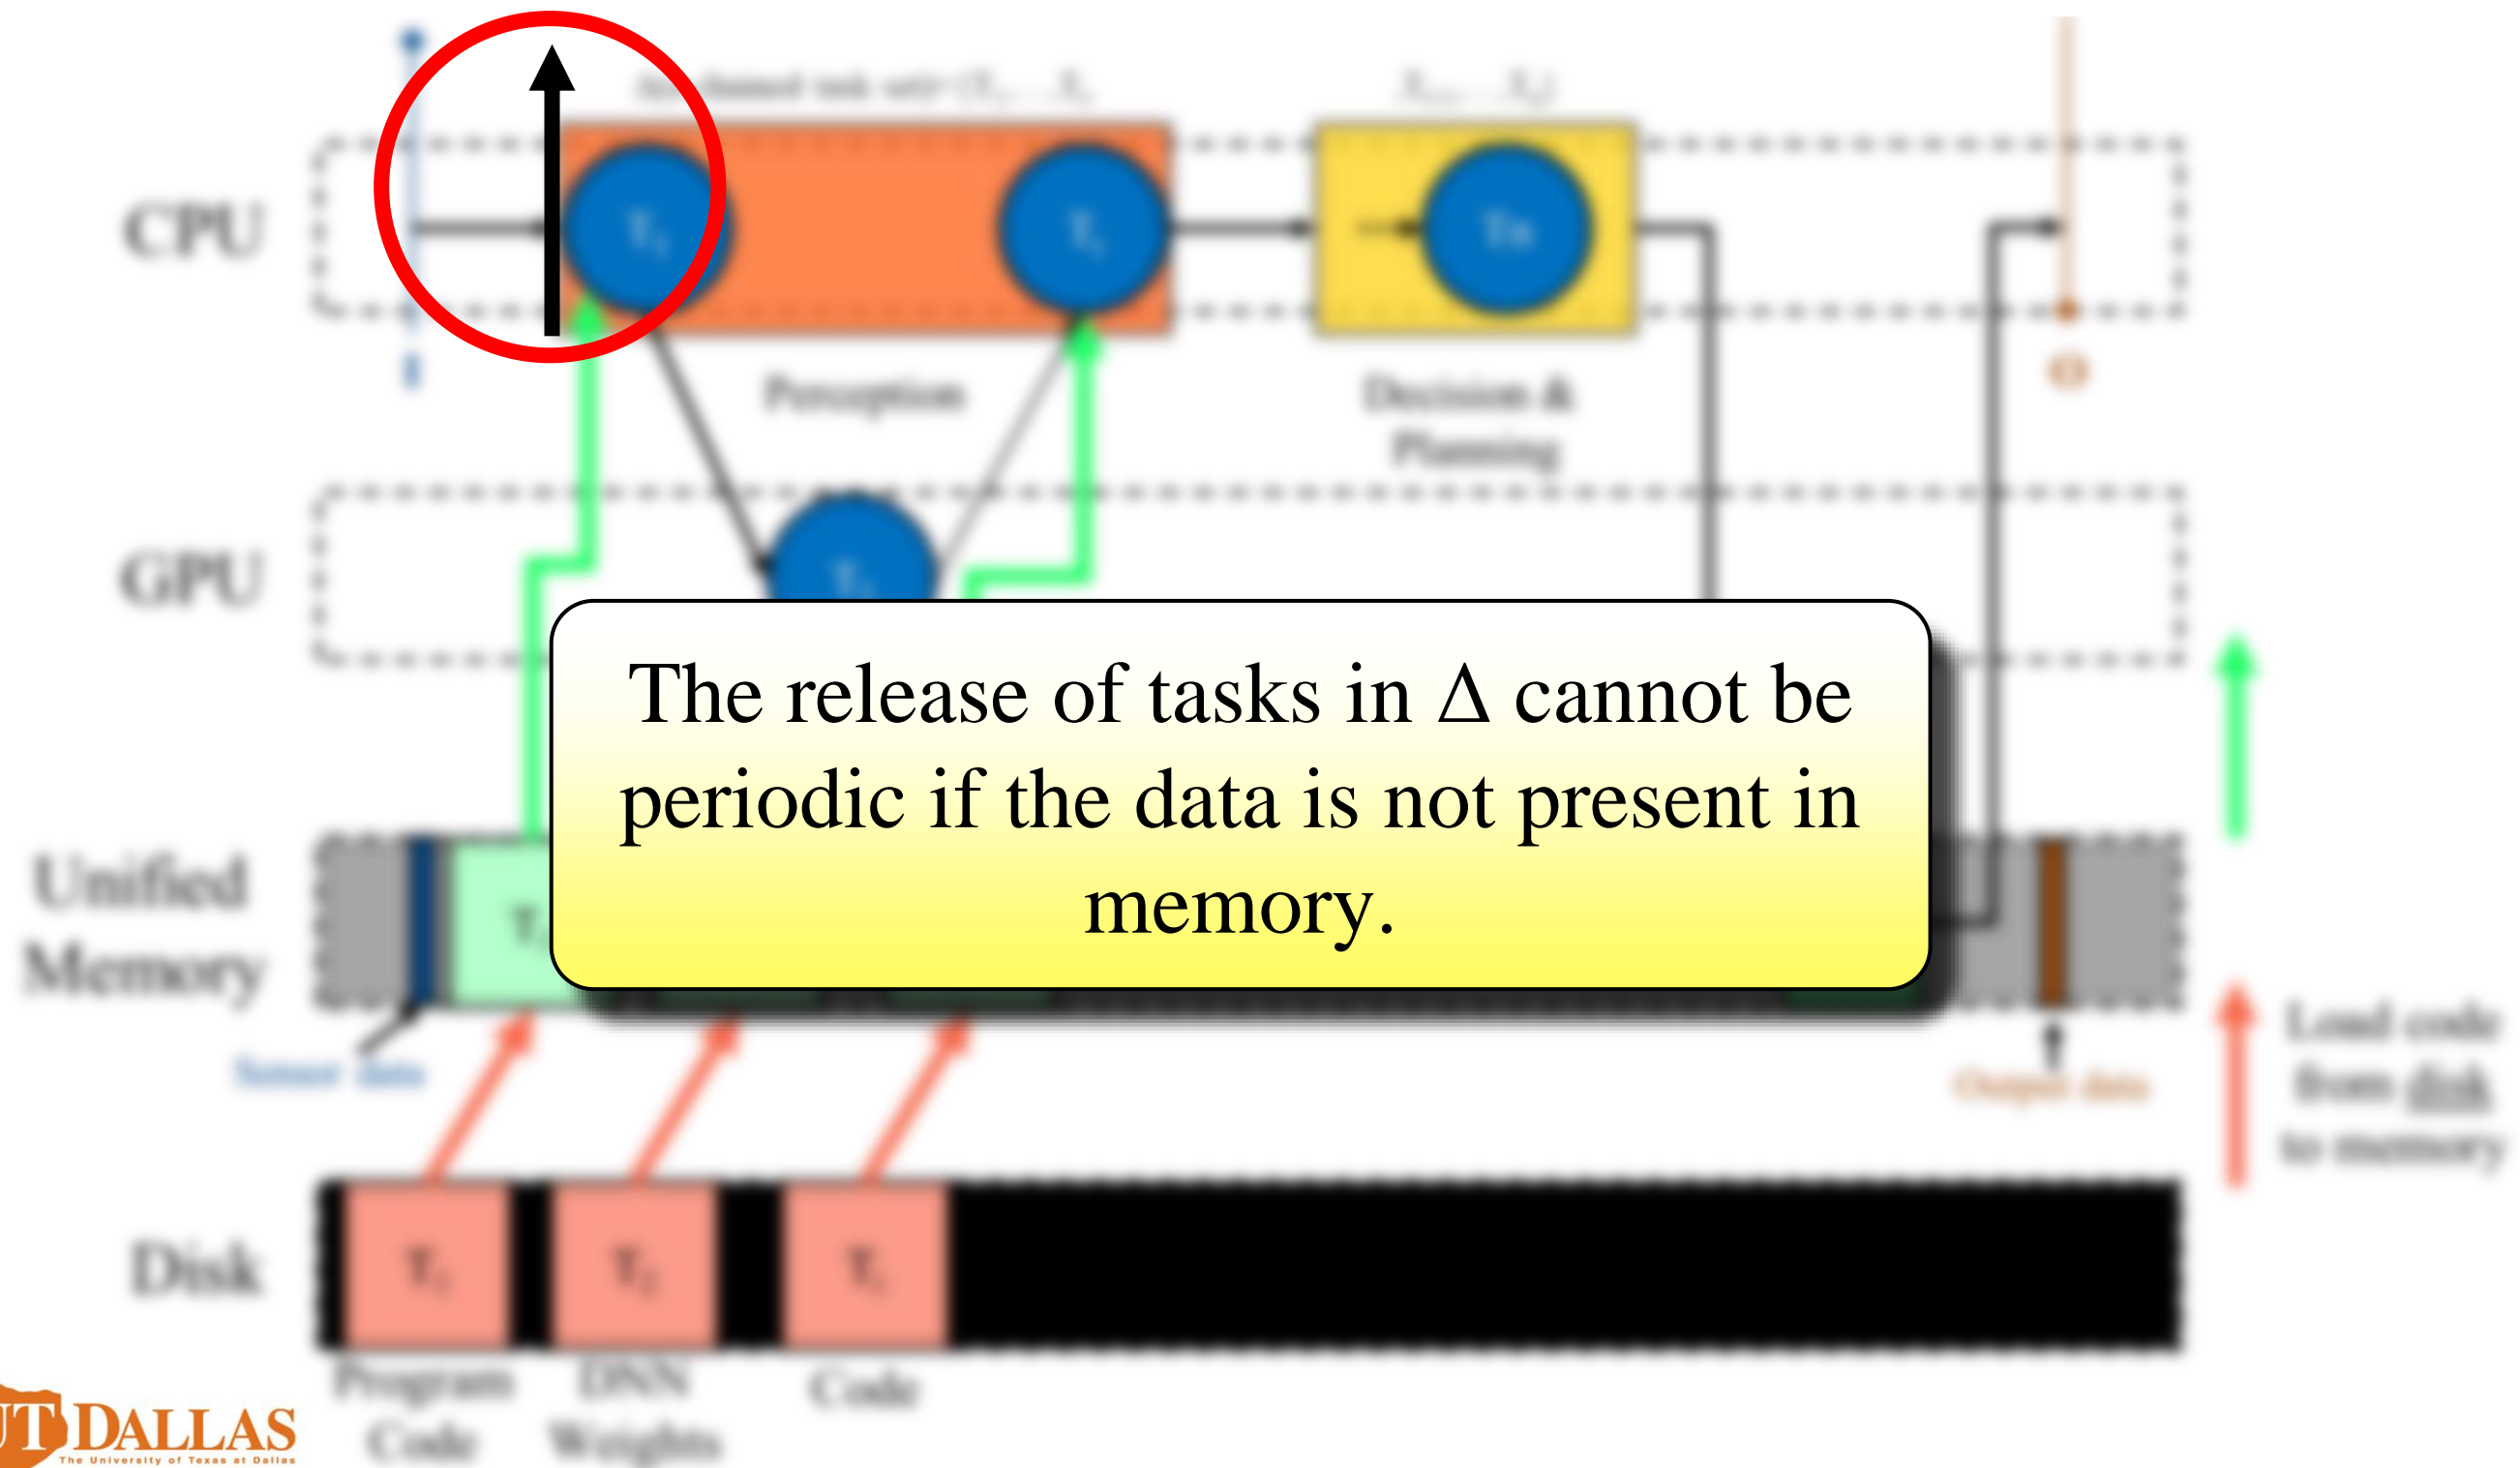
Integrated CPU/GPU Architecture

Focus

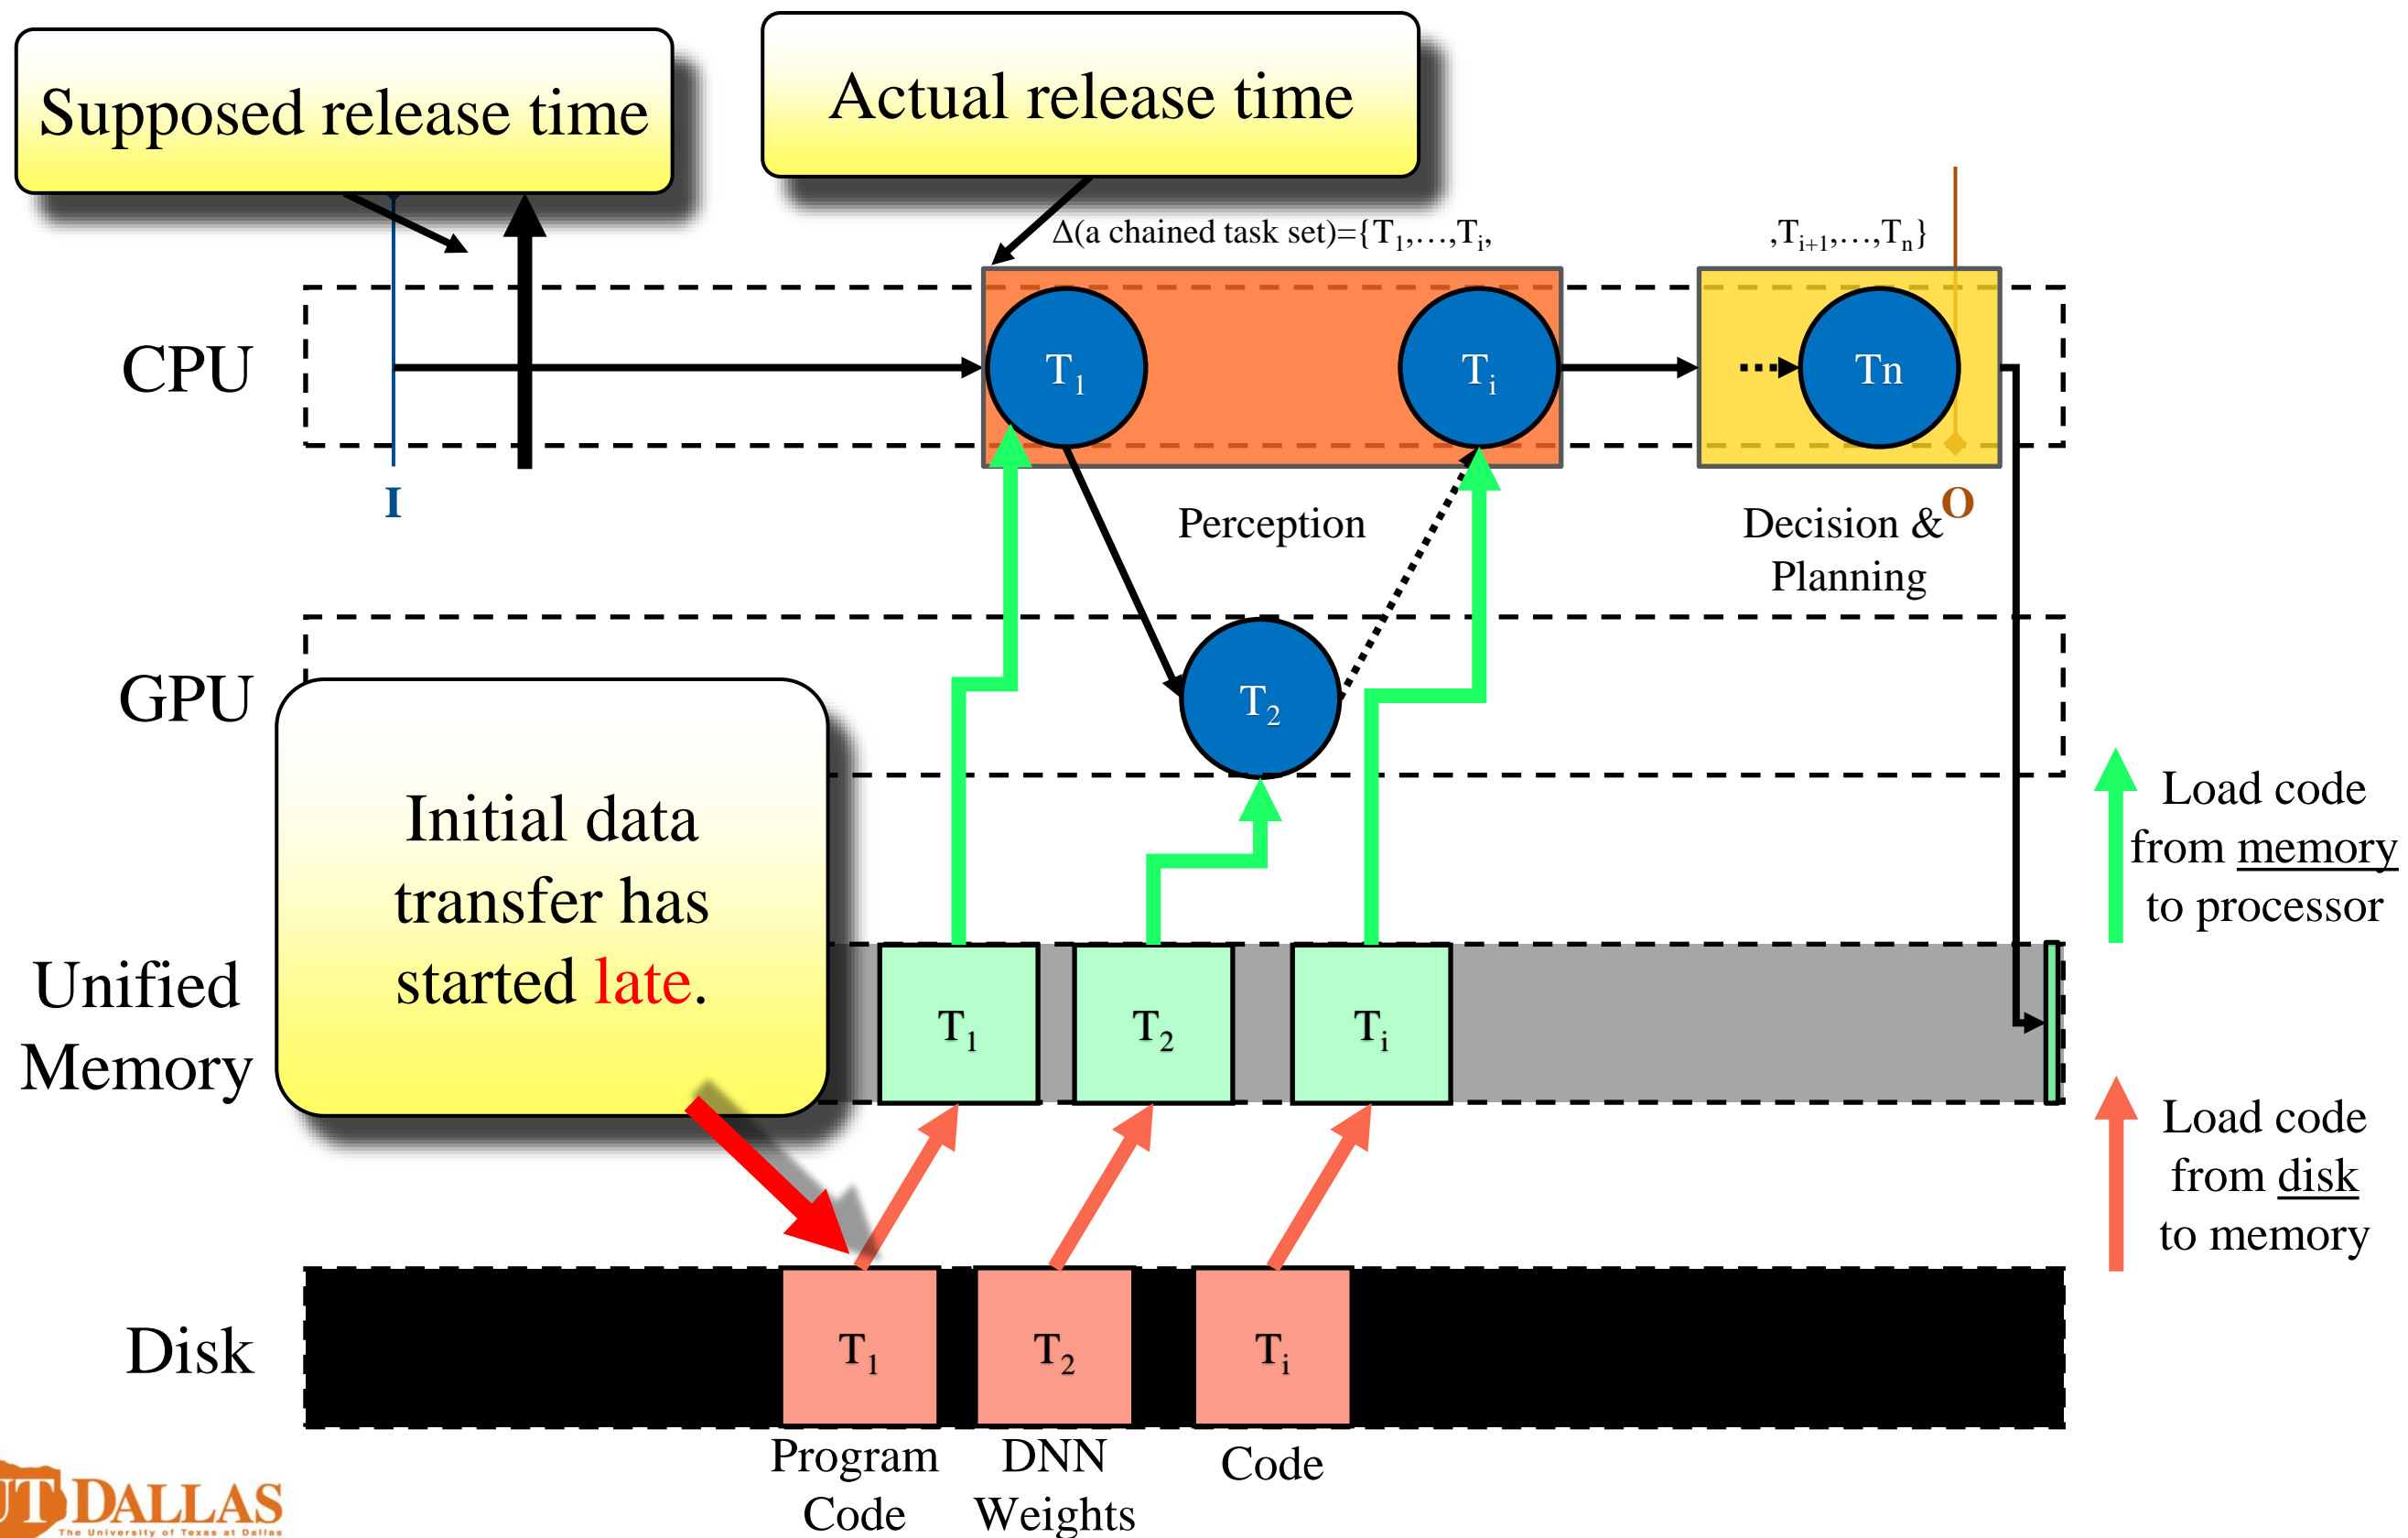
- Sensor data is small compared to program code.
- Loading code from memory to processor cache in a predictable manner has been extensively explored before.
- Memory is limited and must be smartly managed




Logical Execution Time: **Flaw** under Integrated CPU/GPU Architecture



Temporal Data Availability Problem



Spatial Data Availability Problem



The diagram illustrates a task chain $\Delta(a \text{ chained task set}) = \{T_1, \dots, T_i, T_{i+1}, \dots, T_n\}$ across four layers:

- CPU:** Contains tasks T_1 , T_i , and T_n . Task T_1 is highlighted in orange, and T_n is in yellow. A blue arrow labeled I points to T_1 . A black arrow connects T_1 to T_i , and another connects T_i to T_n . A dotted arrow labeled "Perception" points from T_1 to T_2 in the GPU layer. A black arrow labeled "Decision & Planning" points from T_i to T_n .
- GPU:** Contains task T_2 . A yellow box labeled "Subsequent data transfer has started **late**." is positioned between the CPU and GPU layers, with a black arrow pointing from T_1 to it. Green arrows point from T_2 to T_1 and from T_2 to T_i .
- Unified Memory:** Contains a buffer and data blocks for T_2 and T_i . A yellow box labeled "Subsequent data transfer has started **late**." is also present here. A blue arrow labeled "Sensor data" points to the start of the buffer. A black arrow labeled "Output data" points from the end of the buffer. Green arrows point from the data blocks in memory to the corresponding tasks in the GPU.
- Disk:** Contains blocks for "Program Code" (T_1), "DNN Weights" (T_2), and "Code" (T_i). Red arrows point from these blocks to the Unified Memory layer.

Annotations on the right side of the diagram:

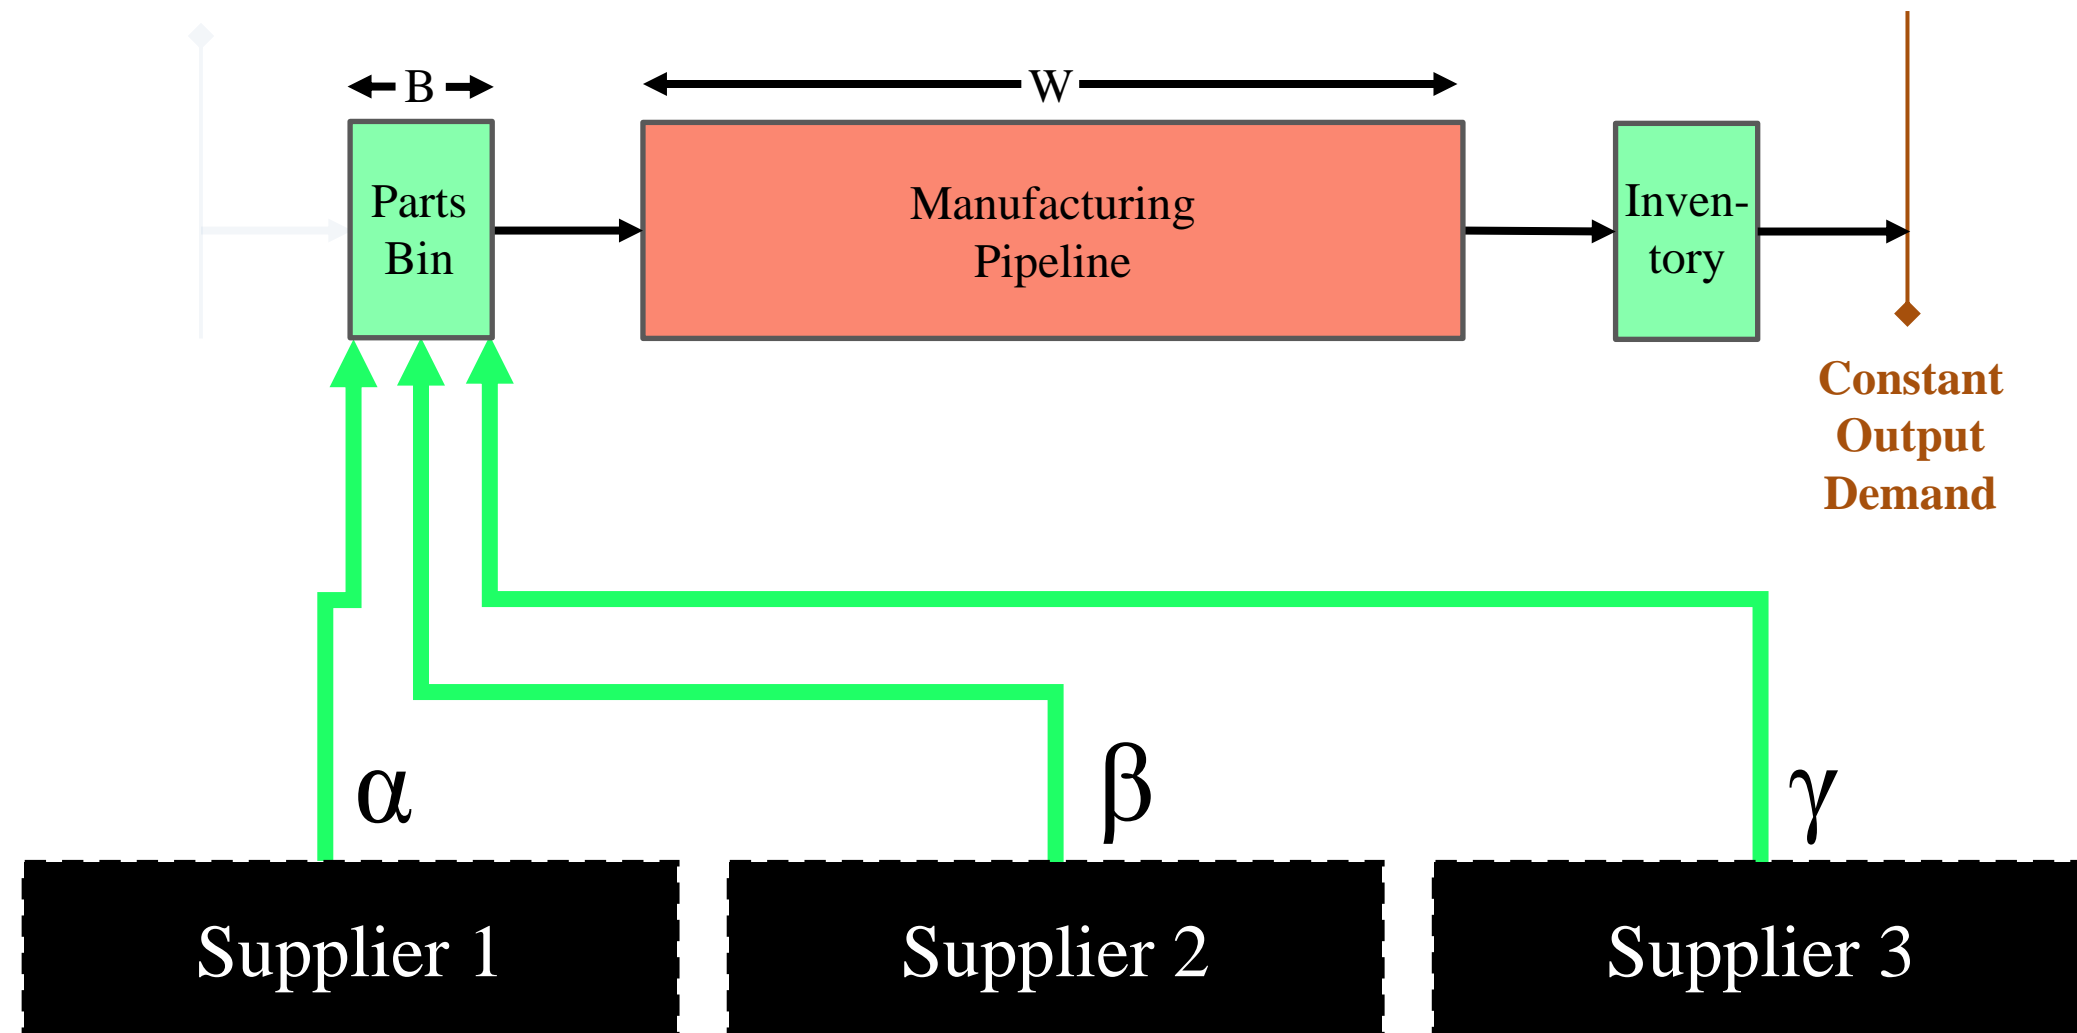
- Green arrow: "Load code from memory to processor"
- Red arrow: "Load code from disk to memory"

UT DALLAS
The University of Texas at Dallas

To Summarize

- ❖ LET is beneficial for autonomous embedded systems
 - ❖ Periodic Input
 - ❖ Predictable and verifiable output
- ❖ Has a flawed assumption: periodic task release
 - ❖ The flaw is due to ignorance to data (when memory is limited)
 - Temporal data availability
 - Spatial data availability
- ❖ Existing work
 - ❖ Has considered aperiodic input interval on LET (sporadic input events)
 - ❖ None to the best of our knowledge has considered aperiodic task release due to data and memory size limitation

Heijunka (Production Leveling)

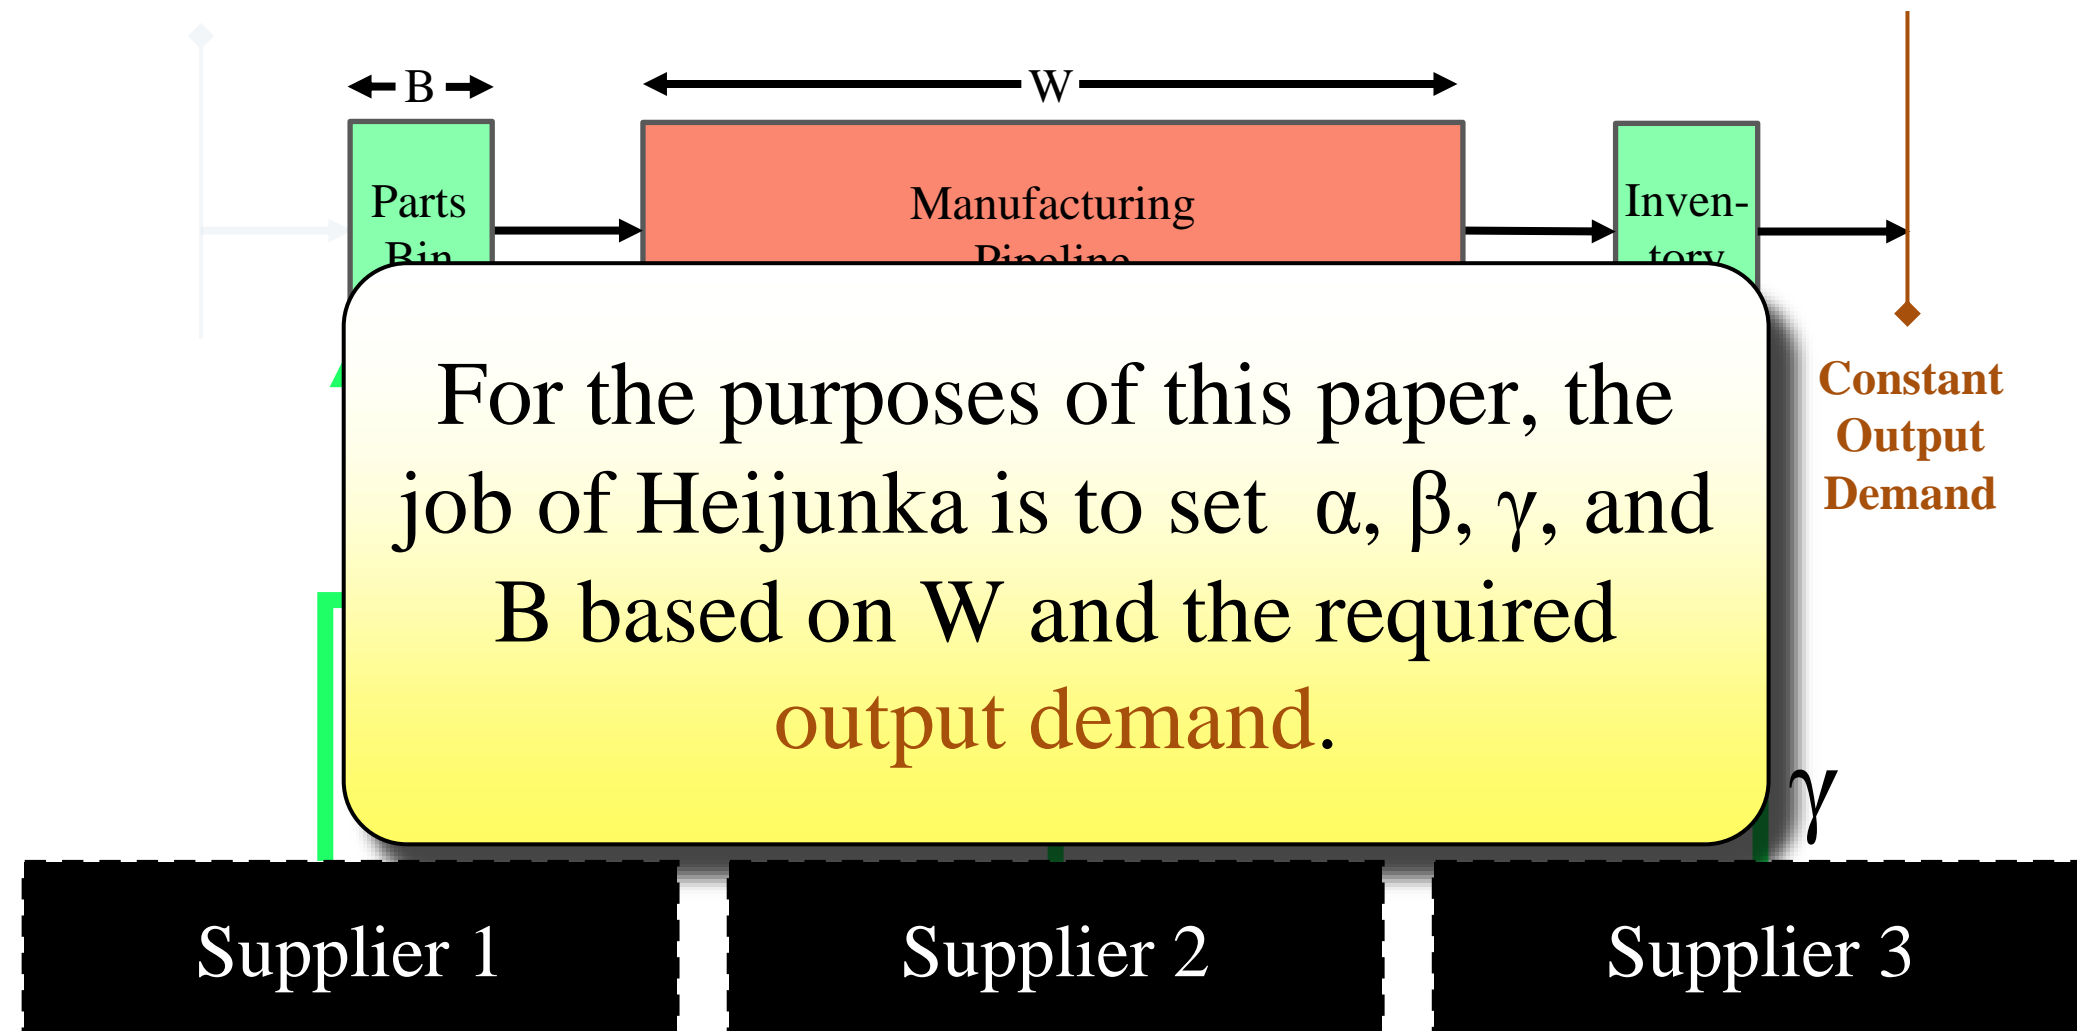


❖ $\underline{\alpha}$, $\underline{\beta}$, and $\underline{\gamma}$ are supply rates for each supplier

❖ \underline{B} is the size of the parts bin to store parts

❖ \underline{W} is the worst-case time it takes to manufacture a car

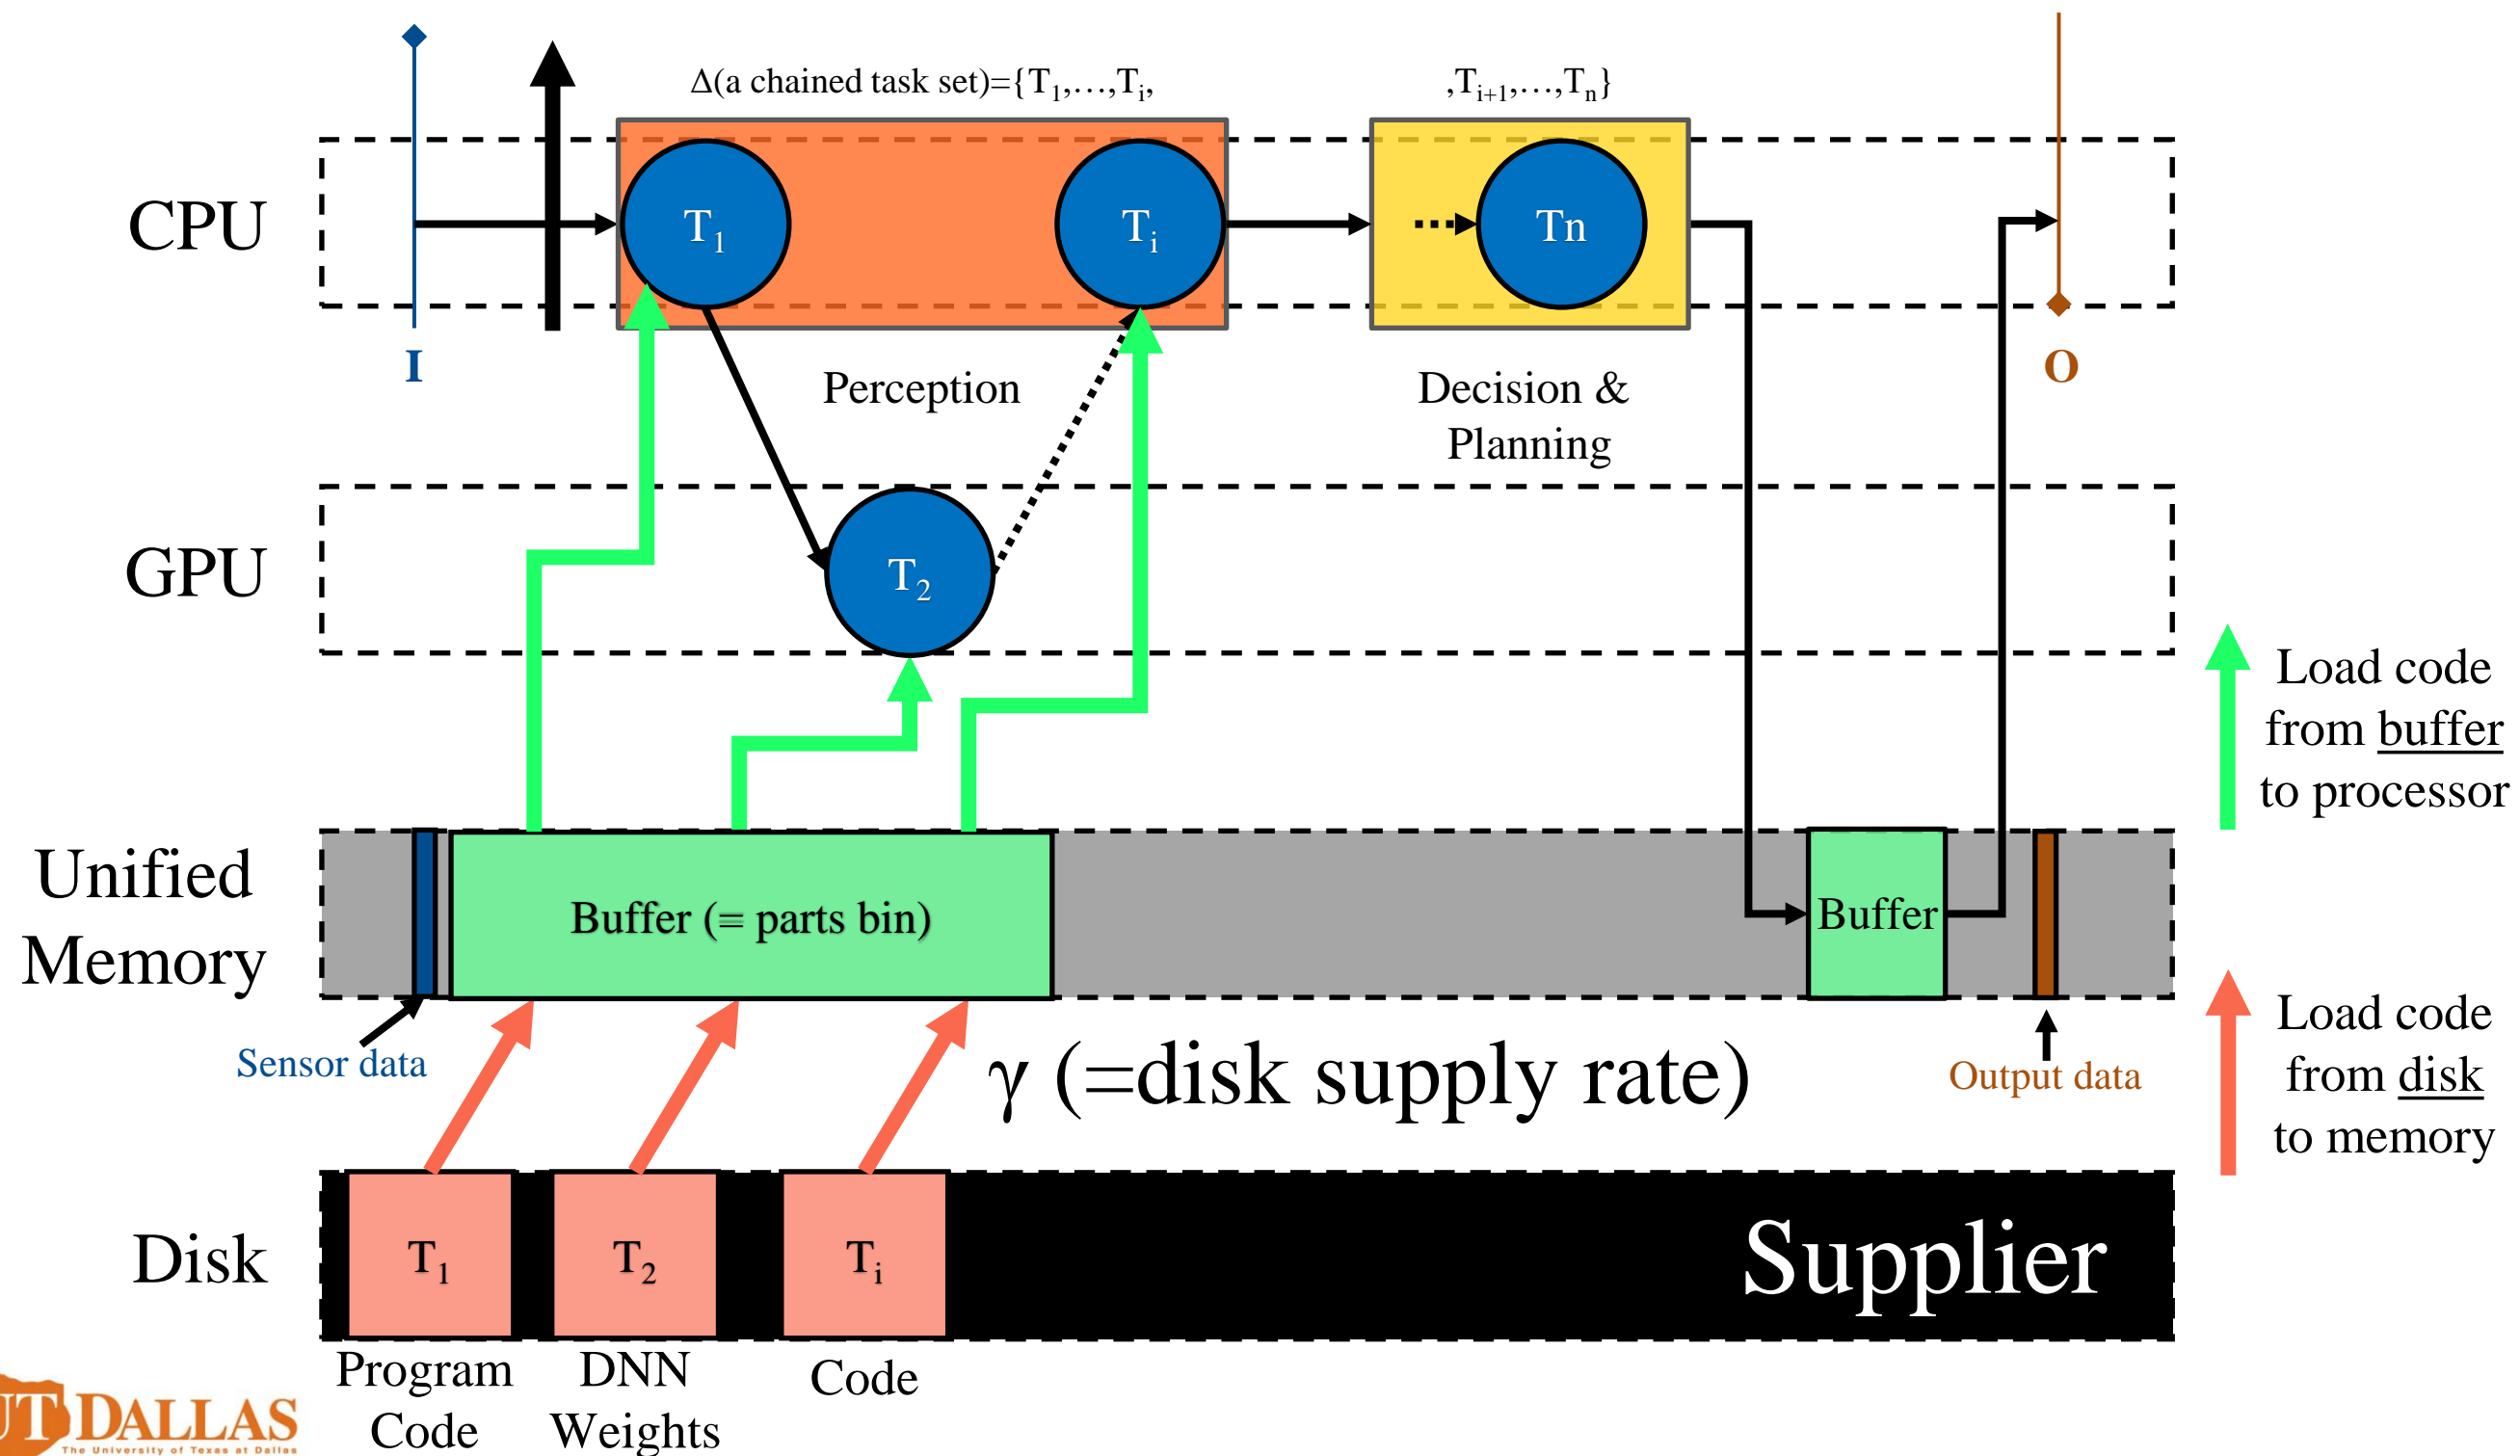
Heijunka (Production Leveling)



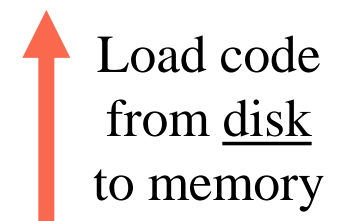
- ❖ α , β , and γ are supply rates for each supplier
- ❖ B is the size of the parts bin to store parts

- ❖ W is the worst-case time it takes to manufacture a car

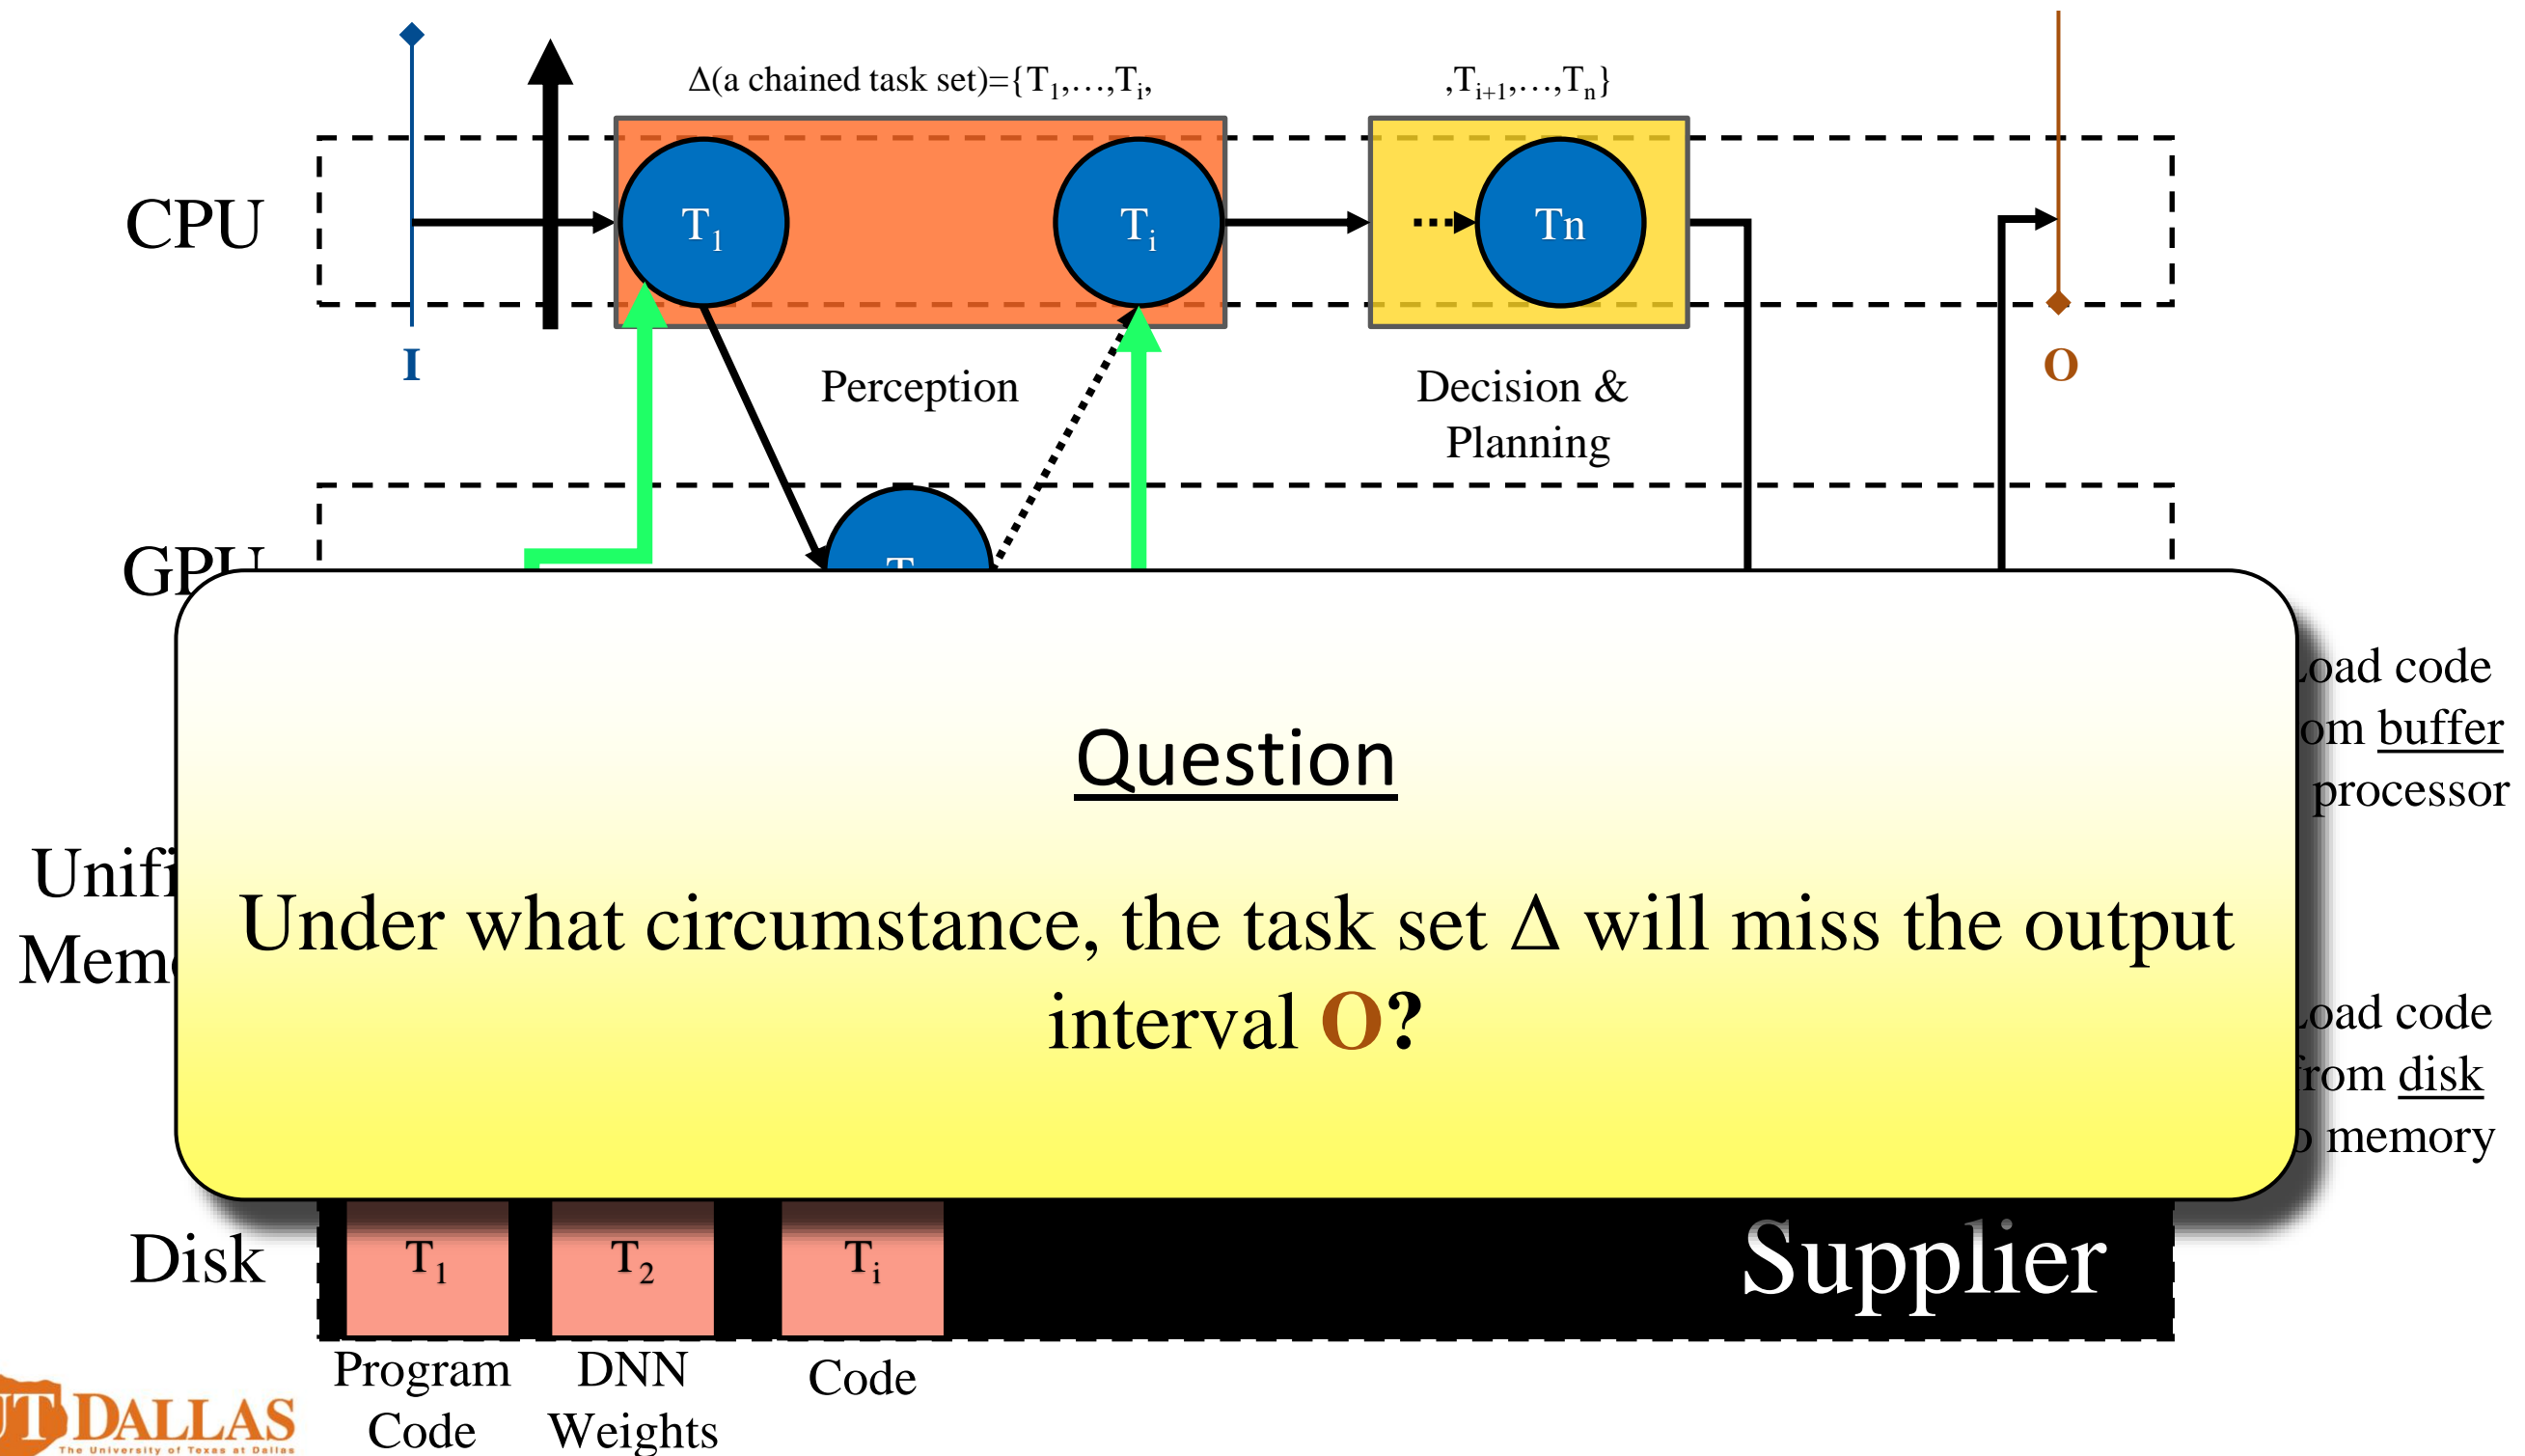
LET-based System Model with Heijunka Expansion




Task set is still sporadic.



Ensuring Temporal Data Availability



Ensuring Temporal Data Availability

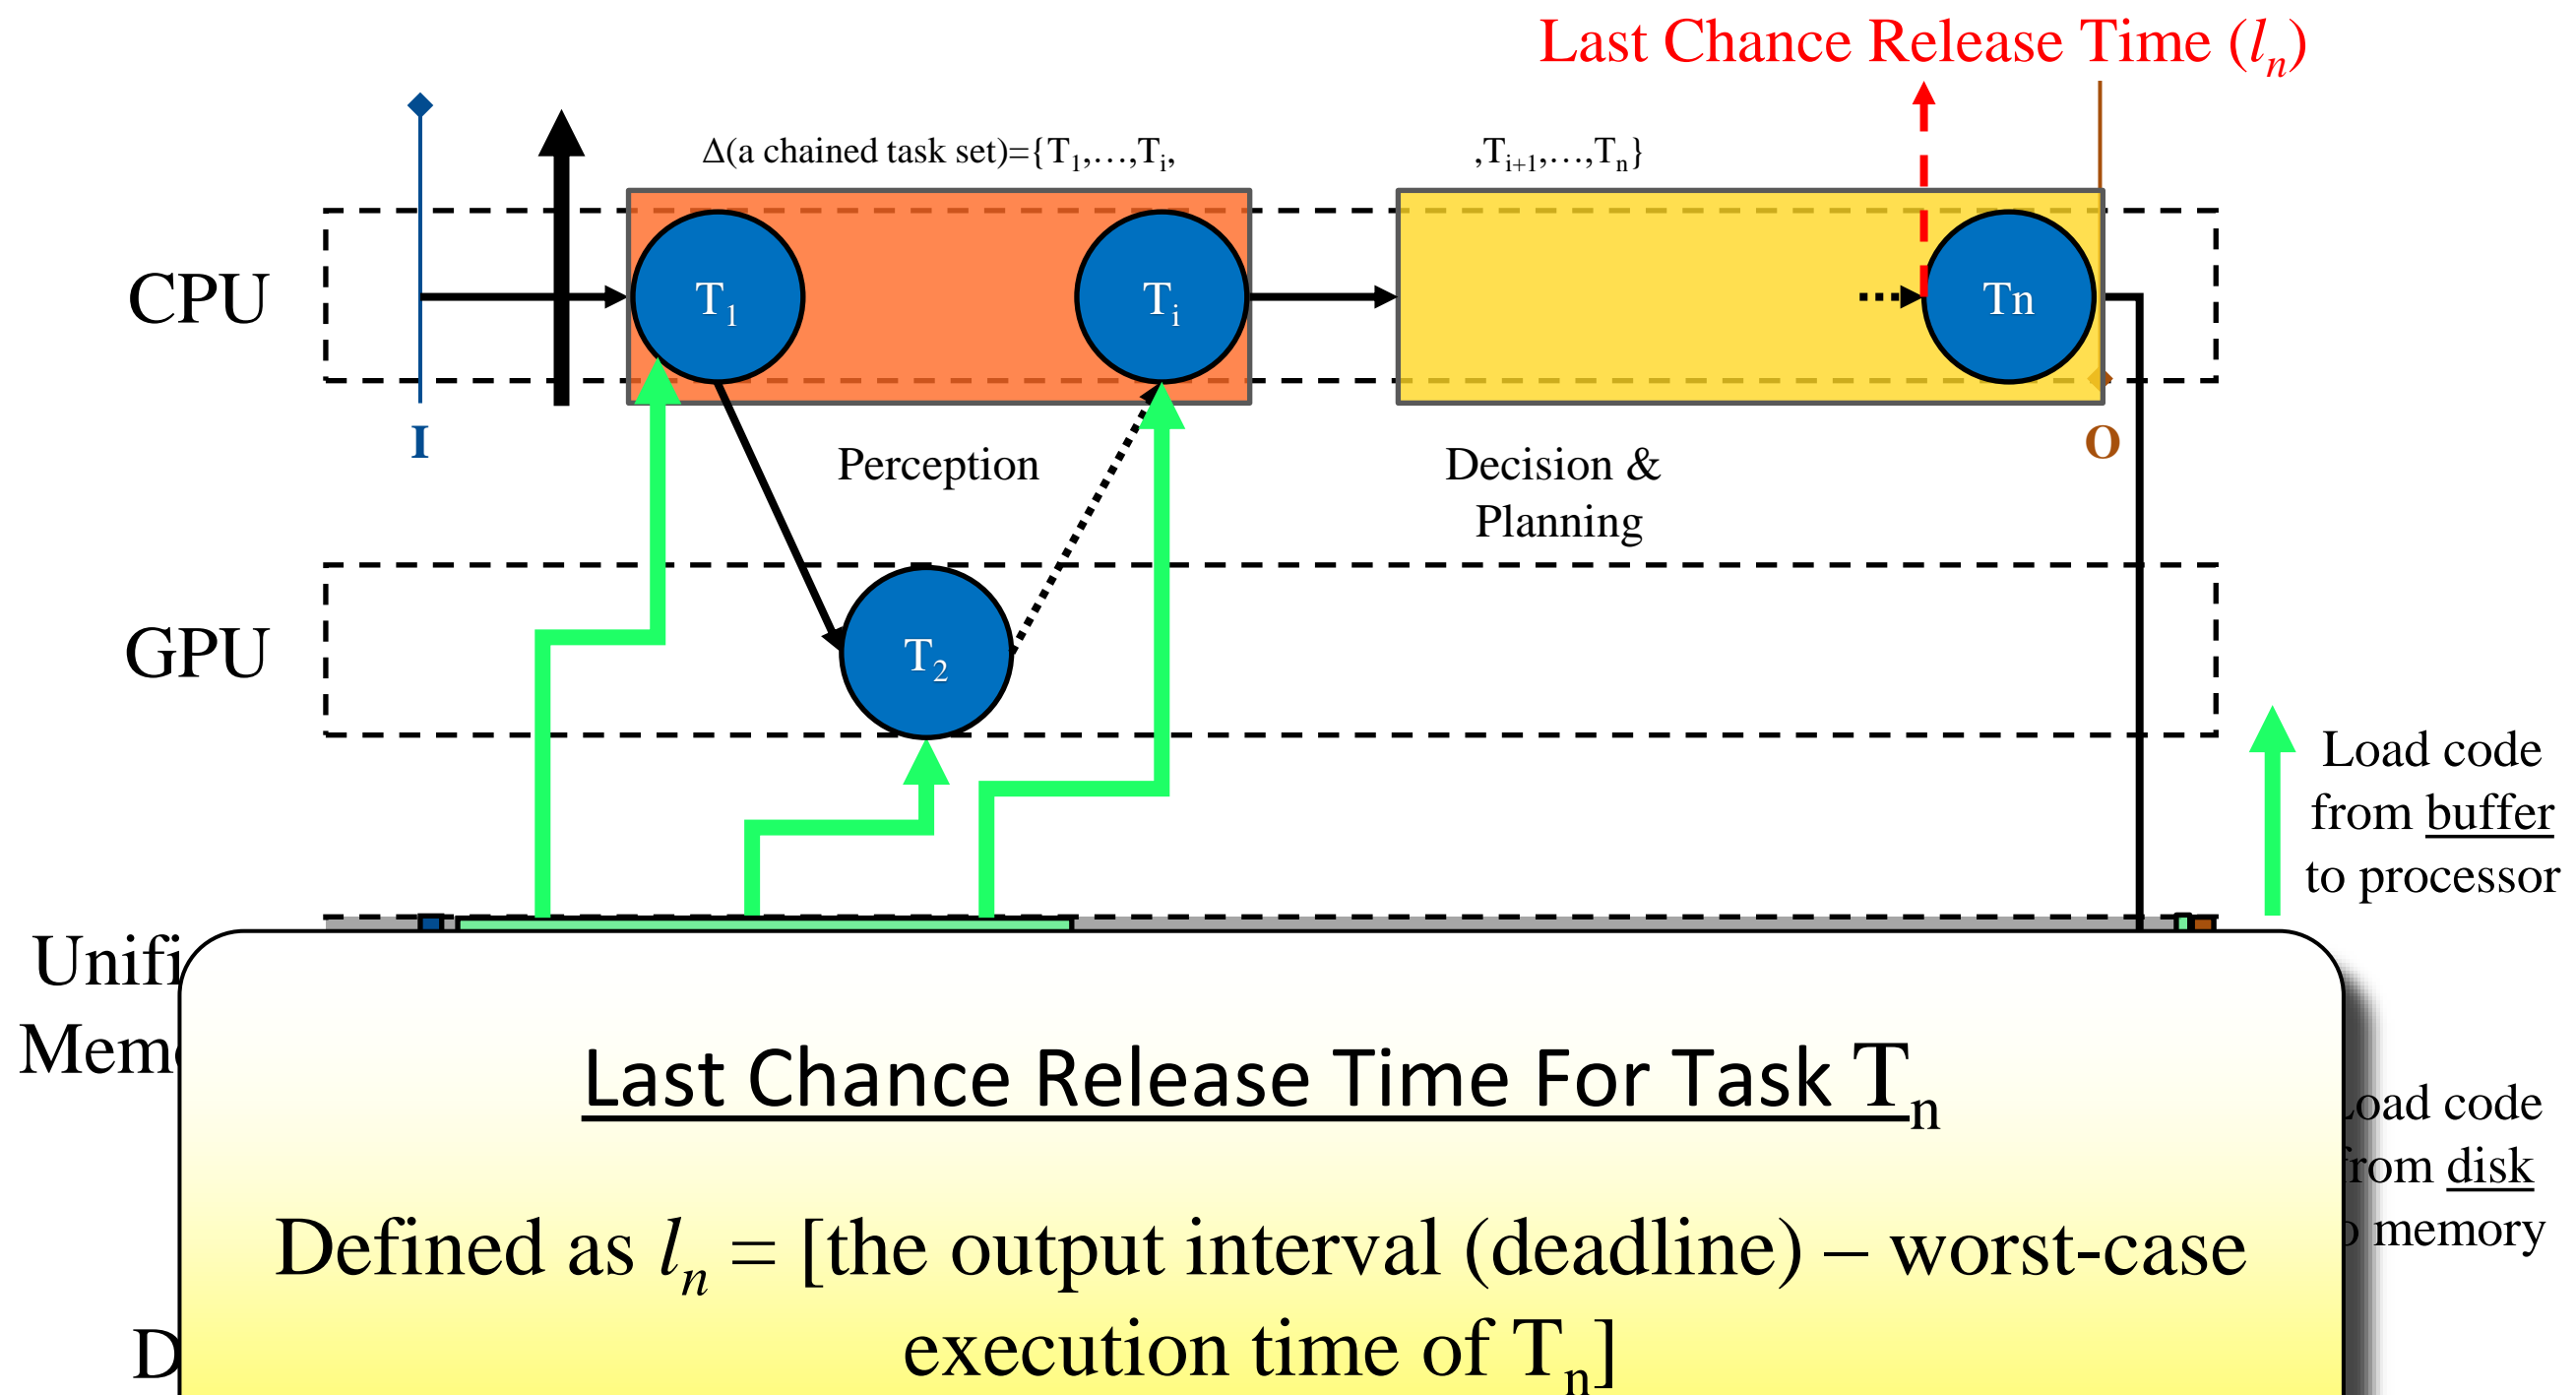


Unifi
Mem

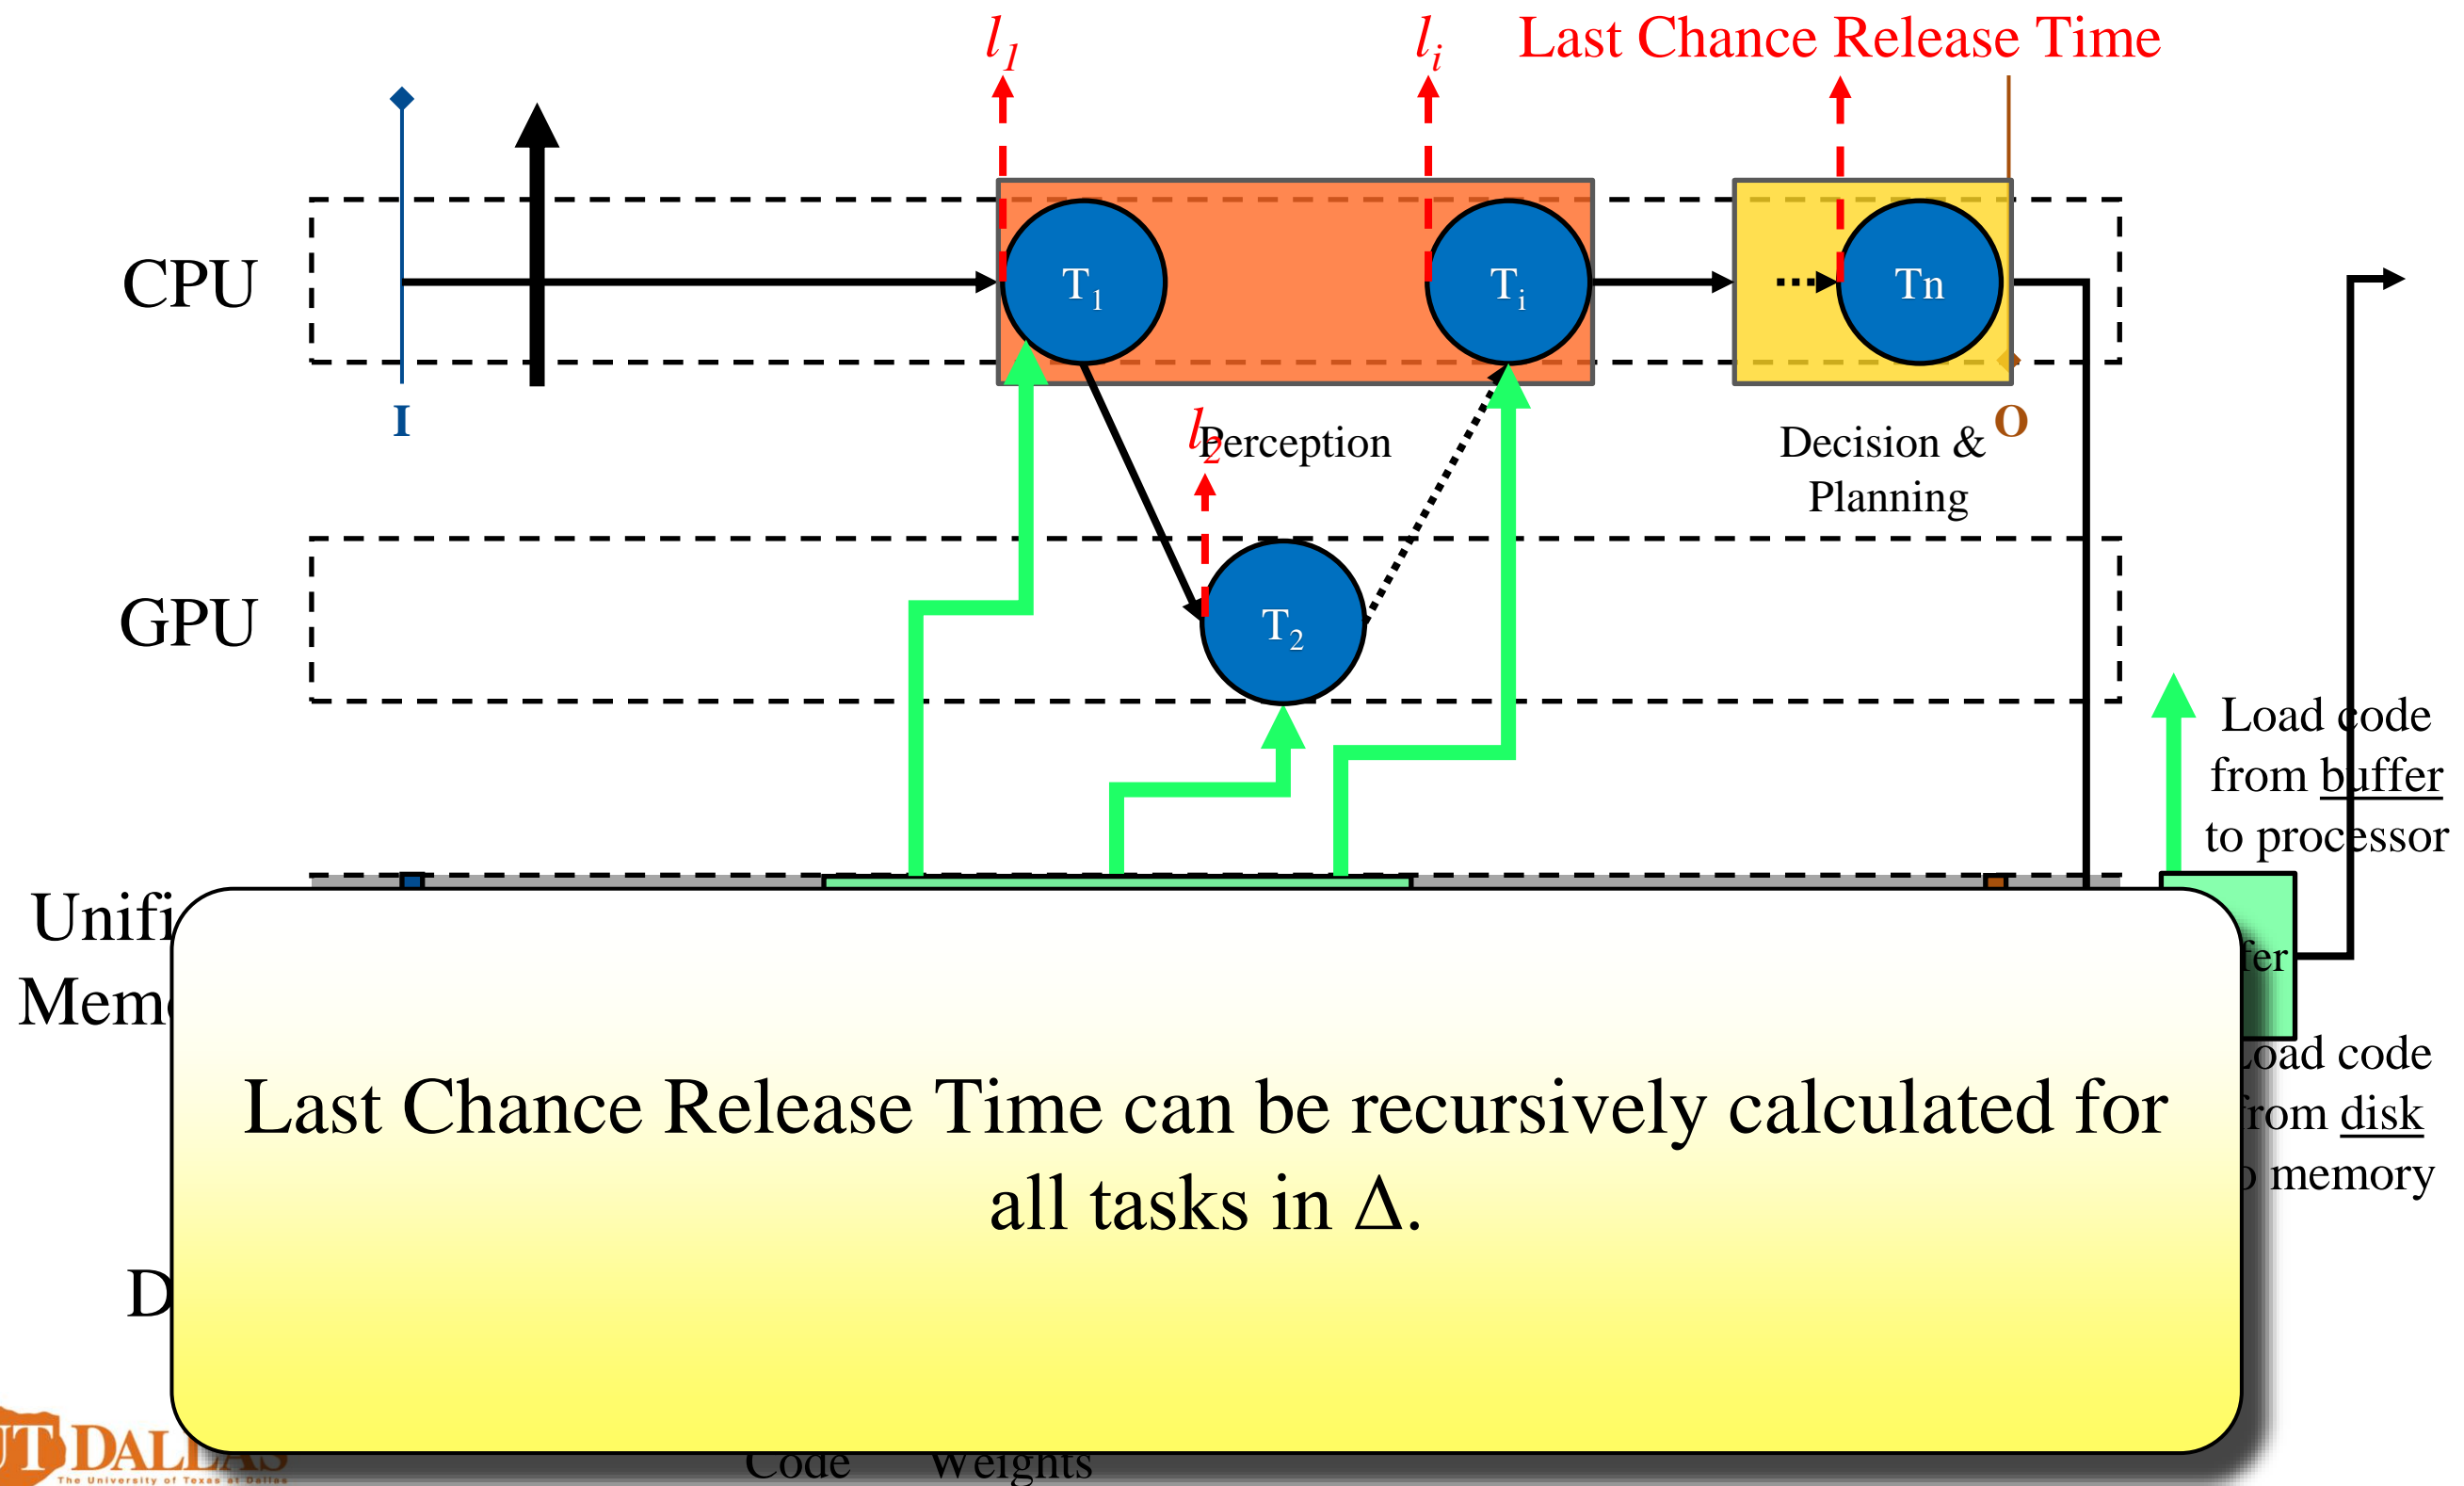
Question

When is the latest time T_n can be released to not miss the deadline? In other words, what is the latest time instance when the data for T_n should be available?

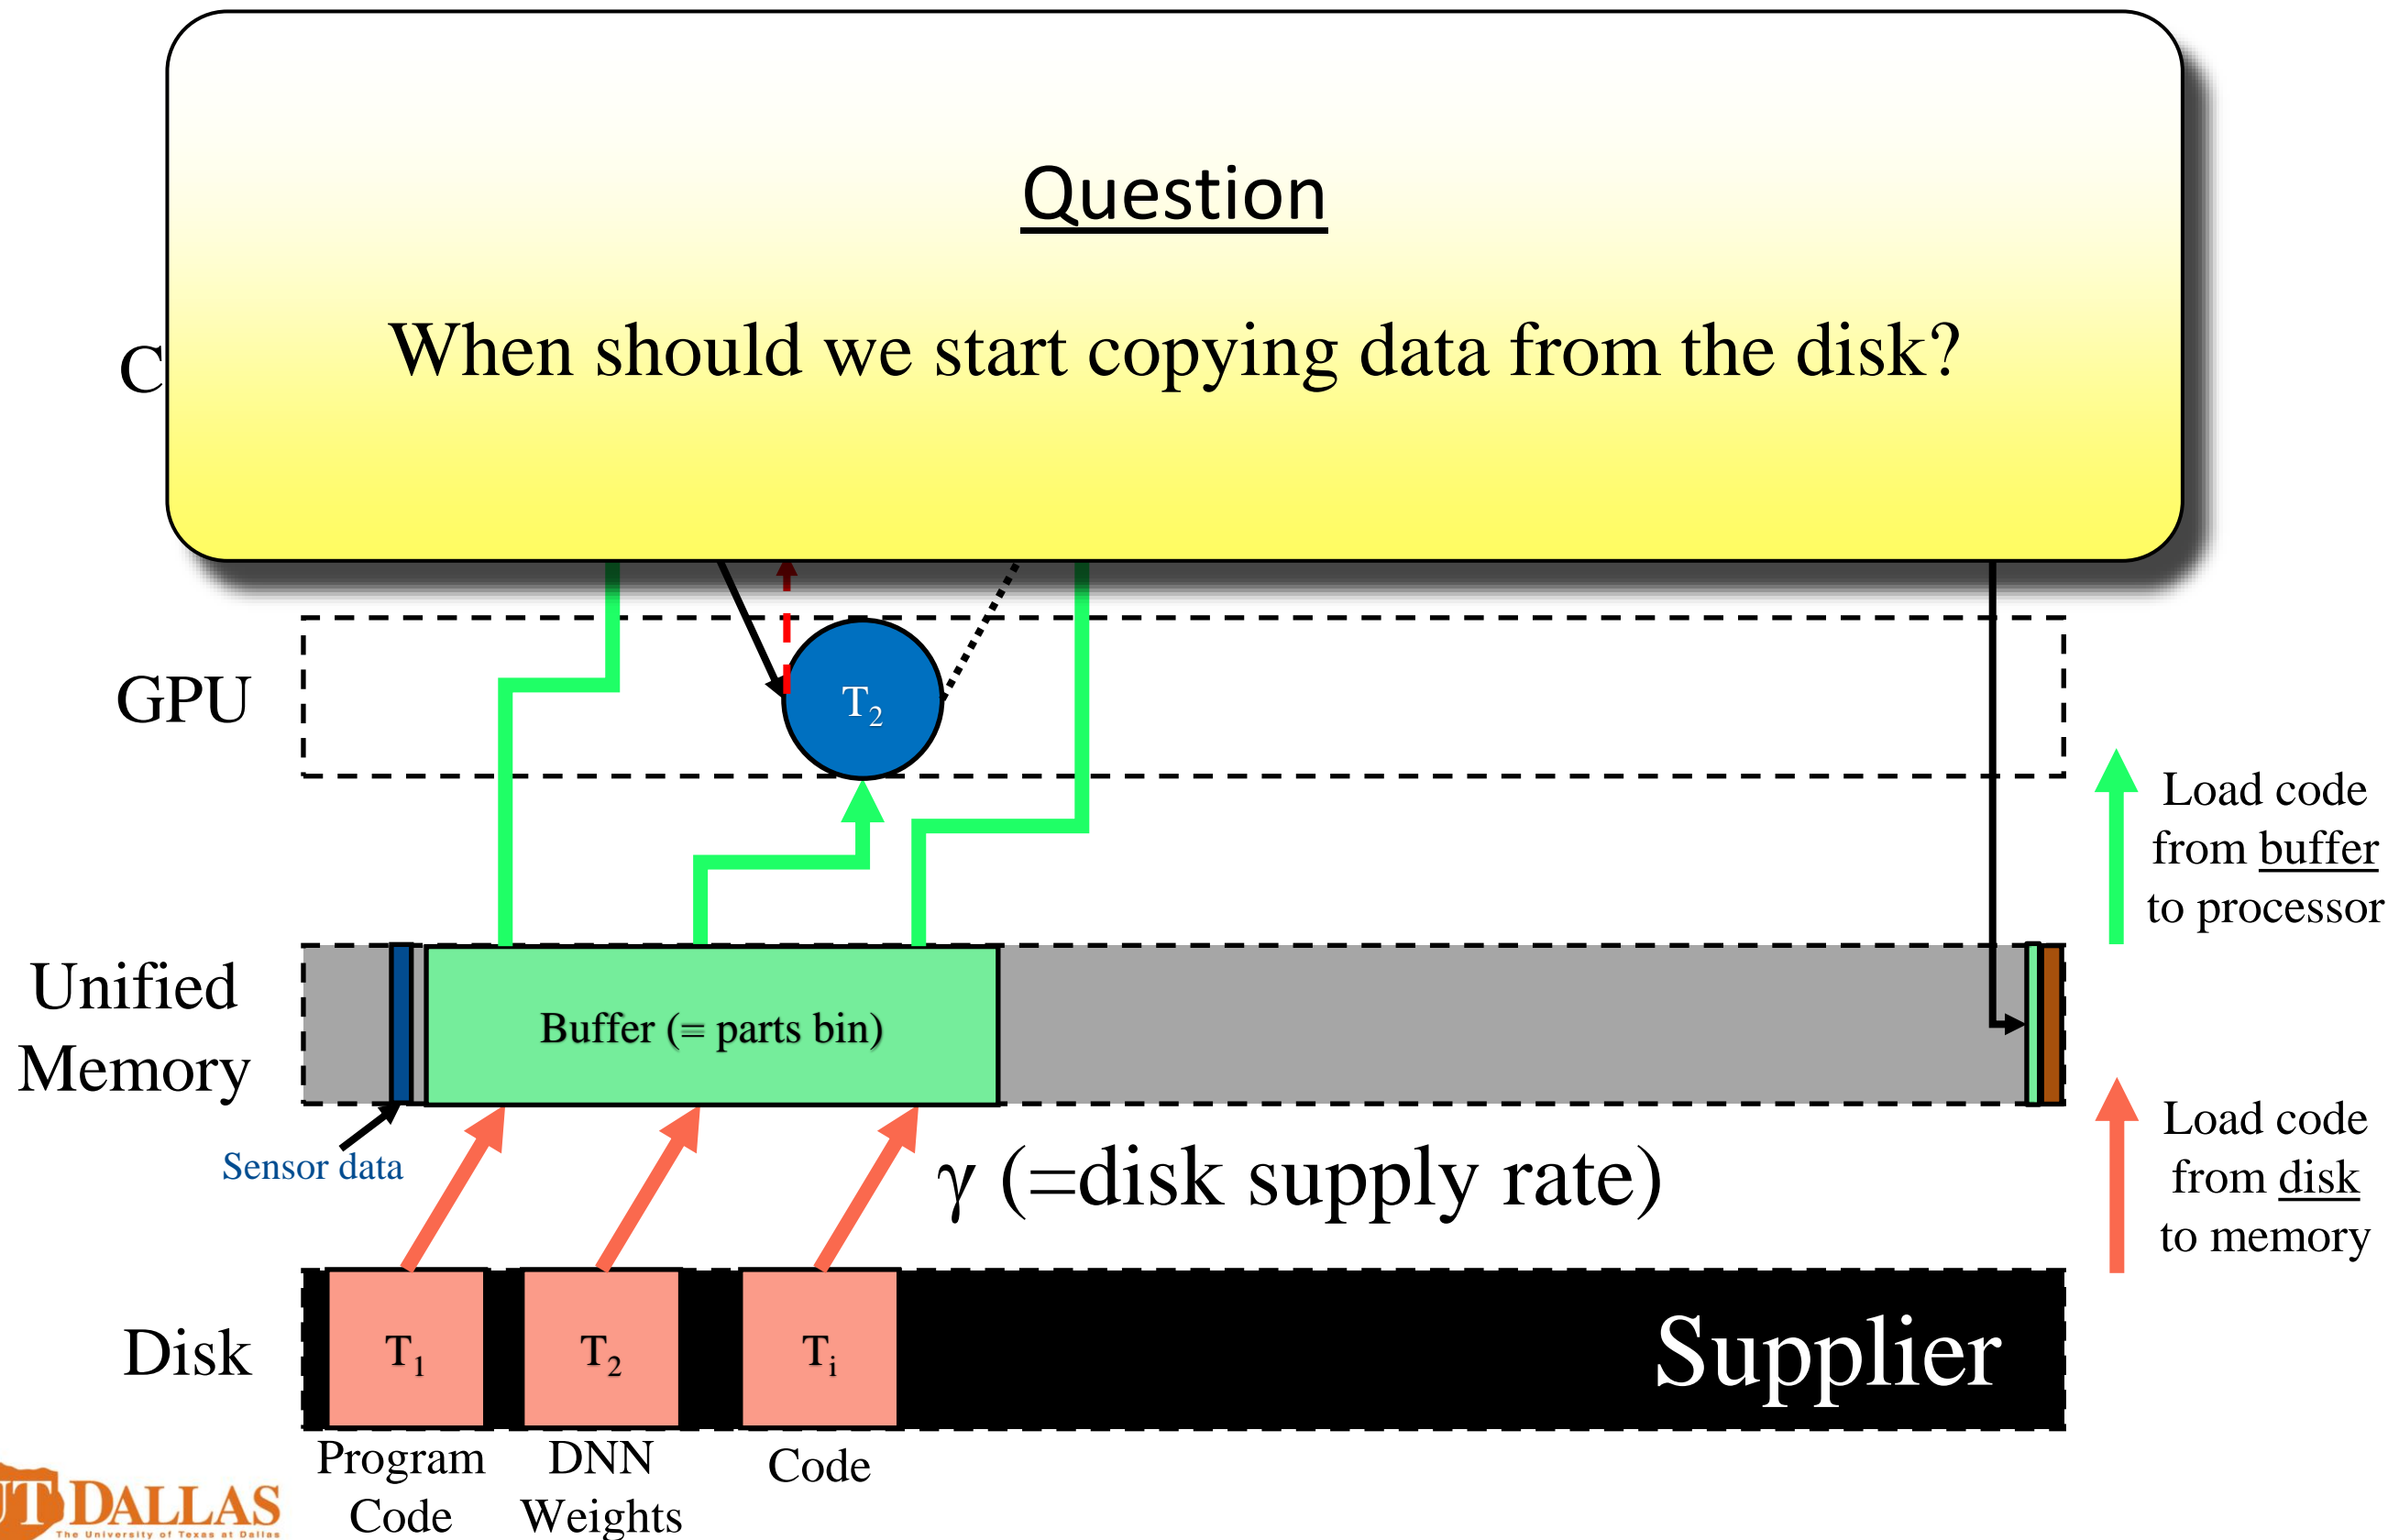
Ensuring Temporal Data Availability



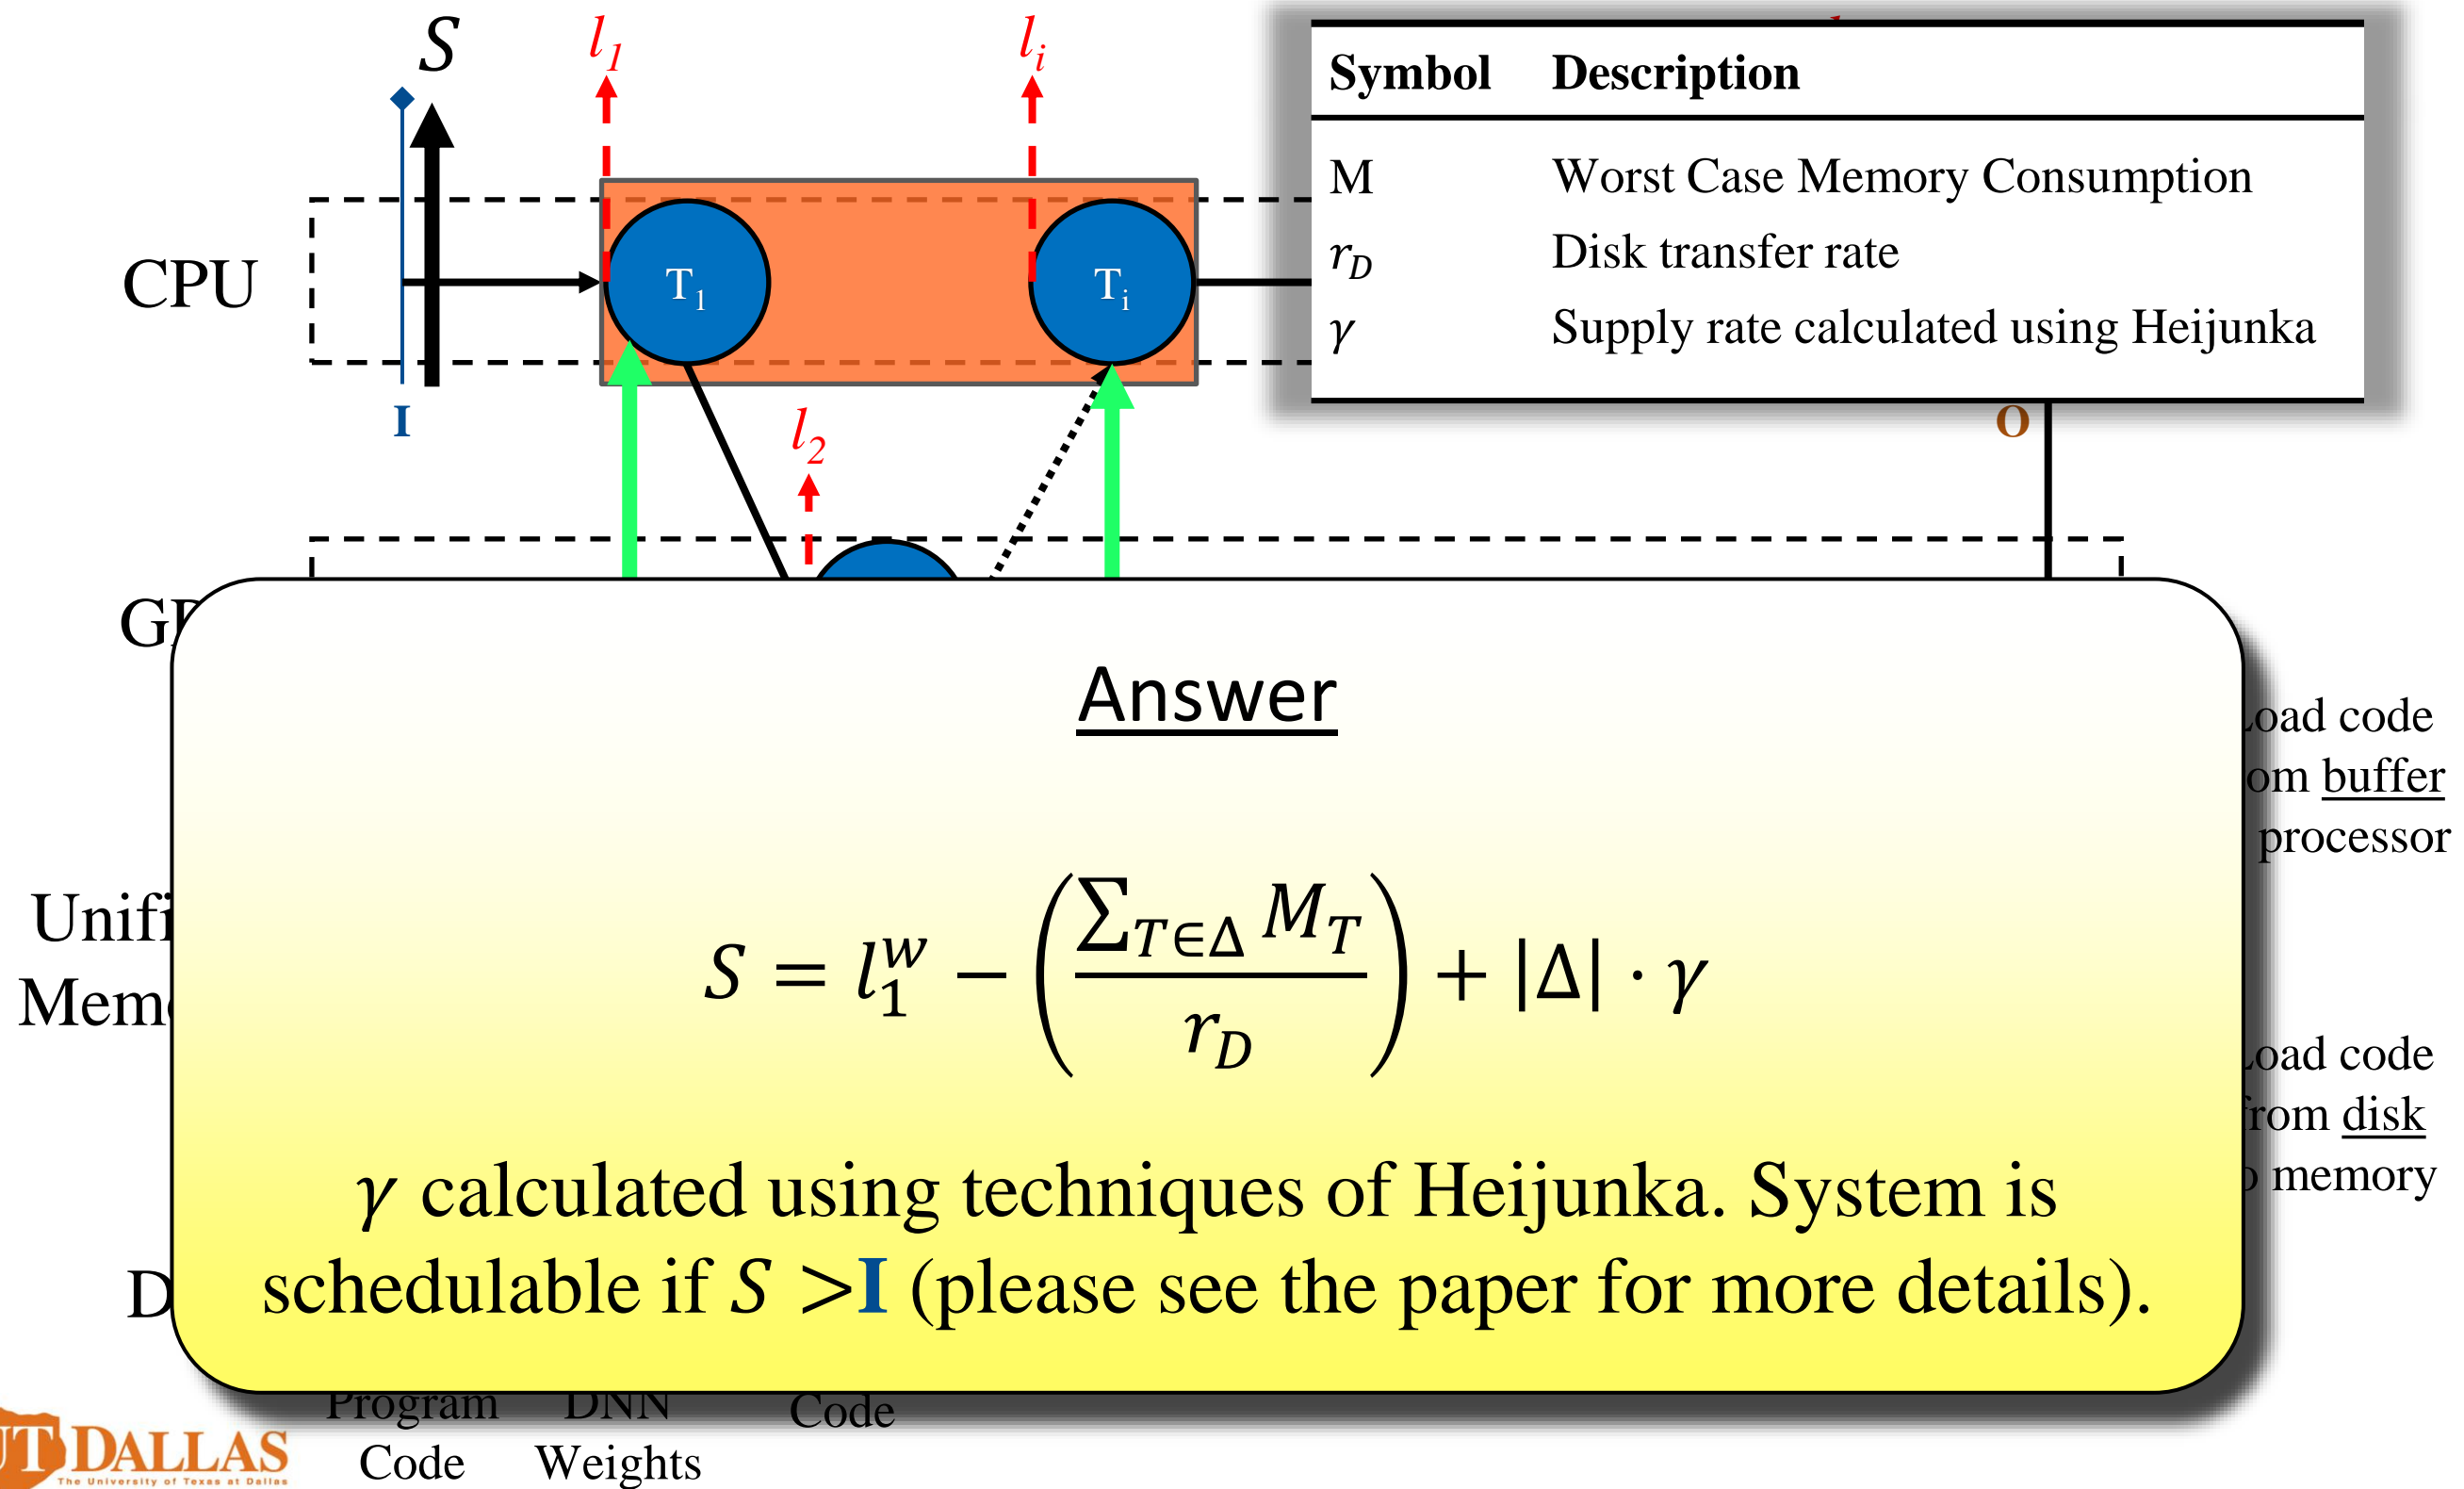
Ensuring Temporal Data Availability



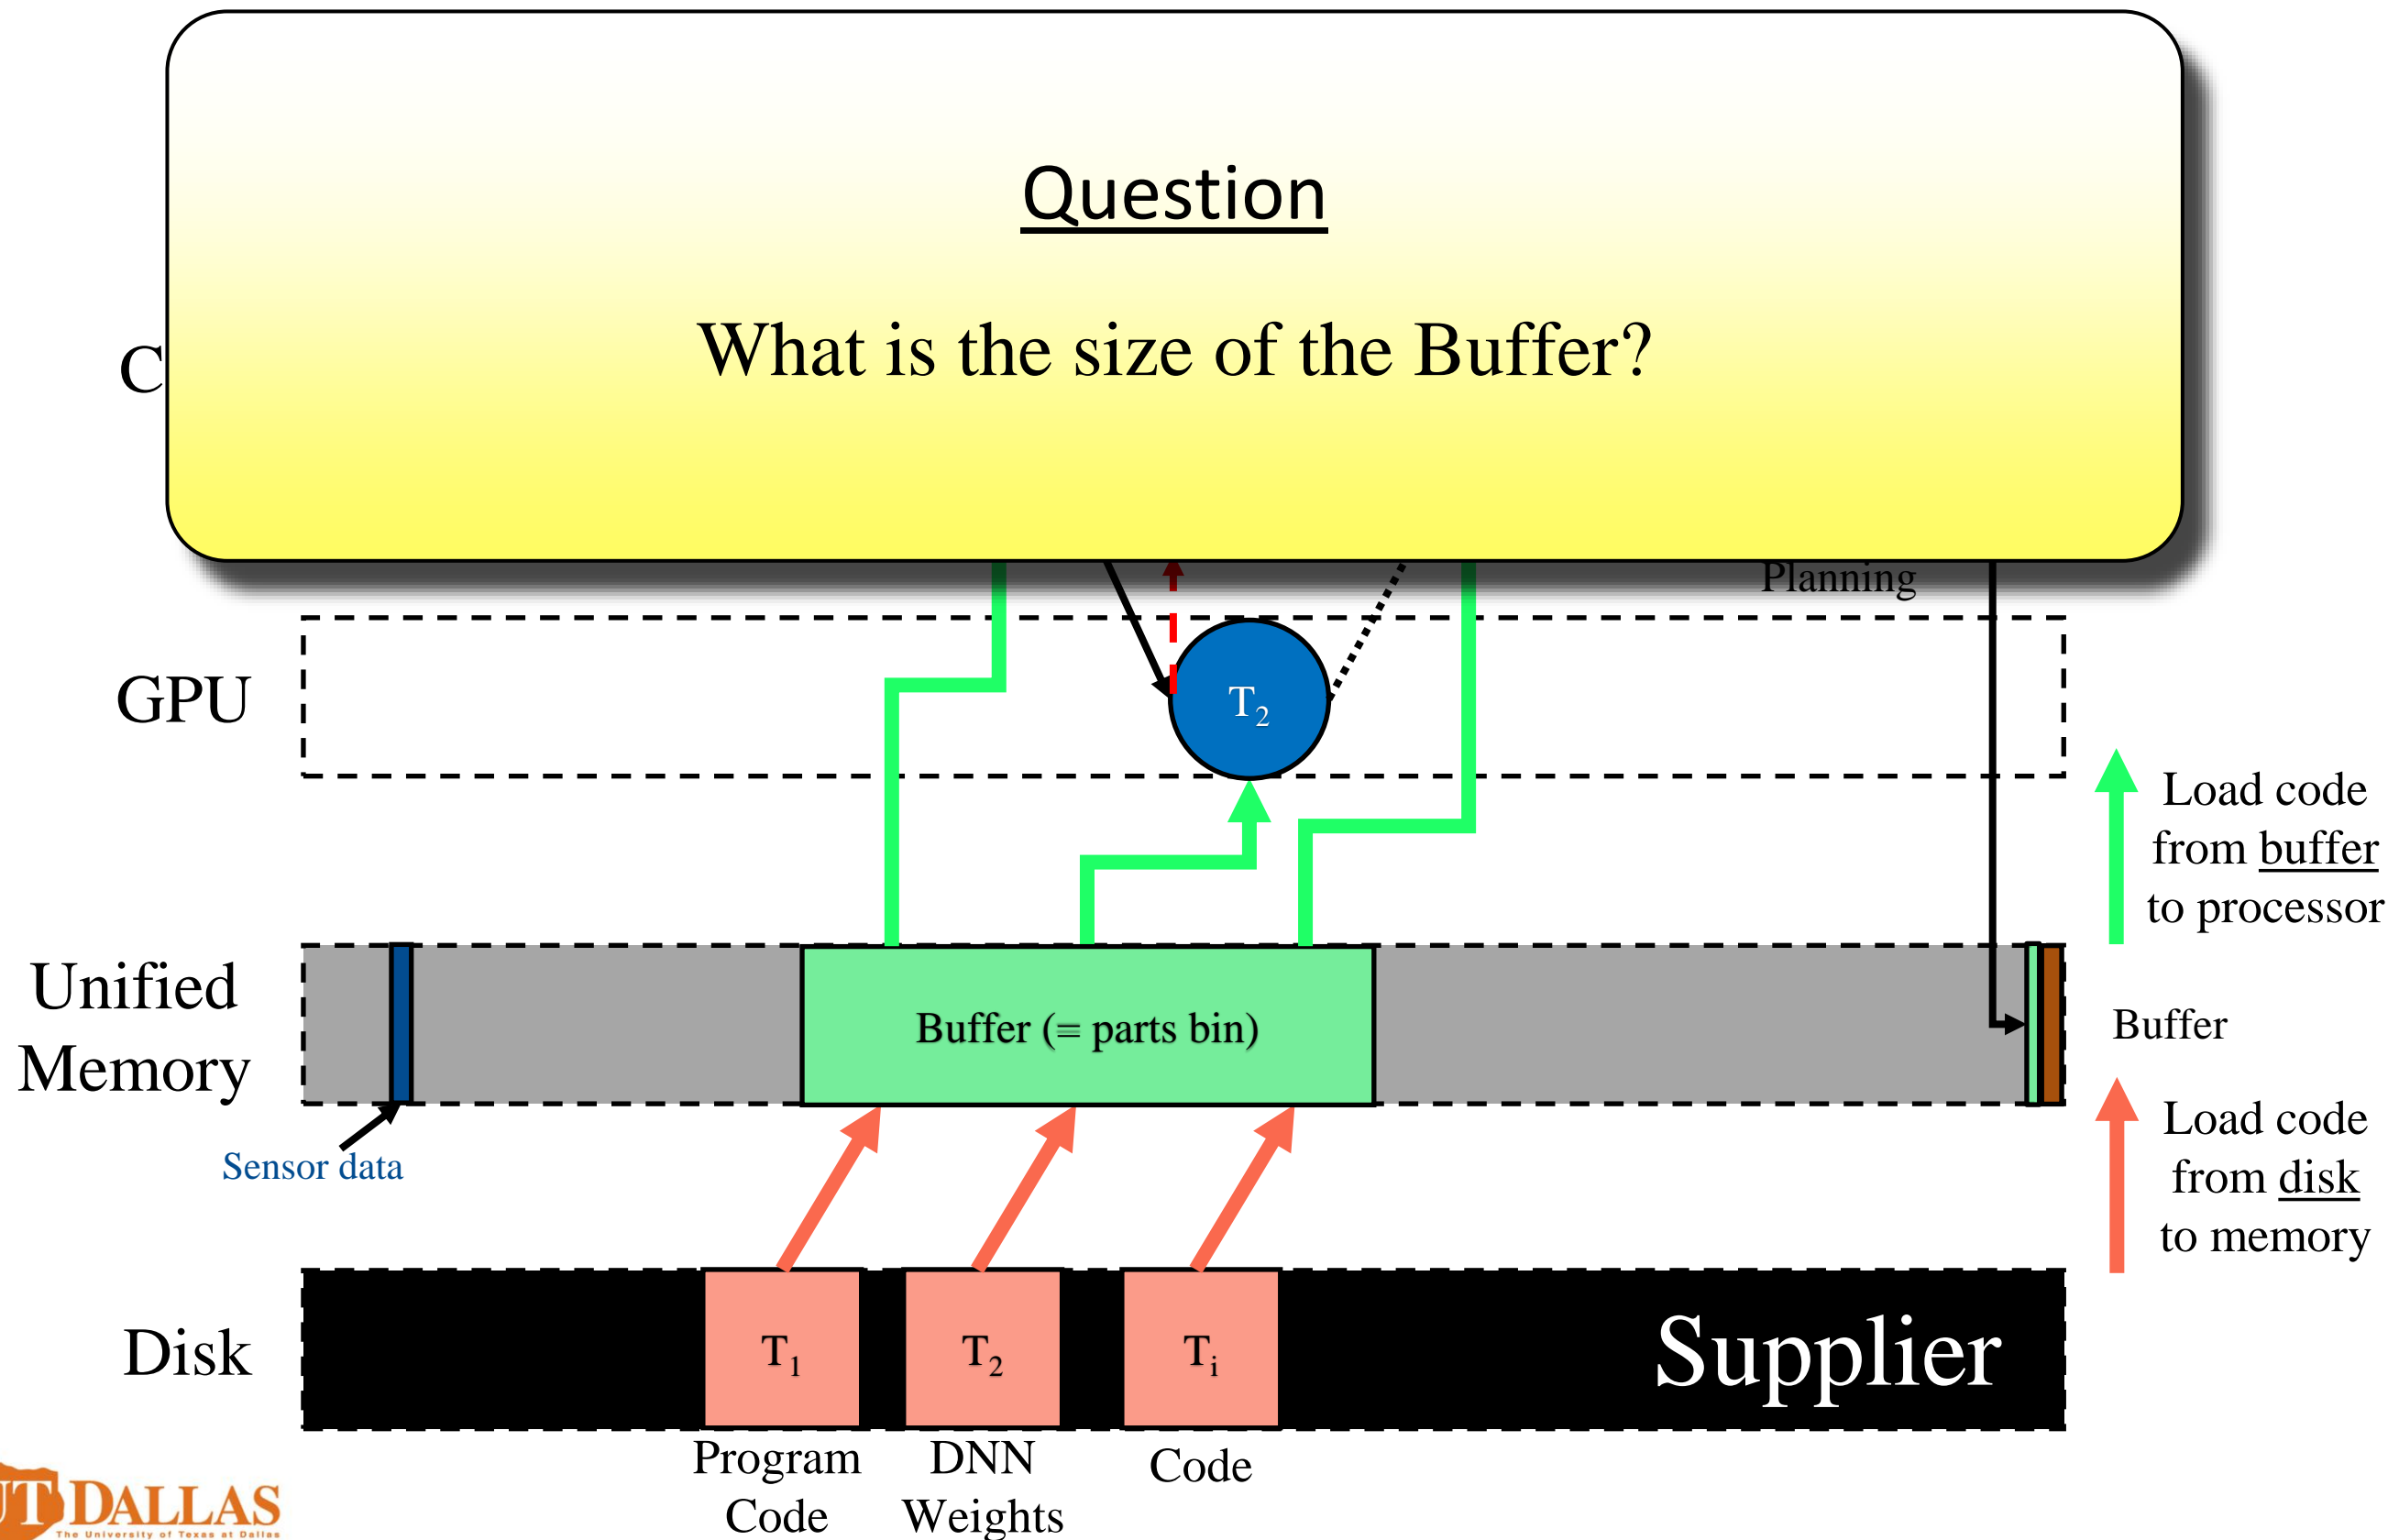
Ensuring Temporal Data Availability



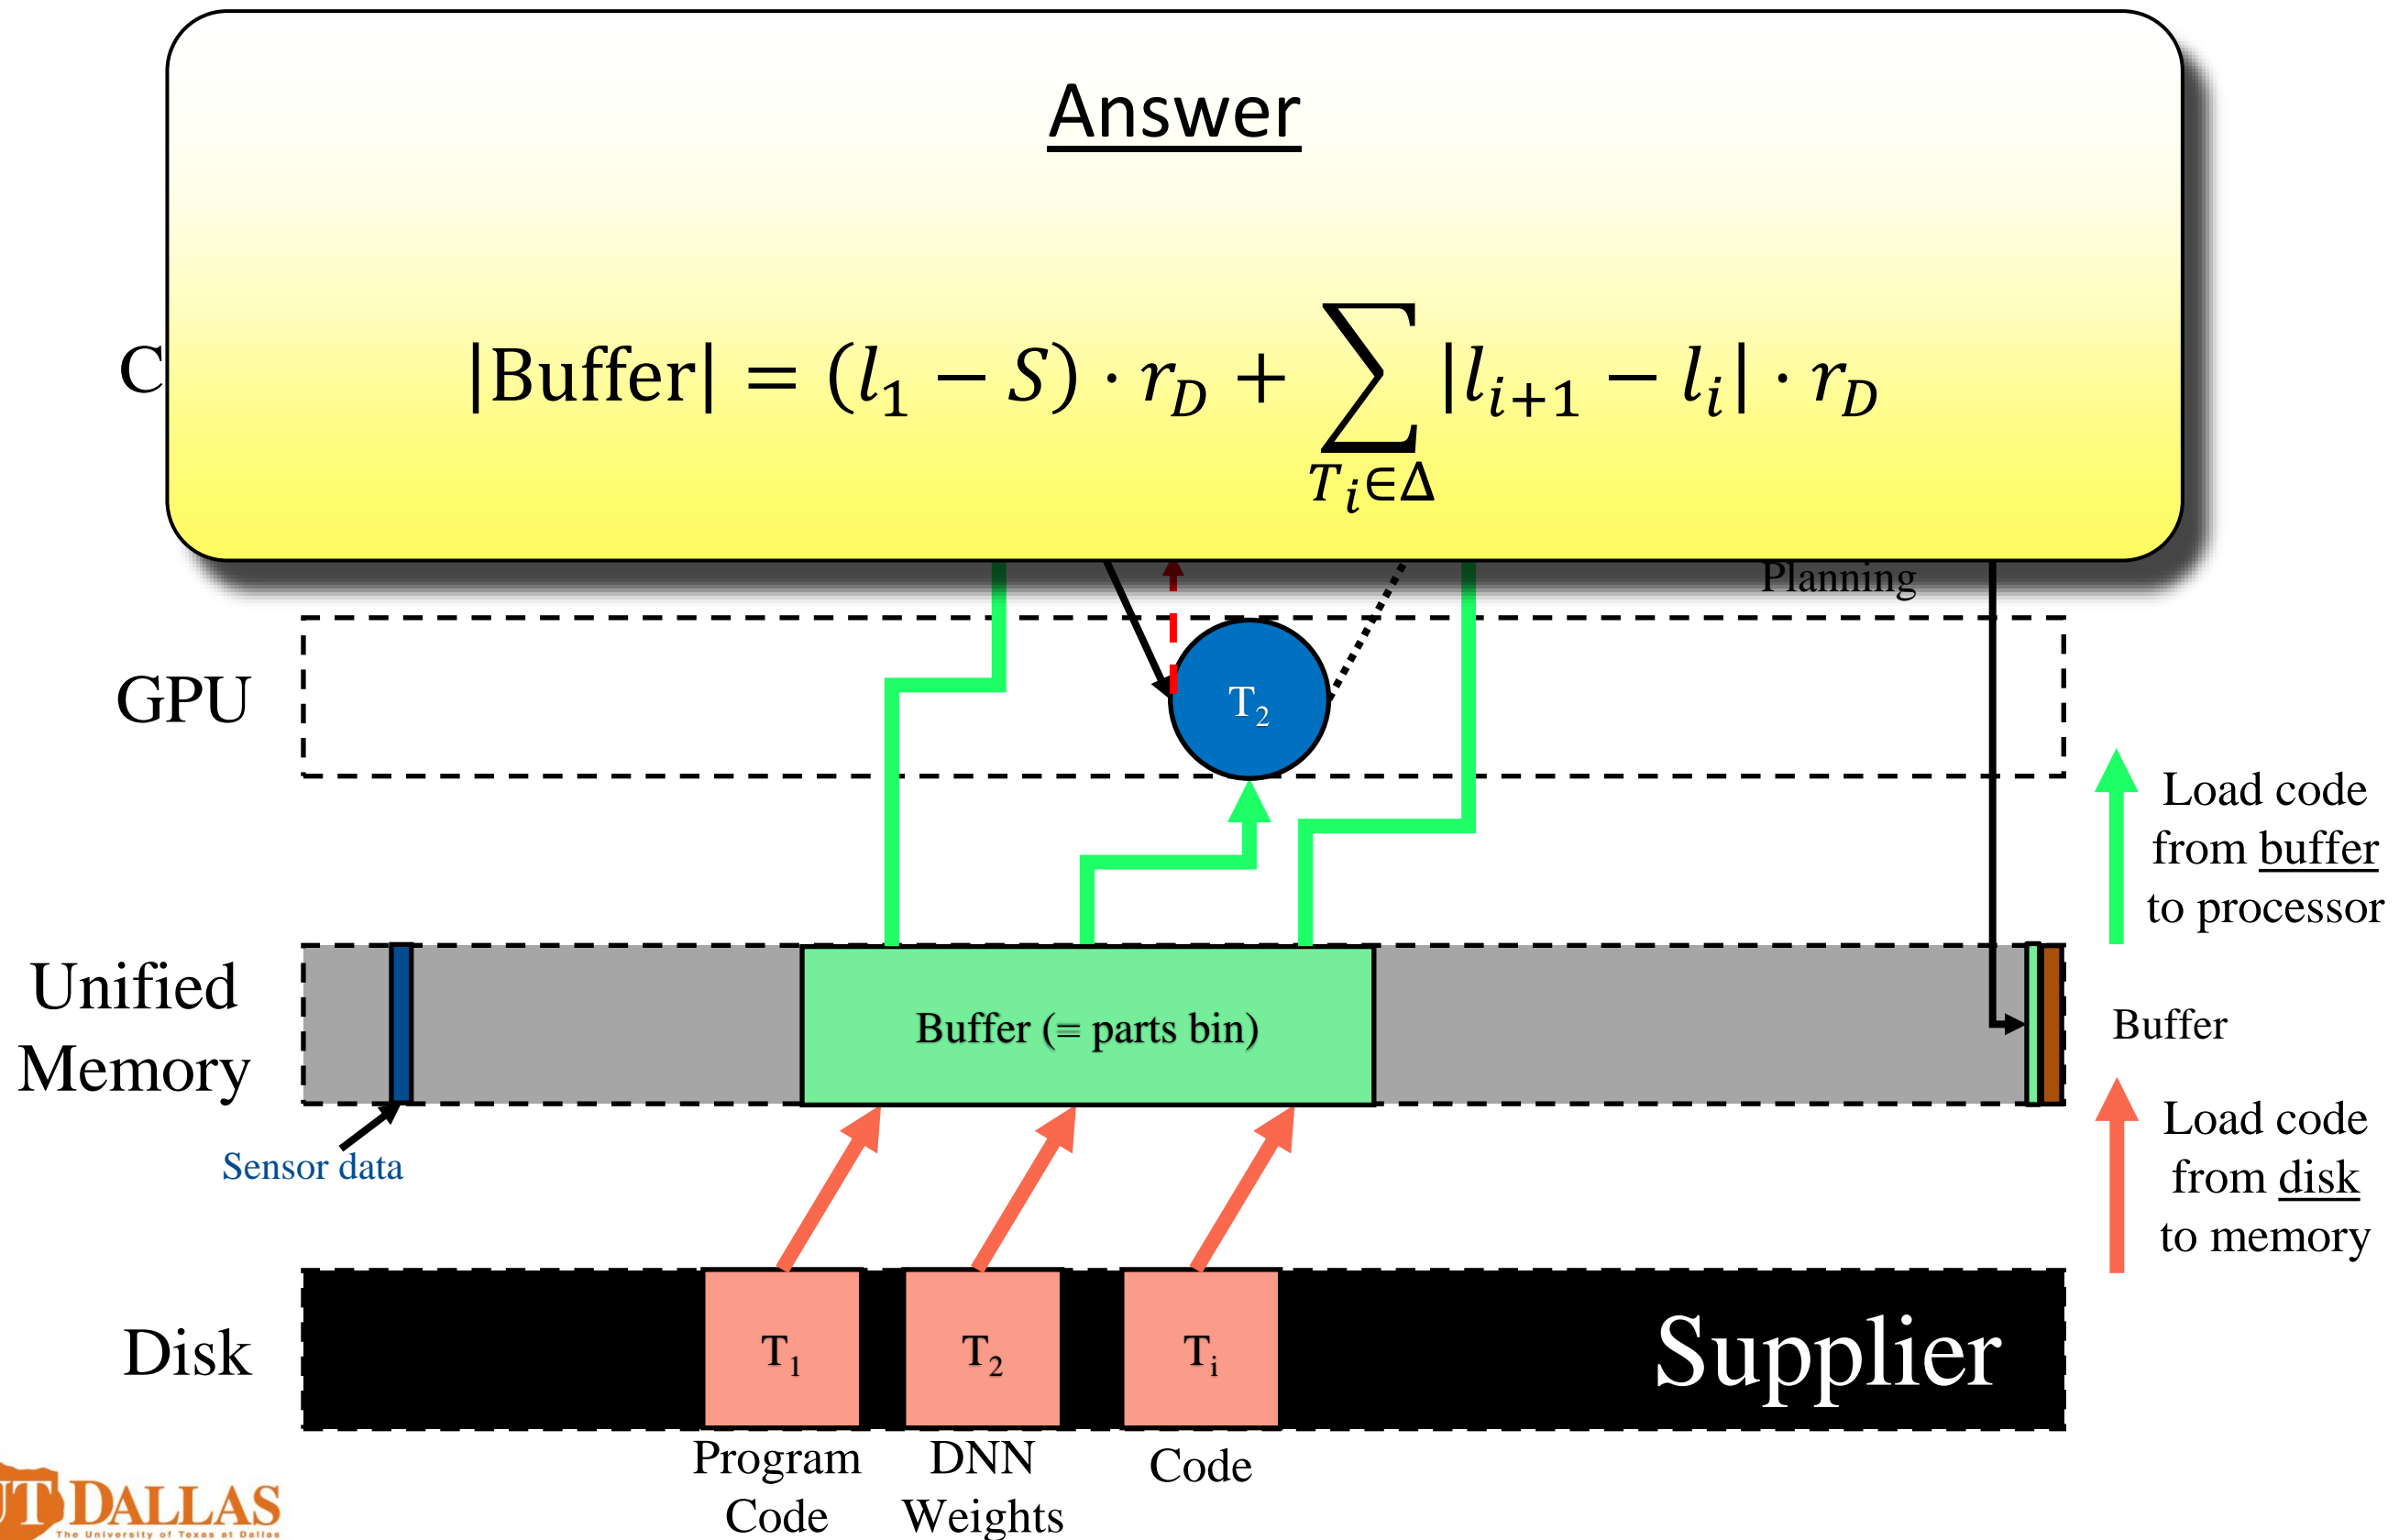
Ensuring Temporal Data Availability



Ensuring Spatial Data Availability



Ensuring Spatial Data Availability



Ensuring Spatial Data Availability

Answer

$$|\text{Buffer}| = (l_1 - S) \cdot r_D + \sum_{T_i \in \Delta} |l_{i+1} - l_i| \cdot r_D$$

Planning

GPU

T_2

The above calculation is very pessimistic, resulting in a large buffer size. By using a technique of Heijunka, we can shrink the reservation window (i.e., the buffer size) to $\gamma \cdot r_D$.

(Please see the paper for details and proof)

Disk

T_1

T_2

T_i

Supplier

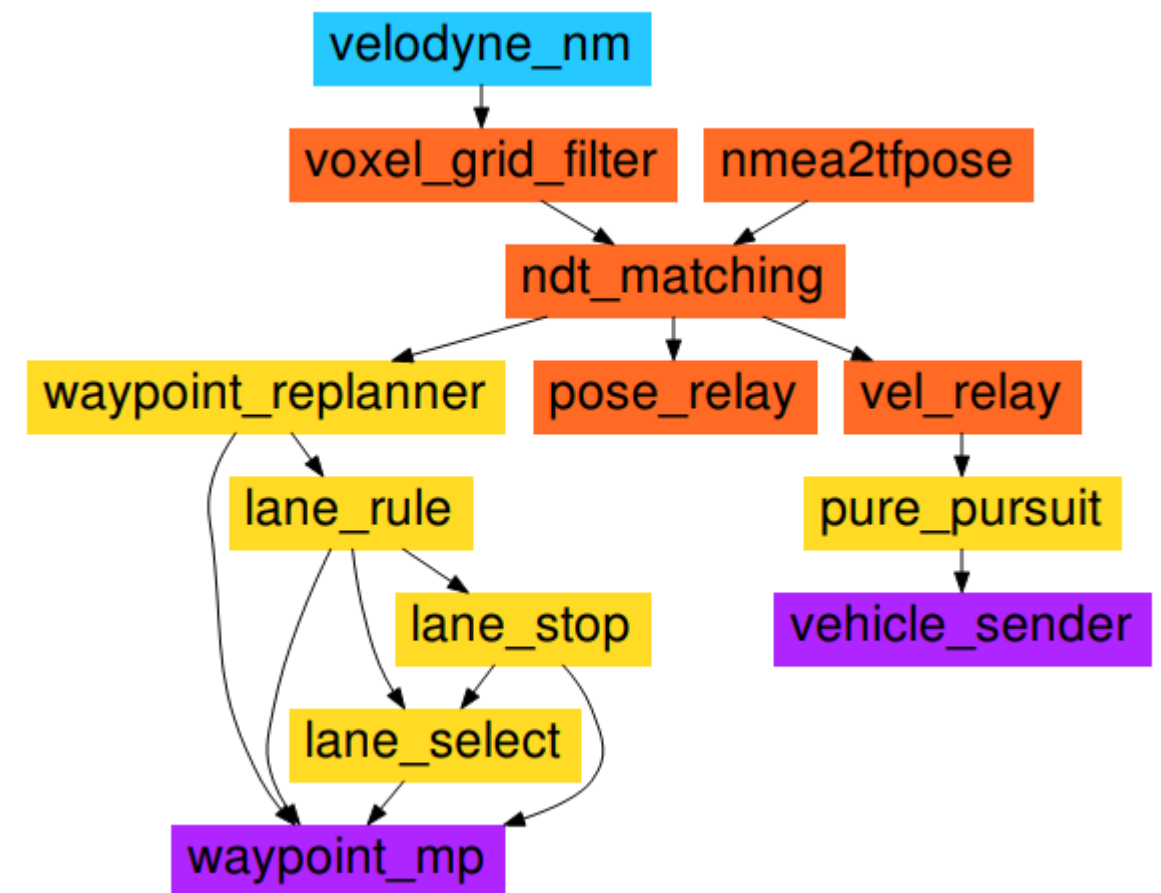
Program
Code

DNN
Weights

Code

Evaluation

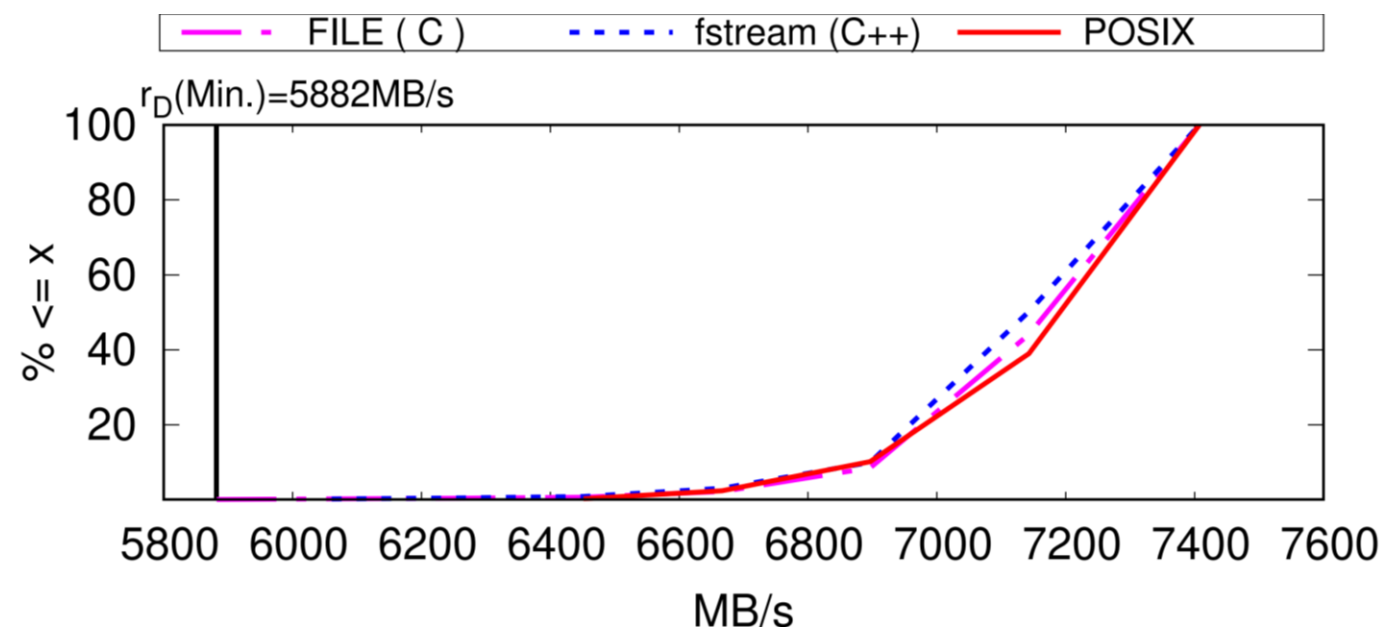
- ❖ We implement ResCue, a runtime solution with:
 - ❖ A dynamic data scheduler
 - ❖ A dynamic memory reservation
- ❖ Evaluation Platform
 - ❖ Autoware
 - Open-source autonomous driving full-stack software
 - ❖ NVIDIA Jetson AGX Xavier
 - 8-core CPU, 512-core Volta-based GPU, 16GB of RAM



Autoware

Results

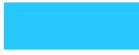











- ❖ First, we calculate r_D for the on-board SSD of AGX Xavier



- ❖ $r_D = 5882$, the low 99.9th percentile

Calculated Parameters for Autoware

TABLE I: Components of Autoware sorted in order of execution, and their WCET, WCMC, and the value of l for each component under various input and output periods (LCR value is for the first period). All times are in milliseconds and memory size is in megabytes.

Color	Category	Process	WCET(ms)	WCMC(MB)	LCR(I=42,O=300)	LCR(I=33,O=200)	LCR(I=171,O=171)
	LIDAR (Sensing)	velodyne_nodelet_manager	2.64	56.01	206.100	106.100	77.100
	Downsampling (Perception)	voxel_grid_filter	3.96	52.75	208.740	108.740	79.740
	Localizer (Perception)	ndt_matching	45.37	2063.05 (static)	212.700	112.700	83.700
		nmea2tfpose	20.7	13.32	258.070	158.070	129.070
	Positioning (Perception)	pose_relay (Positioning)	12.78	11.98	278.770	178.770	149.770
		vel_relay (Velocity)		11.97			
	Planning	waypoint_replanner	2.06	9.84	291.550	191.550	162.550
		lane_rule	1.2	28.29	293.610	193.610	164.610
		lane_stop	0.86	11.76	294.810	194.810	165.810
		lane_select	0.84	12.18	295.670	195.670	166.670
		pure_pursuit	1.31	10.09	296.510	196.510	167.510
	Output (rviz)	waypoint_marker_publisher	1.35	275.1	297.820	197.820	168.820
		vehicle_sender	0.83	9.48	299.170	199.170	170.170
Summary		SUM	93.9	483.2			
Parameters		S			129.994	29.994	0.994
		γ			0.83	0.83	0.83
Memory Reservation					9.76	9.76	9.76
Deadline Miss (% over 1e09 steps)					0%	0%	0%

Calculated Parameters for Autoware

TABLE I: Components of Autoware sorted in order of execution, and their WCET, WCMC, and the value of γ for each component under various input and output periods (LCR value is for the first period). All times are in ms. **Different input and output periods**

Color	Category	Process	WCET(ms)	WCMC	Scenarios:	LCR(I=42,O=300)	LCR(I=33,O=200)	LCR(I=171,O=171)
						S_1	S_2	S_3
	LIDAR (Sensing)	velodyne_nodelet_manager	2.64	56.01		206.100	106.100	77.100
	Downsampling (Perception)	voxel_grid_filter	3.96	52.75		208.740	108.740	79.740
	Localizer (Perception)	ndt_matching	45.37	2063.05 (static)		212.700	112.700	83.700
		nmea2tfpose	20.7	13.32		258.070	158.070	129.070
	Positioning (Perception)	pose_relay (Positioning)		11.98				
		vel_relay (Velocity)	12.78	11.97		278.770	178.770	149.770
	Planning	waypoint_replanner	2.06	9.84		291.550	191.550	162.550
		lane_rule	1.2	28.29		293.610	193.610	164.610
		lane_stop	0.86	11.76		294.810	194.810	165.810
		lane_select	0.84	12.18		295.670	195.670	166.670
		pure_pursuit	1.31	10.09		296.510	196.510	167.510
	Output (rviz)	waypoint_marker_publisher	1.35	275.1		297.820	197.820	168.820
		vehicle_sender	0.83	9.48				
	Summary	SUM	93.9	483.2				
	Parameters	S				129.994	29.994	0.994
		γ				0.83	0.83	0.83
	Memory Reservation					9.76	9.76	9.76
	Deadline Miss (% over 1e09 steps)					0%	0%	0%

Values for S and γ

Calculated Parameters for Autoware

TABLE I: Components of Autoware sorted in order of execution, and their WCET, WCMC, and the value of γ for each component under various input and output periods (LCR value is for the first period). All times are in ms

Calculated LCR for each task

Color	Category	Process	WCET(ms)	WCMC(MB)	LCR(I=42,O=300)	LCR(I=33,O=200)	LCR(I=171,O=171)
	LIDAR (Sensing)	velodyne_nodelet_manager	2.64	56.01	206.100	106.100	77.100
	Downsampling (Perception)	voxel_grid_filter	3.96	52.75	208.740	108.740	79.740
	Localizer (Perception)	ndt_matching	45.37	2063.05 (static)	212.700	112.700	83.700
		nmea2tfpose	20.7	13.32	258.070	158.070	129.070
	Positioning (Perception)	pose_relay (Positioning)	12.78	11.98	278.770	178.770	149.770
		vel_relay (Velocity)		11.97			
	Planning	waypoint_replanner	2.06	9.84	291.550	191.550	162.550
		lane_rule	1.2	28.29	293.610	193.610	164.610
		lane_stop	0.86	11.76	294.810	194.810	165.810
		lane_select	0.84	12.18	295.670	195.670	166.670
		pure_pursuit	1.31	10.09	296.510	196.510	167.510
	Output (rviz)	waypoint_marker_publisher	1.35	275.1	297.820	197.820	168.820
		vehicle_sender	0.83	9.48	299.170	199.170	170.170
Summary		SUM	93.9	483.2			
Parameters		S			129.994	29.994	0.994
		γ			0.83	0.83	0.83
Memory Reservation					9.76	9.76	9.76
Deadline Miss (% over 1e09 steps)					0%	0%	0%

Response Time Analysis




Fig. 9: Response time under ResCue with the three scenarios.

Fig. 10: Response time under ResCue with no data scheduler.




Fig. 11: Response time under ResCue with reservation disabled.






Fig. 12: Response time with all features of ResCue disabled.

Overhead



Memory consumption.



Execution time.

- ❖ Overhead is relatively negligible compared to savings

Conclusion

- ❖ LET cannot be applied directly to integrated CPU/GPU architectures in Autonomous Embedded Systems
 - ❖ Memory is limited
 - ❖ Task release times can be sporadic
- ❖ We found inspiration from Heijunka, with goals similar to LET (i.e., meeting a constant output) but with a limited supply of parts and storage (e.g., limited disk transfer rate and small unified memory)
- ❖ We presented a data-aware expanded system model based on a combination of LET and Heijunka
- ❖ Based on our model, we presented ResCue, a runtime solution that provides:
 - ❖ Dynamic data scheduling to ensure temporal data availability (e.g., by calculating γ and S)
 - ❖ Dynamic memory reservation to ensure spatial data availability (e.g., by calculating the size of the buffer and handling the reservation)