


Amazon sagemaker debugger: A system for real-time insights into machine learning model training

Debugging in machine learning programs?

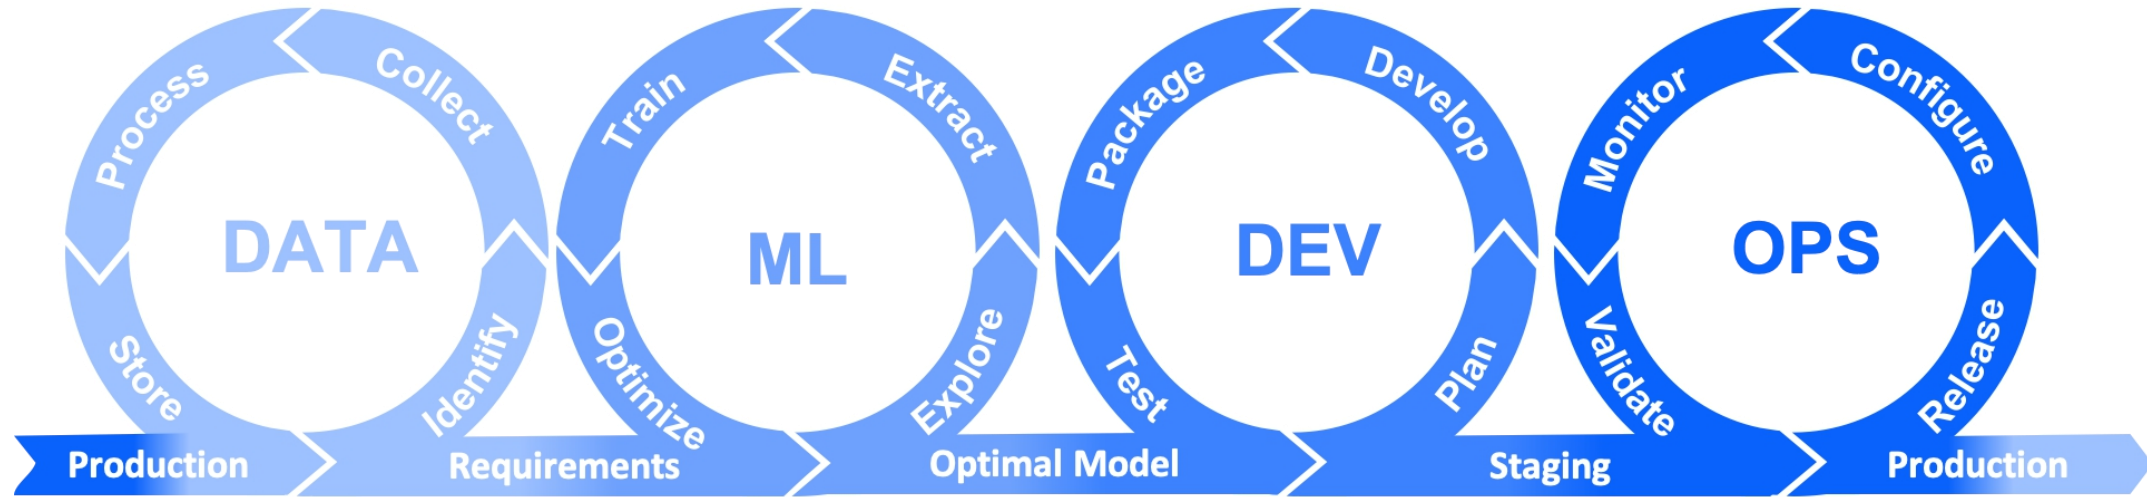


Debugging in machine learning programs?

- Bugs in traditional programs:
 - Segmentation fault
 - Division by zero
 - Business logic related error
 - ...
 - Symptom: Program crash / Error code (Exception thrown) ...
- Bugs in machine learning programs:
 - Low model capacity
 - Bad hyperparameter settings
 - Biased training data
 - Numerically unstable operations
 - ...
 - Symptom: Poor accuracy

Use rule-based method to detect common failure types in all stages of ML lifecycle.

Potential bugs in ML lifecycle




Data preparation: Data cleaning, Data pre-processing, Feature engineering

Data don't contain representative samples: Underfit the dataset

Data imbalance: Overfit on parts of the data

Data not normalized

Potential bugs in ML lifecycle



Model training: Different configurations and model architecture is applied.


The model has too few parameters: Cannot converge

The model has many parameters: Overfit

Inproper {initialization schema, optimizer settings, layer configurations, hyper-parameters}:

Cannot converge

Potential bugs in ML lifecycle

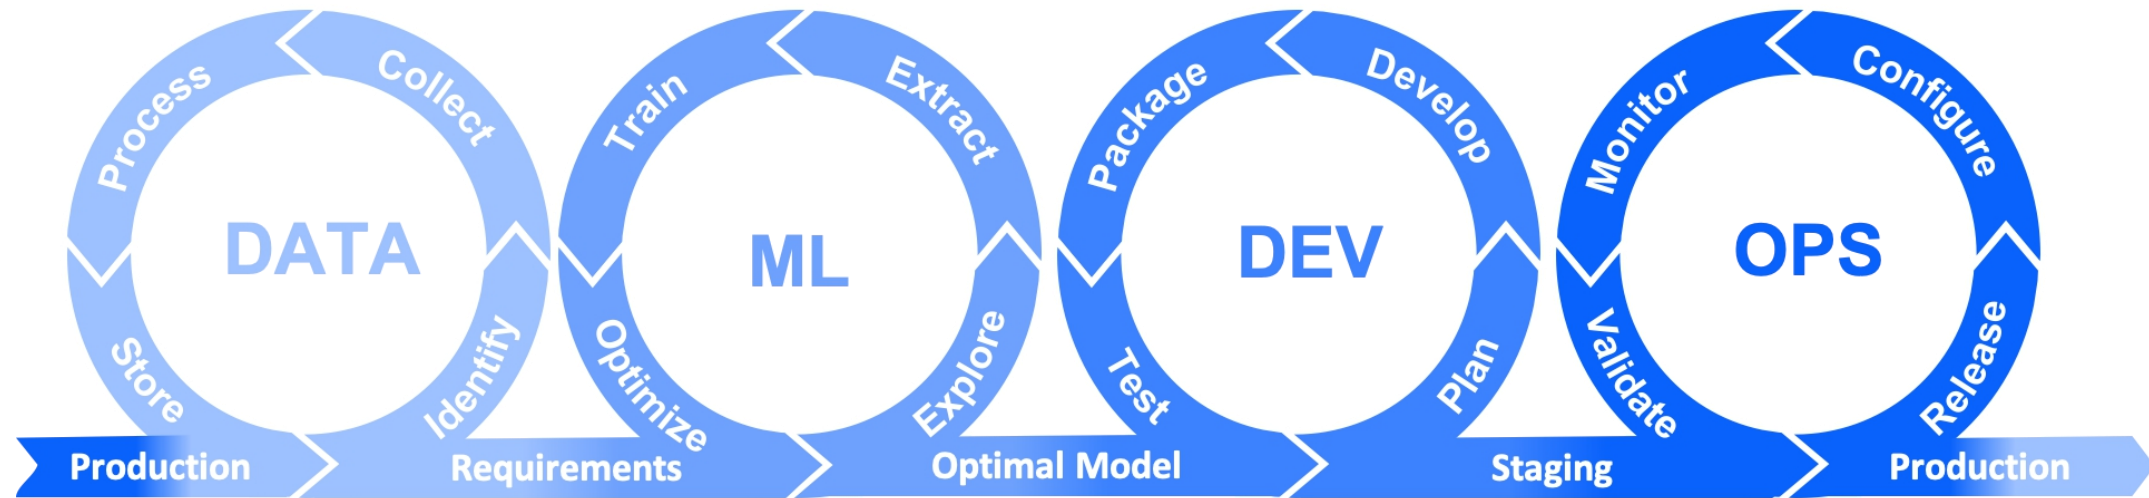


Hyperparameter tuning: Further refine a good model configuration

Non-optimal combination of hyper-parameters: The model has sub-optimal performance

ML model doesn't use full computation resources: Sub-optimal performance (a new feature that wasn't mentioned in the paper)

Potential bugs in ML lifecycle




Deployment: Deploy ML model into a computation cluster

Distribution of the data inference is significantly different from the distribution of training data:
Wrong prediction results

Rule-based bug-detection: Data preparation

- Problem: Data imbalance




Rule-based bug-detection: Data preparation

- Problem: Data imbalance

	Predicted label class 1	Predicted label class 2
True label class 1	correct true positive for class 1	wrong false positive for class 2
True label class 2	wrong false positive for class 1	correct true positive for class 2

$$\text{accuracy} = \frac{\text{orange} + \text{blue}}{\text{orange} + \text{yellow} + \text{blue} + \text{green}}$$

class 1 precision	$= \frac{\text{orange}}{\text{orange} + \text{yellow}}$	class 1 recall	$= \frac{\text{orange}}{\text{orange} + \text{green}}$
class 2 precision	$= \frac{\text{blue}}{\text{blue} + \text{green}}$	class 2 recall	$= \frac{\text{blue}}{\text{blue} + \text{yellow}}$



Rule-based bug-detection: Data preparation

- Solution to Data imbalance
Imbalance ratio+threshold

$$IR = \frac{\max_i \zeta_i}{\min_j \zeta_j}$$

Rule-based bug-detection: Data preparation


- Data not normalized

Feature Scaling

Idea: Make sure features are on a similar scale.


E.g. $x_1 = \text{size (0-2000 feet}^2\text{)}$ ←

$x_2 = \text{number of bedrooms (1-5)}$ ←



→ $x_1 = \frac{\text{size (feet}^2\text{)}}{2000}$ ↖

→ $x_2 = \frac{\text{number of bedrooms}}{5}$ ↖




Rule-based bug-detection: Data preparation

- Solution to Data not normalized
 - Check if the data has zero mean
 - Check if the data has zero variance


Rule-based bug-detection: Model training

- Activation function chosen improperly



Rule-based bug-detection: Model training


- Activation function chosen improperly



$$\Delta w = -\alpha \frac{\partial Loss}{\partial w}$$

Rule-based bug-detection: Model training

- Activation function chosen improperly
- Vanishing gradient (Activation function saturation), Gradient explode (Gradient too large)



$$\frac{\partial C}{\partial b_1} = \sigma'(z_1) \times w_2 \times \sigma'(z_2) \times w_3 \times \sigma'(z_3) \times w_4 \times \sigma'(z_4) \times \frac{\partial C}{\partial a_4}$$

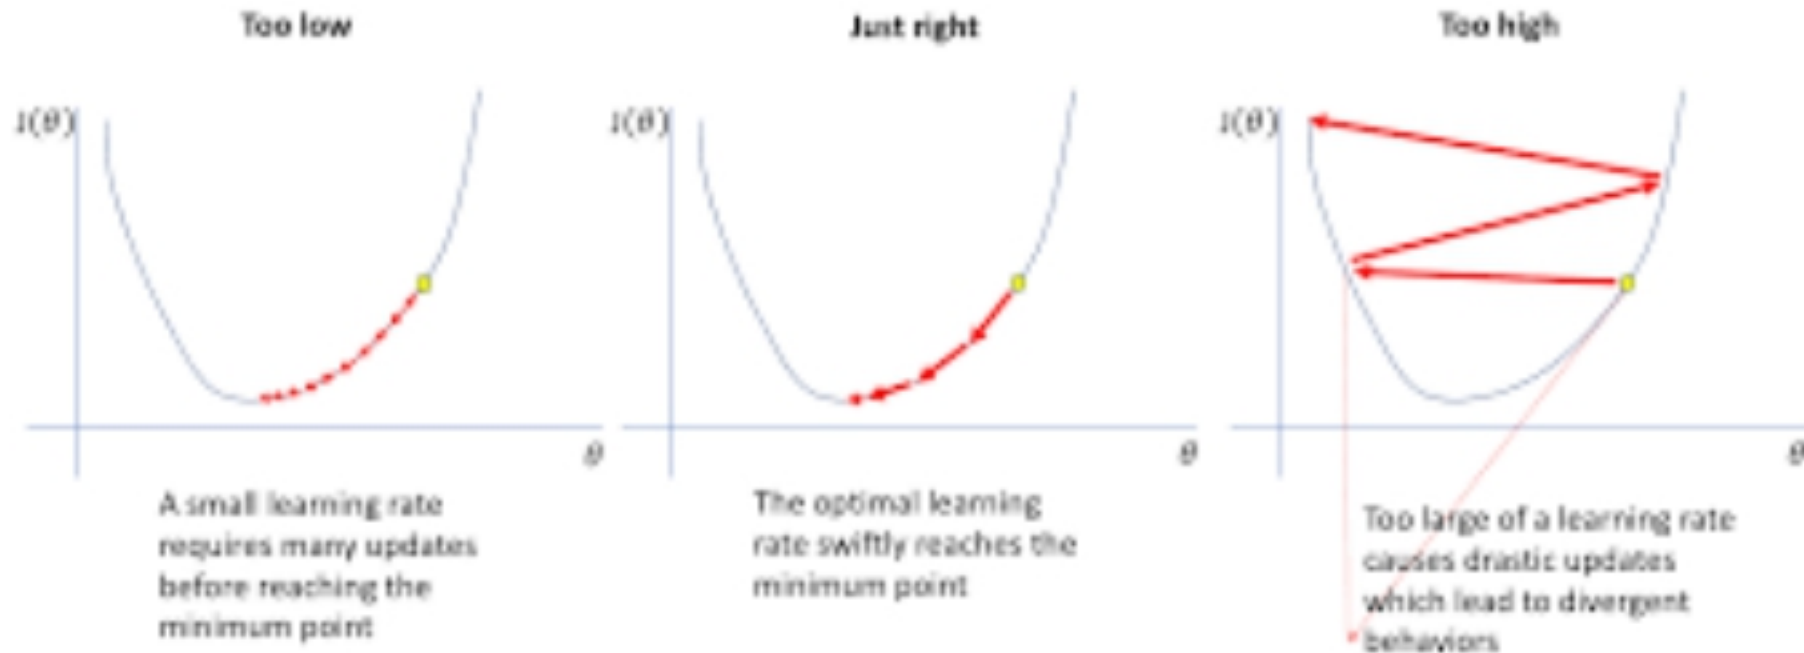


Rule-based bug-detection: Model training

- Solution to Gradient Vanishing: Retrieve activation outputs across steps and determine how many neurons in a model output zero values
- Solution to Gradient Explode: Retrieve activation outputs across steps and determines how many neurons have very large gradients

Rule-based bug-detection: Model training

- Optimizer options not chosen properly




Rule-based bug-detection: Model training

- Solution to optimizer options not chosen properly:
 - Check if output changes too fast

Rule-based bug-detection: Model training

- Model won't converge/model already converge but are still training




Rule-based bug-detection: Model training

- Solution to model won't converge/model already converged but are still training
 - Stop training when loss exceed certain limit
 - Stop training when loss doesn't change very often

Rule-based bug-detection: Hyperparameter tuning

- Parameter initialization wrong
 - eg: Initialization all weights to zero



$$z_1 = w_{10} * x_0 + w_{11} * x_1 + w_{12} * x_2 + w_{13} * x_3$$

$$z_2 = w_{20} * x_0 + w_{21} * x_1 + w_{22} * x_2 + w_{23} * x_3$$

Rule-based bug-detection: Hyperparameter tuning

- Solution to parameter initialization wrong
 - Check variance of output layers. A large variance may mean incorrect results.

Rule-based bug-detection: Tree model related

- Problem1: Tree model will overfit dataset when the tree depth will reach certain height.
- Problem2: Feature redundancy will occur if features are linearly dependent
- Solution:
 - Check tree depth
 - Check whether features are linearly dependent

System design

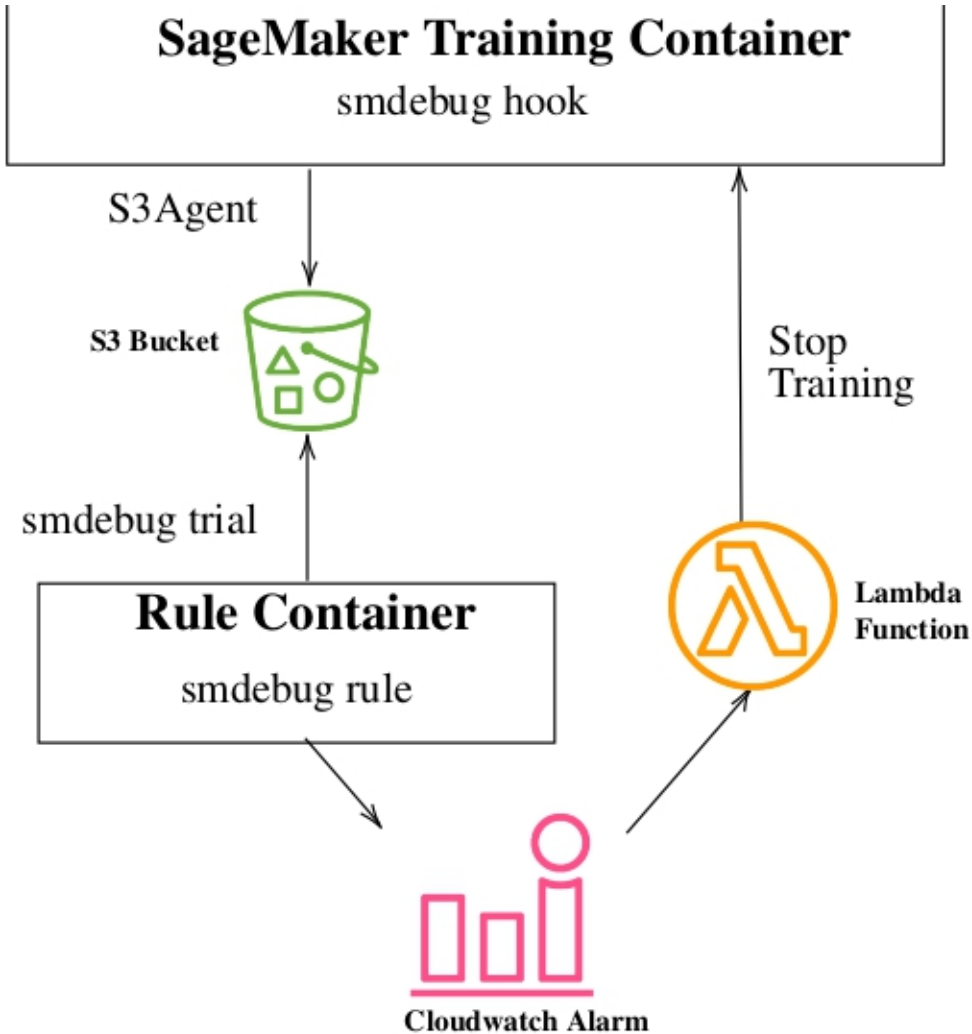


Figure 2. Debugger workflow

Technical challenges

- 1. Scale rule analysis by offloading into separate containers
- 2. Reduce overhead when recording and fetching tensors
- 3. Separate compute and storage and minimize impact on training

Case study

Check rule status

```
[ ]: 1 pytorch_estimator.latest_training_job.rule_job_summary()
```

Read Debugger data

```
[ ]: 1 ! pip install smdebug
```

```
[ ]: 1 from smdebug.trials import create_trial  
2  
3 path = pytorch_estimator.latest_job_debugger_artifacts_path()  
4 print('Tensors are stored in: {}'.format(path))  
5  
6 trial = create_trial(path)
```

Case study

Experiment: reinvent-demo

Trial: Trial-2020-11-15-203637-kciu

Trial Component Created: 5 days ago

Trial Component Status: Completed

Trial Component Debug Status: Fail

Training job detailed status: Completed

Trial components: Training ▼


Charts	Metrics	Parameters	Artifacts	AWS settings	Debugger	Trial Mappings
Status	Last modified	Rule name	Job ARN			
No Issues Found	5 days ago	ClassImbalance				
No Issues Found	5 days ago	DeadRelu				
No Issues Found	5 days ago	LossNotDecreasing				
Issues Found	5 days ago	LowGPUUtilization				
Issues Found	5 days ago	ProfilerReport				

Case study

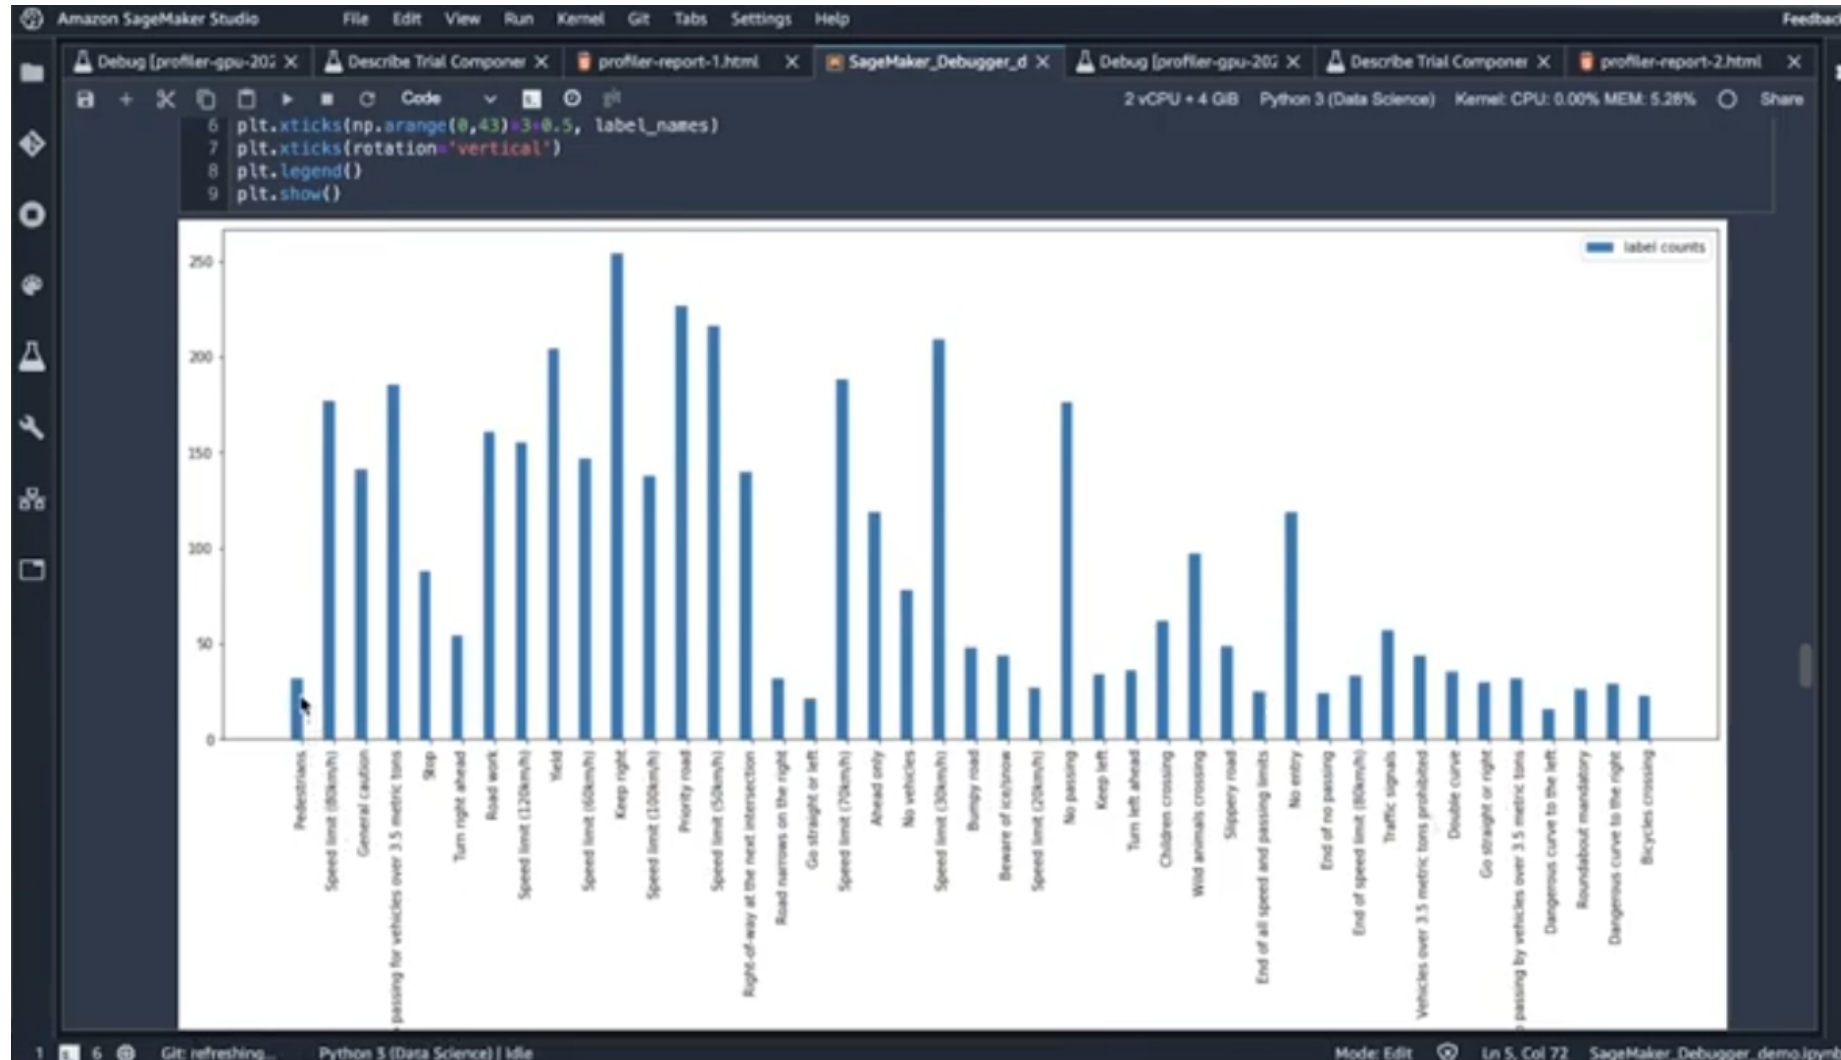
Investigate why 'Loss not decreasing' rule triggered

We can now easily visualize the loss values as training is still in progress.

```
[186]: 1 import matplotlib.pyplot as plt
2 from smdebug import modes
3
4 plt.ylabel('Train Loss')
5 plt.xlabel('Steps')
6 plt.plot(trial.steps(mode=modes.TRAIN),
7          list(trial.tensor('CrossEntropyLoss_output_0').values(mode=modes.TRAIN).values()))
8 plt.show()
```



Case study



Case study


```
6  
7     for prediction, label, image in zip(predictions, labels, images):  
8         if prediction != label:  
9             plot(image)  
10
```

Predicted: 'Turn left ahead' Groundtruth: 'Turn right ahead'

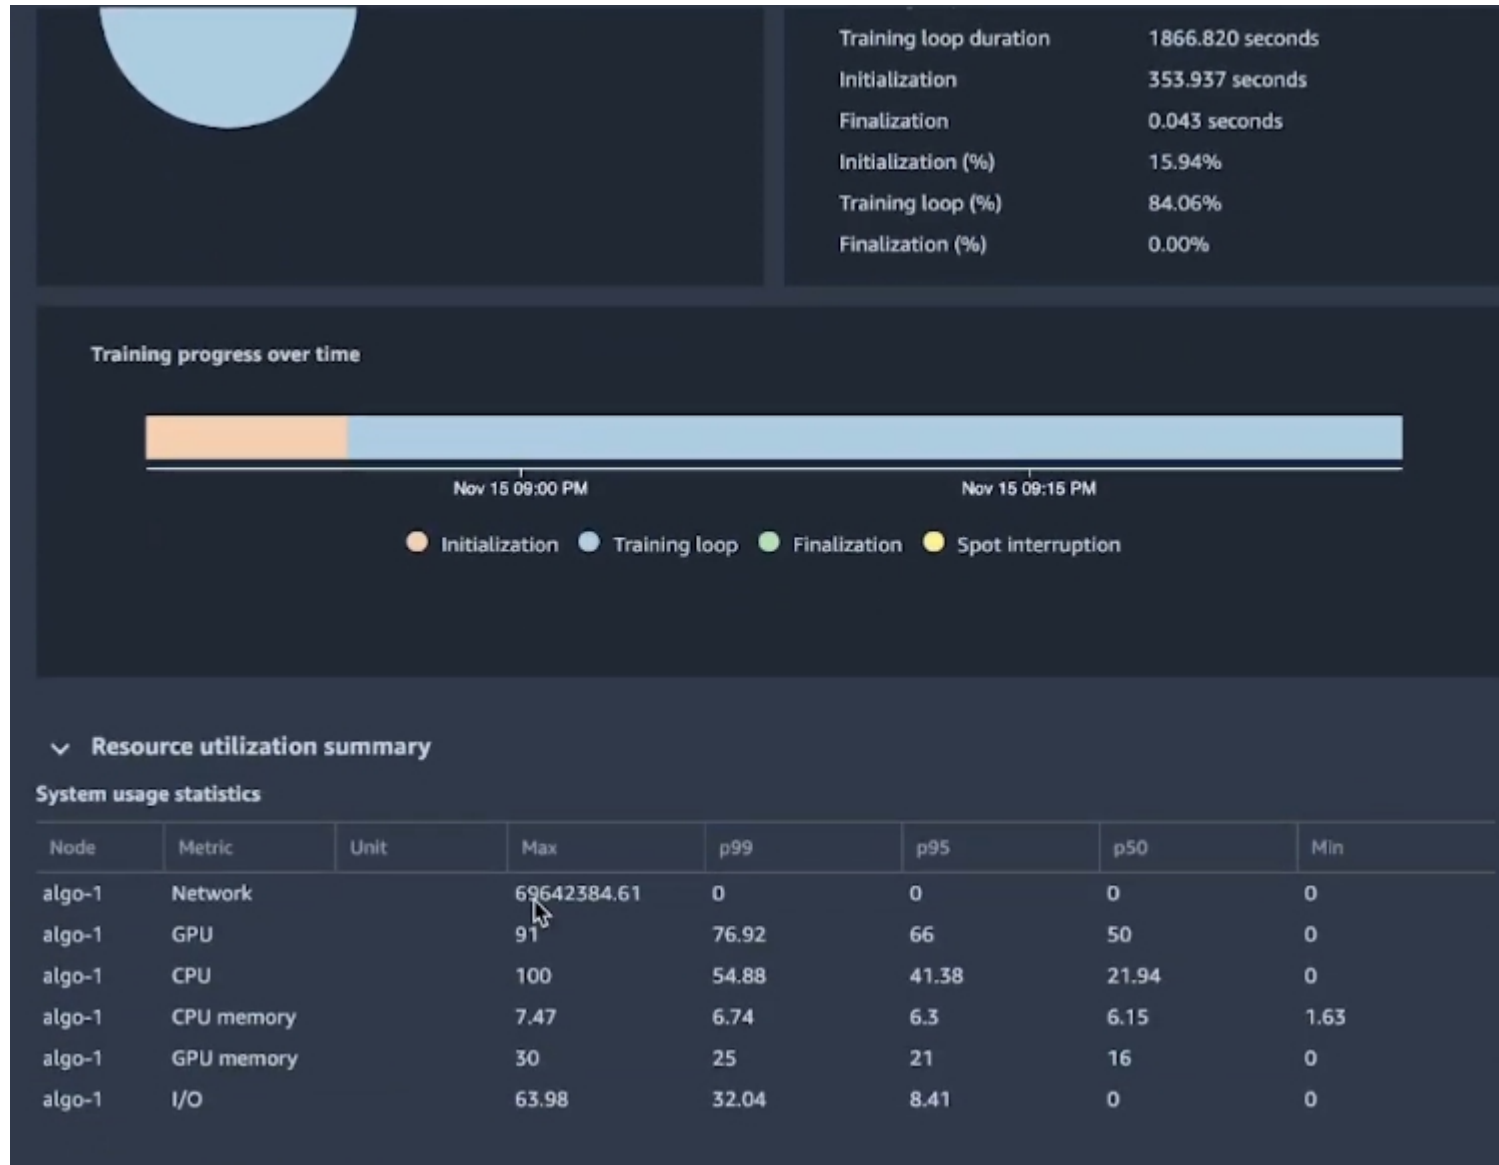


Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Predicted: 'Keep right' Groundtruth: 'Keep left'



Case study




Case study

▼ Resource intensive operations

Top operations on GPU

Percentage (%)	Cumulative time	GPU operator
22.81	209926	CudnnConvolutionBackward
22.74	209312	cudnn_convolution_backward
10.43	96020	conv2d
10.34	95166	convolution
10.24	94253	_convolution
9.97	91768	cudnn_convolution
3.82	35201	to
2.51	23121	CudnnBatchNormBackward
2.42	22239	batch_norm
2.38	21903	cudnn_batch_norm_backward
2.34	21542	_batch_norm_impl_index

Case study

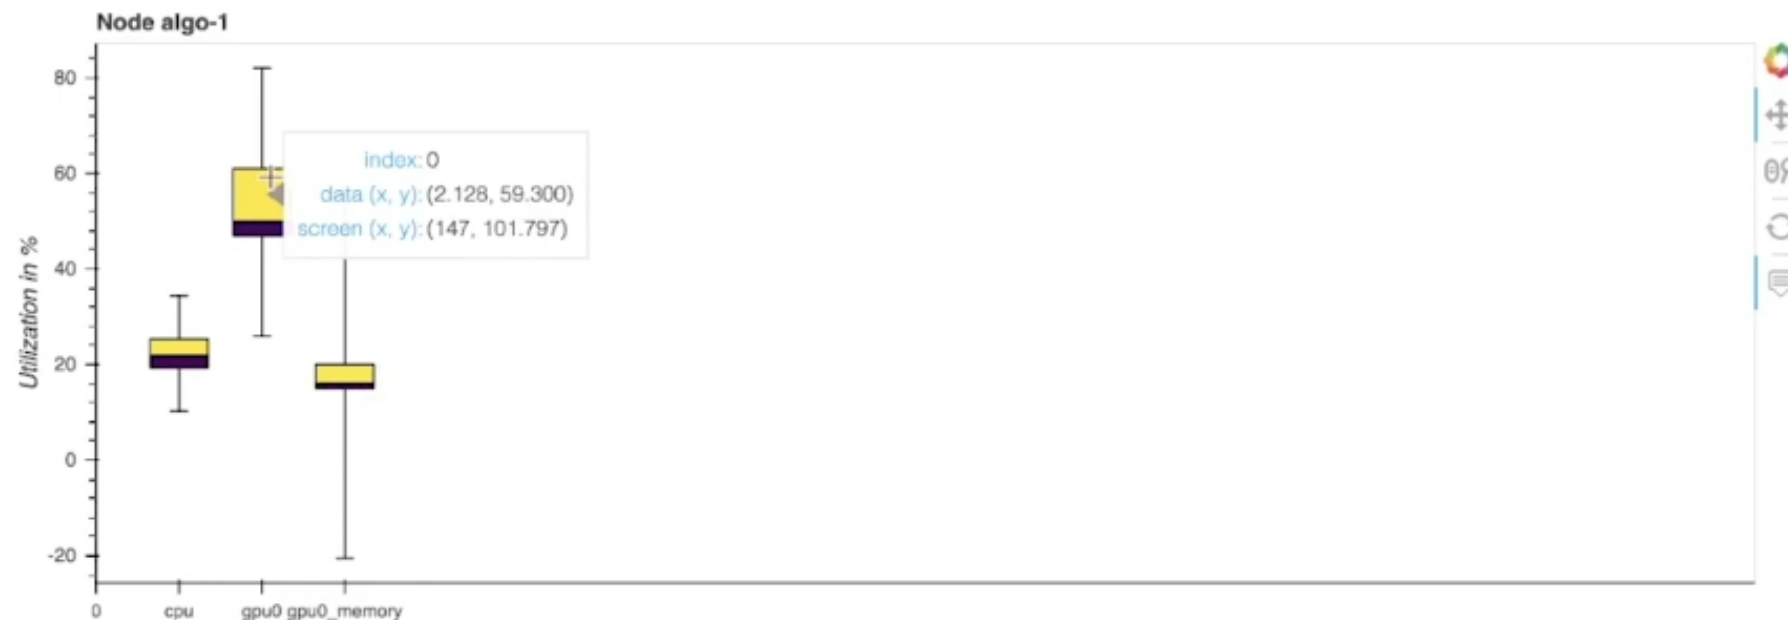


Case study

Batch size

The BatchSize rule helps to detect if GPU is underutilized because of the batch size being too small. To detect this the rule analyzes the GPU memory footprint, CPU and GPU utilization. The rule checked if the 95th percentile of CPU utilization is below `cpu_threshold_p95` of 70%, the 95th percentile of GPU utilization is below `gpu_threshold_p95` of 70% and the 95th percentile of memory footprint below `gpu_memory_threshold_p95` of 70%. In your training job this happened 31 times. The rule skipped the first 4000 datapoints. The rule computed the percentiles over window size of 1000 continuous datapoints. The rule analysed 22208 datapoints and triggered 31 times.

Your training job is underutilizing the instance. You may want to consider either switch to a smaller instance type or to increase the batch size. The last time the BatchSize rule triggered in your training job was on 11/16/2020 at 04:48:00. The following boxplots are a snapshot from the timestamps. They the total CPU utilization, the GPU utilization, and the GPU memory usage per GPU (without outliers).



Case study

Rules summary

The following table shows a profiling summary of the Debugger built-in rules. The table is sorted by the rules that triggered the most frequently. During your training job, the BatchSize rule was the most frequently triggered. It processed 22208 datapoints and was triggered 31 times.

		Description	Recommendation	Number of times rule triggered	Number of datapoints	Rule parameters
BatchSize	Checks if GPUs are underutilized because the batch size is too small. To detect this problem, the rule analyzes the average GPU memory footprint, the CPU and the GPU utilization.	The batch size is too small, and GPUs are underutilized. Consider running on a smaller instance type or increasing the batch size.	31	22208	cpu_threshold_p95:70 gpu_memory_threshold_p95:70 patience:4000 window:1000	
LowGPUUtilization	Checks if the GPU utilization is low or fluctuating. This can happen due to bottlenecks, blocking calls for synchronizations, or a small batch size.	Check if there are bottlenecks, minimize blocking calls, change distributed training strategy, or increase the batch size.	31	22209	threshold_p95:70 threshold_p5:10 window:500 patience:4000	
StepOutlier	Detects outliers in step duration. The step duration for forward and backward pass should be roughly the same throughout the training. If there are significant outliers, it may indicate a system stall or bottleneck issues.	Check if there are any bottlenecks (CPU, I/O) correlated to the step outliers.	25	12519	threshold:3 mode:None n_outliers:10 stddev:3	
Dataloader	Checks how many data loaders are running in parallel and whether the total number is equal the number of available CPU cores. The rule triggers if number is much smaller or larger than the number of available cores. If too small, it might lead to low GPU utilization. If too large, it might impact other compute intensive operations on CPU.	Change the number of data loader processes.	1	84	min_threshold:40 max_threshold:200	
IOBottleneck	Checks if the data I/O wait time is high and the GPU utilization is low. It might indicate IO bottlenecks where GPU is waiting for data to arrive from storage. The rule evaluates the I/O and GPU utilization rates and triggers the issue if the time spent on I/O is high.	Pre-fetch data or choose different file formats, such as binary formats that improve I/O	0	22249	threshold:50 io_threshold:50 cpu_threshold:10	

Performance of recording tensors

